

Apache Kafka

Daniel Hinojosa



About Me...

Daniel Hinojosa

Programmer, Consultant, Trainer

Testing in Scala (Book)

Beginning Scala Programming (Video)

Scala Beyond the Basics (Video)

Contact:

dhinojosa@evolutionnext.com

@dhinojosa

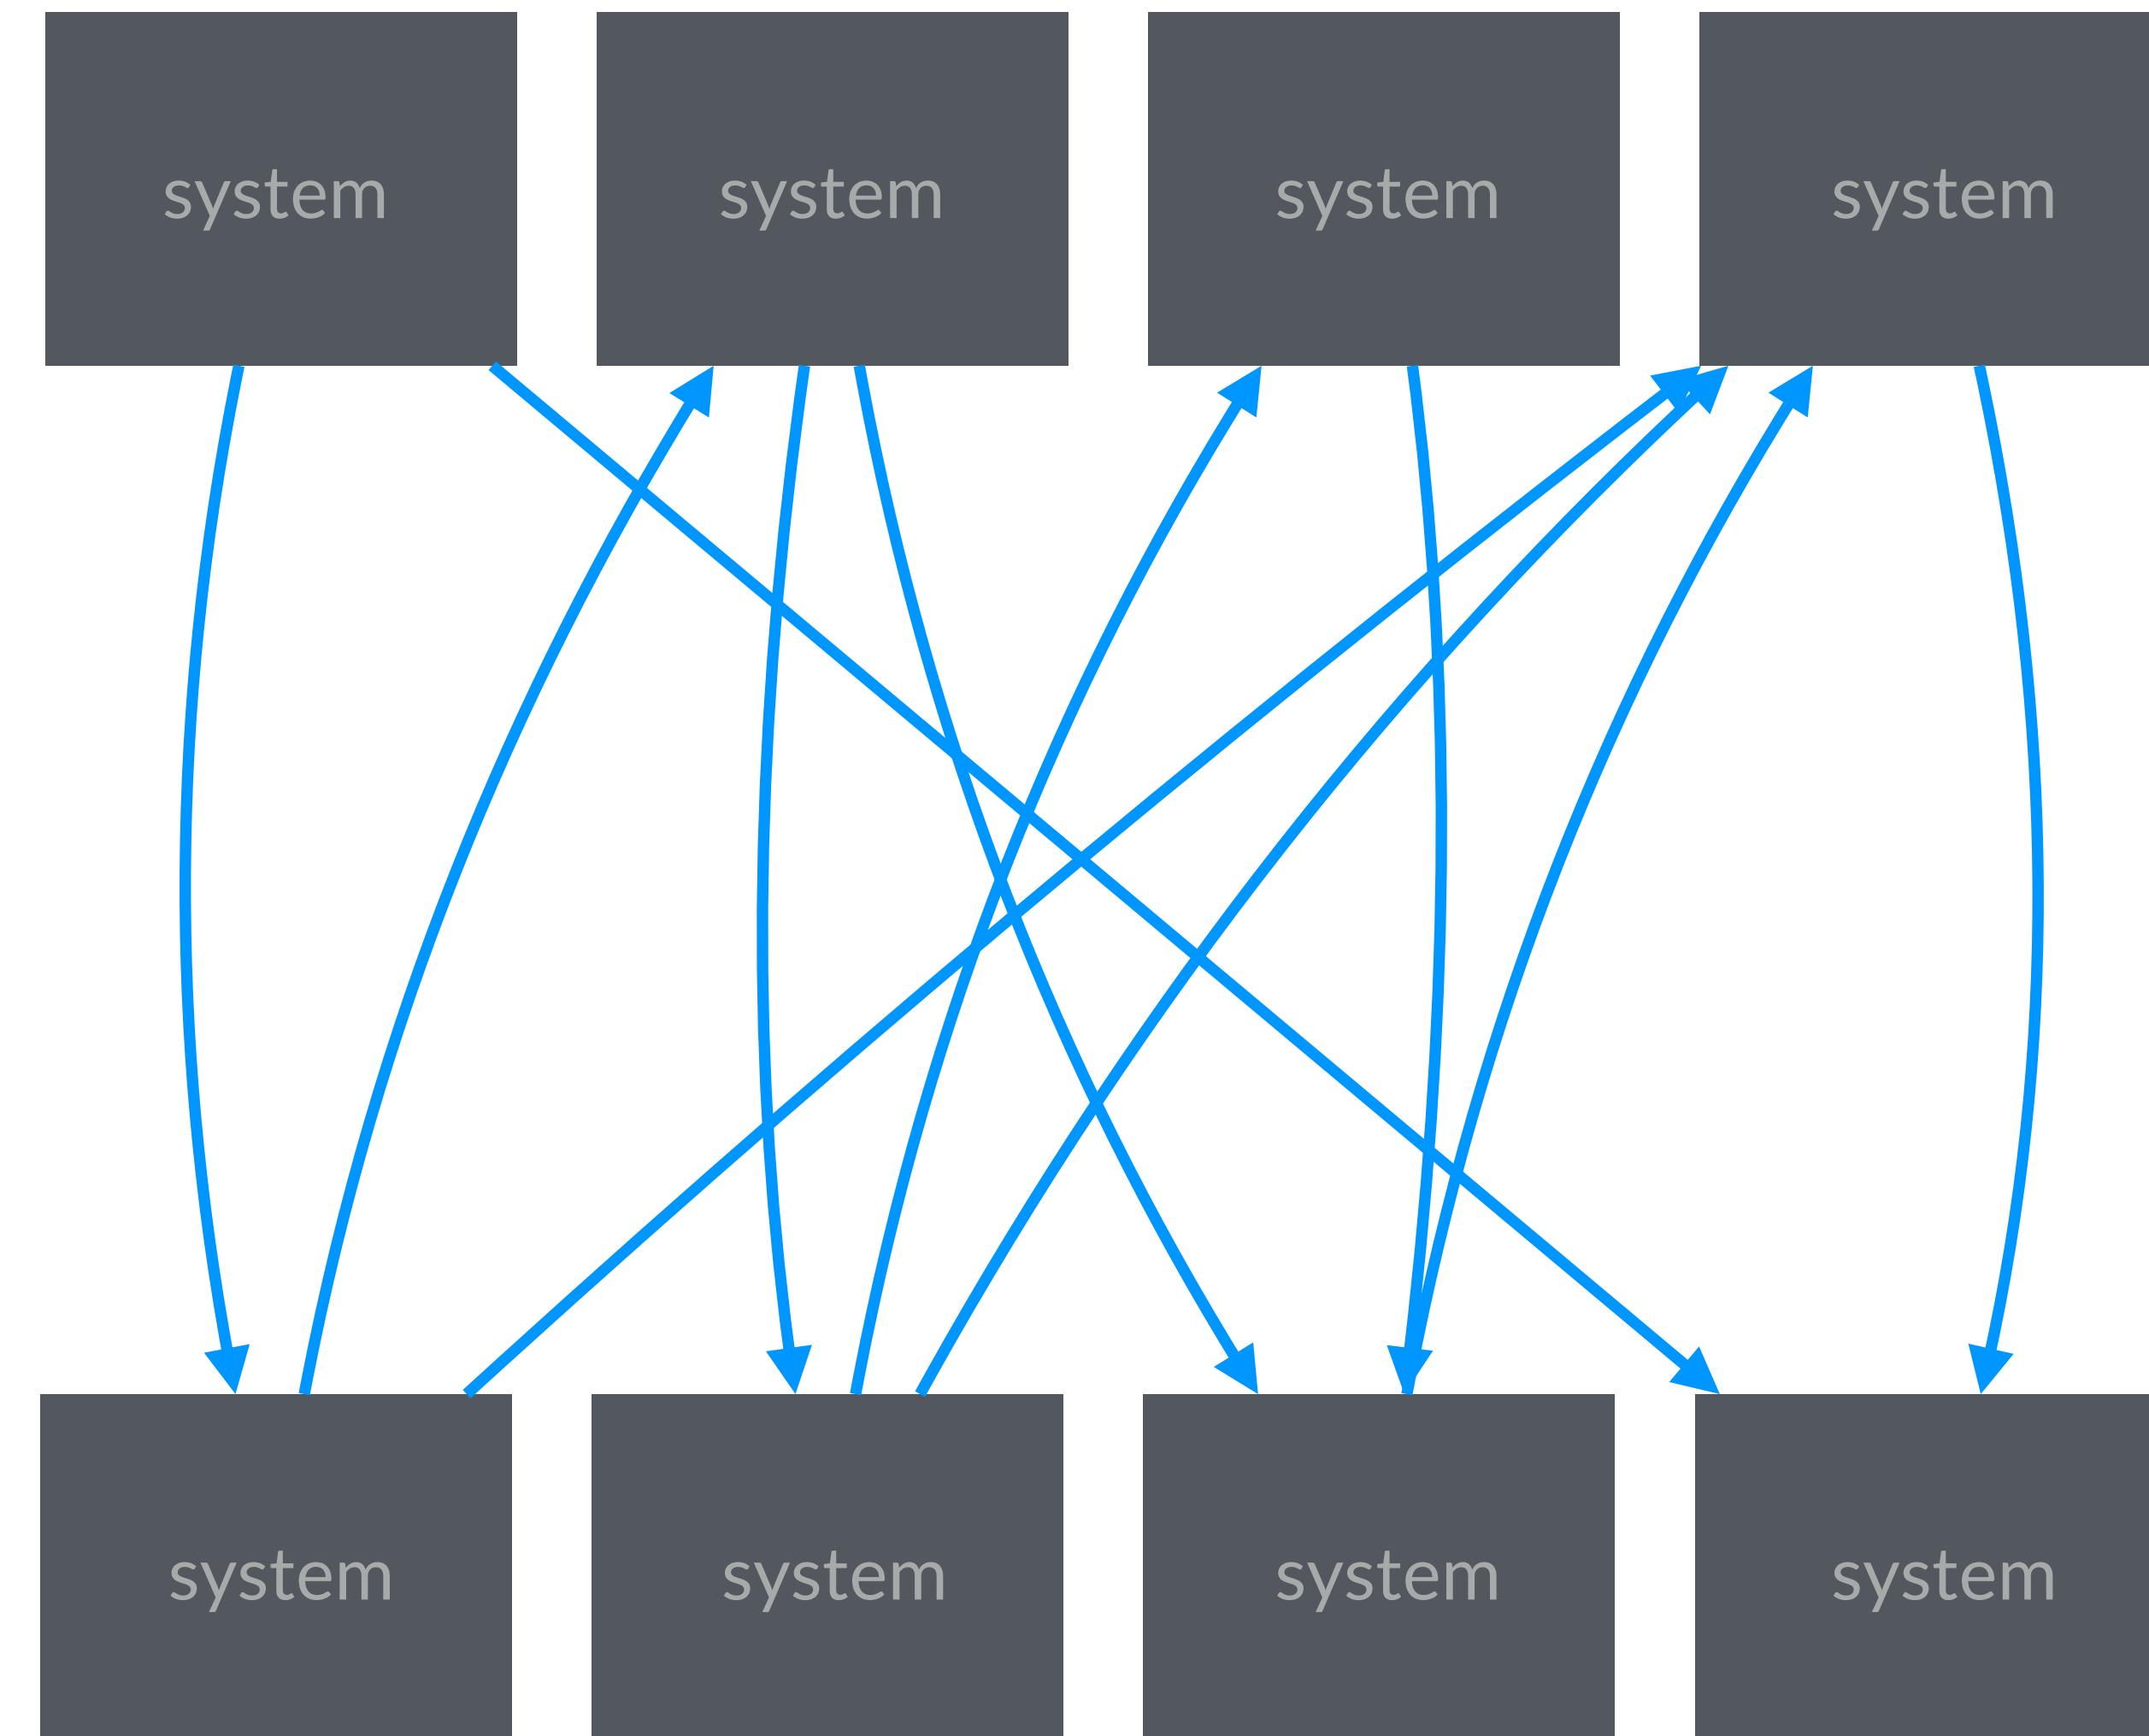


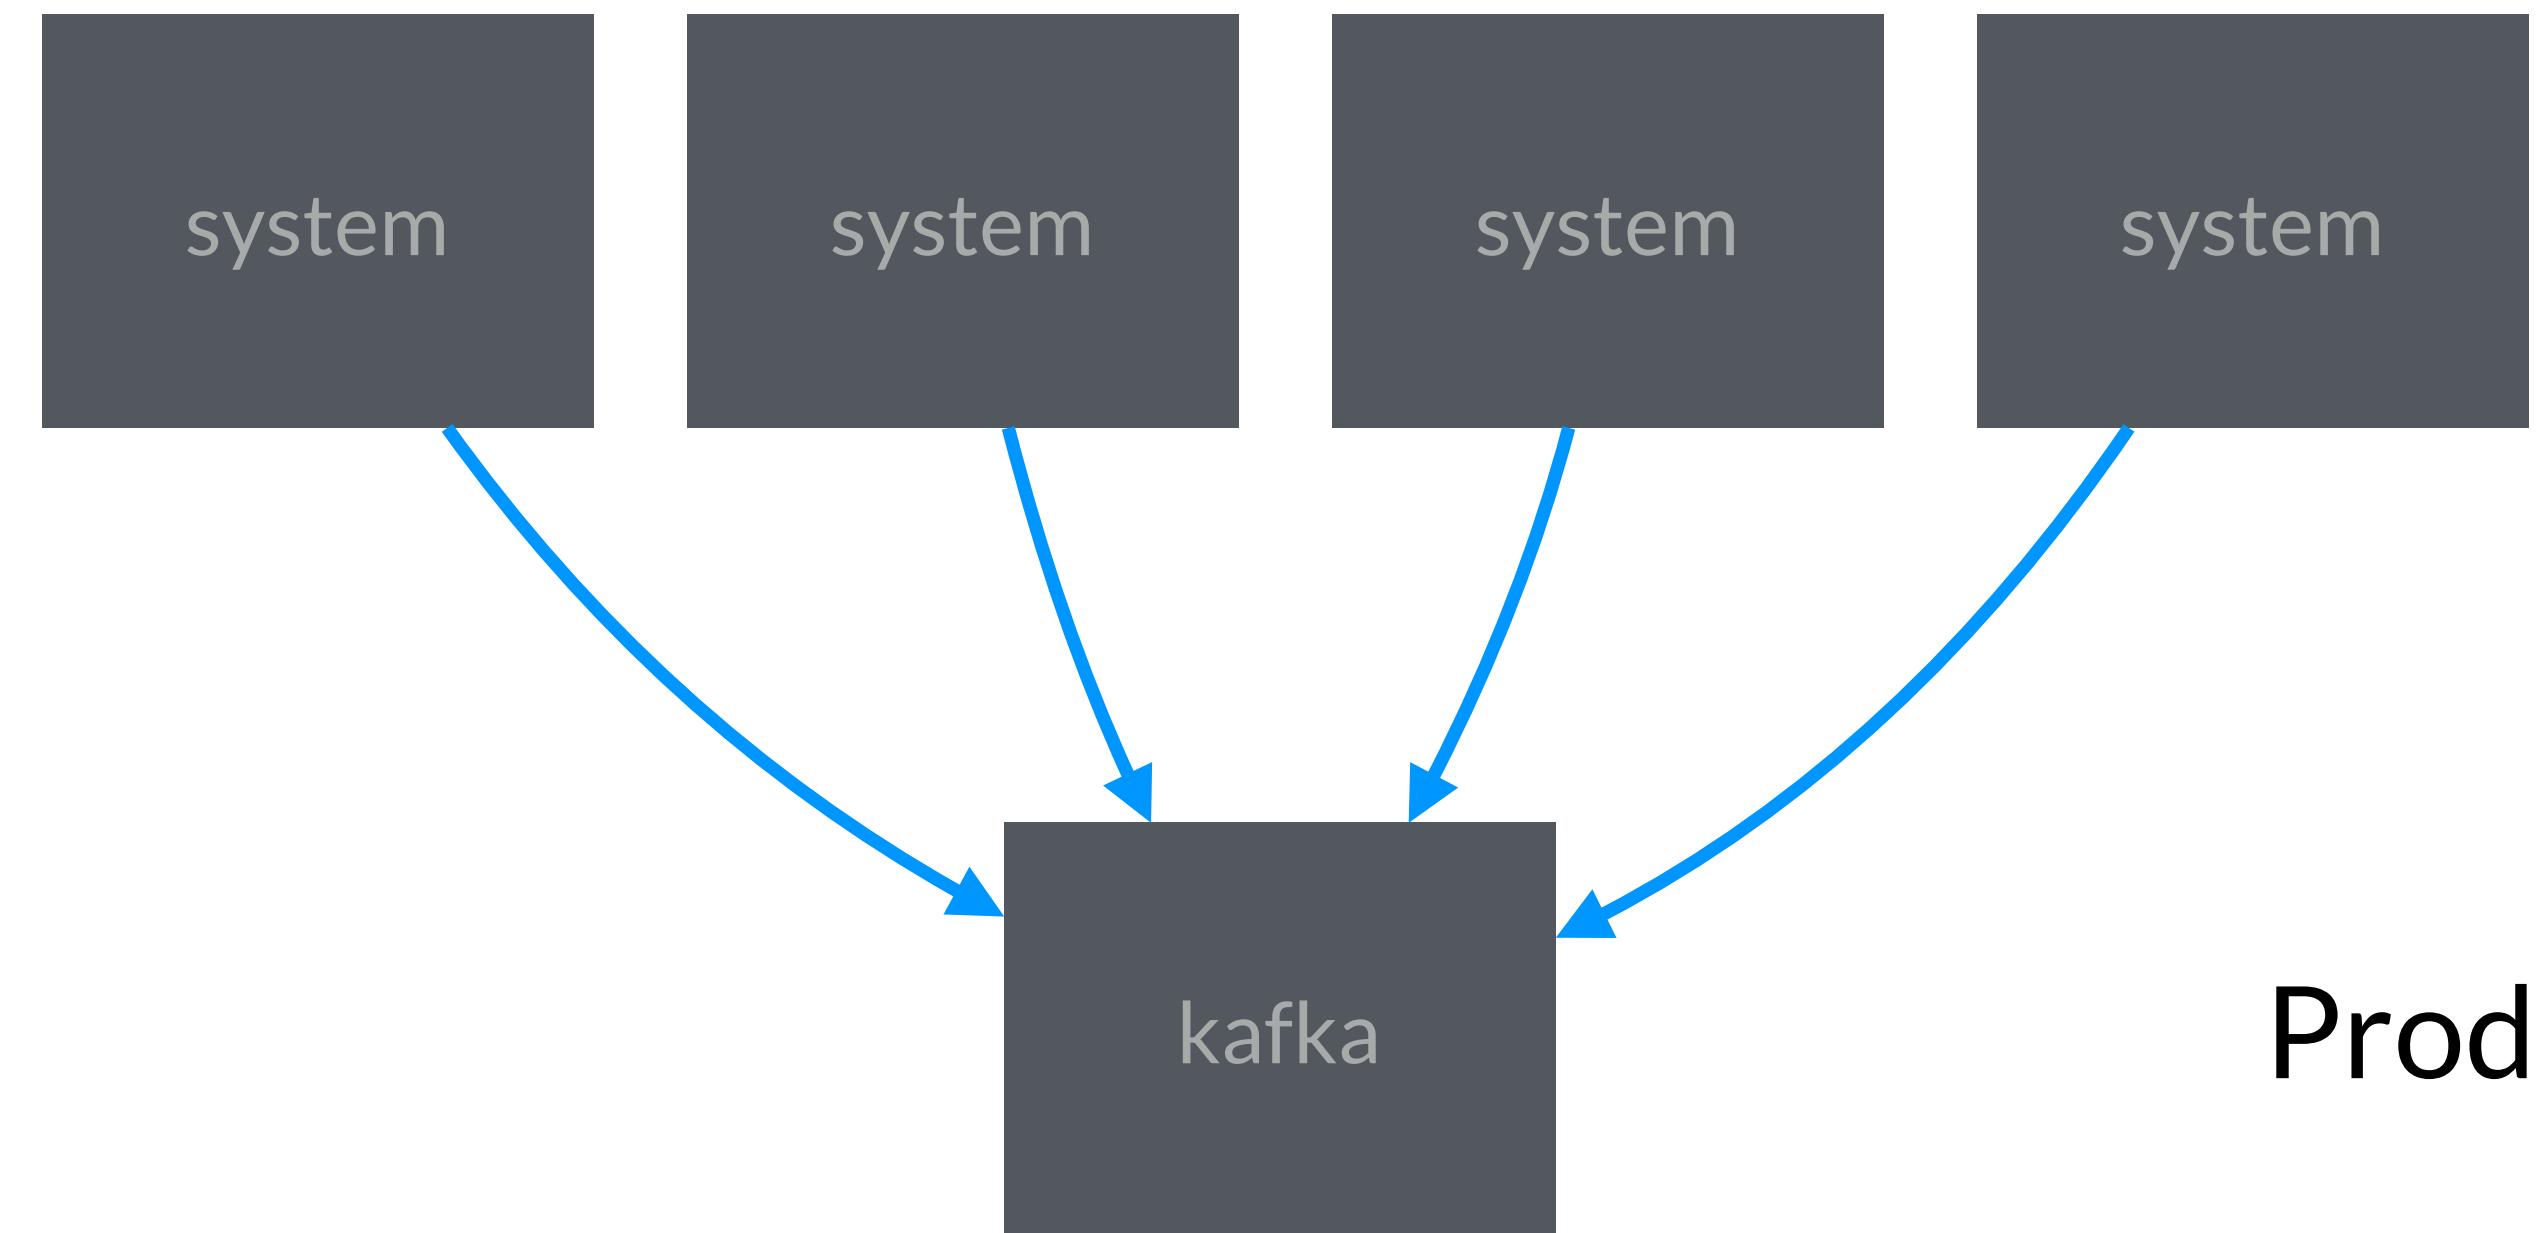
Agenda

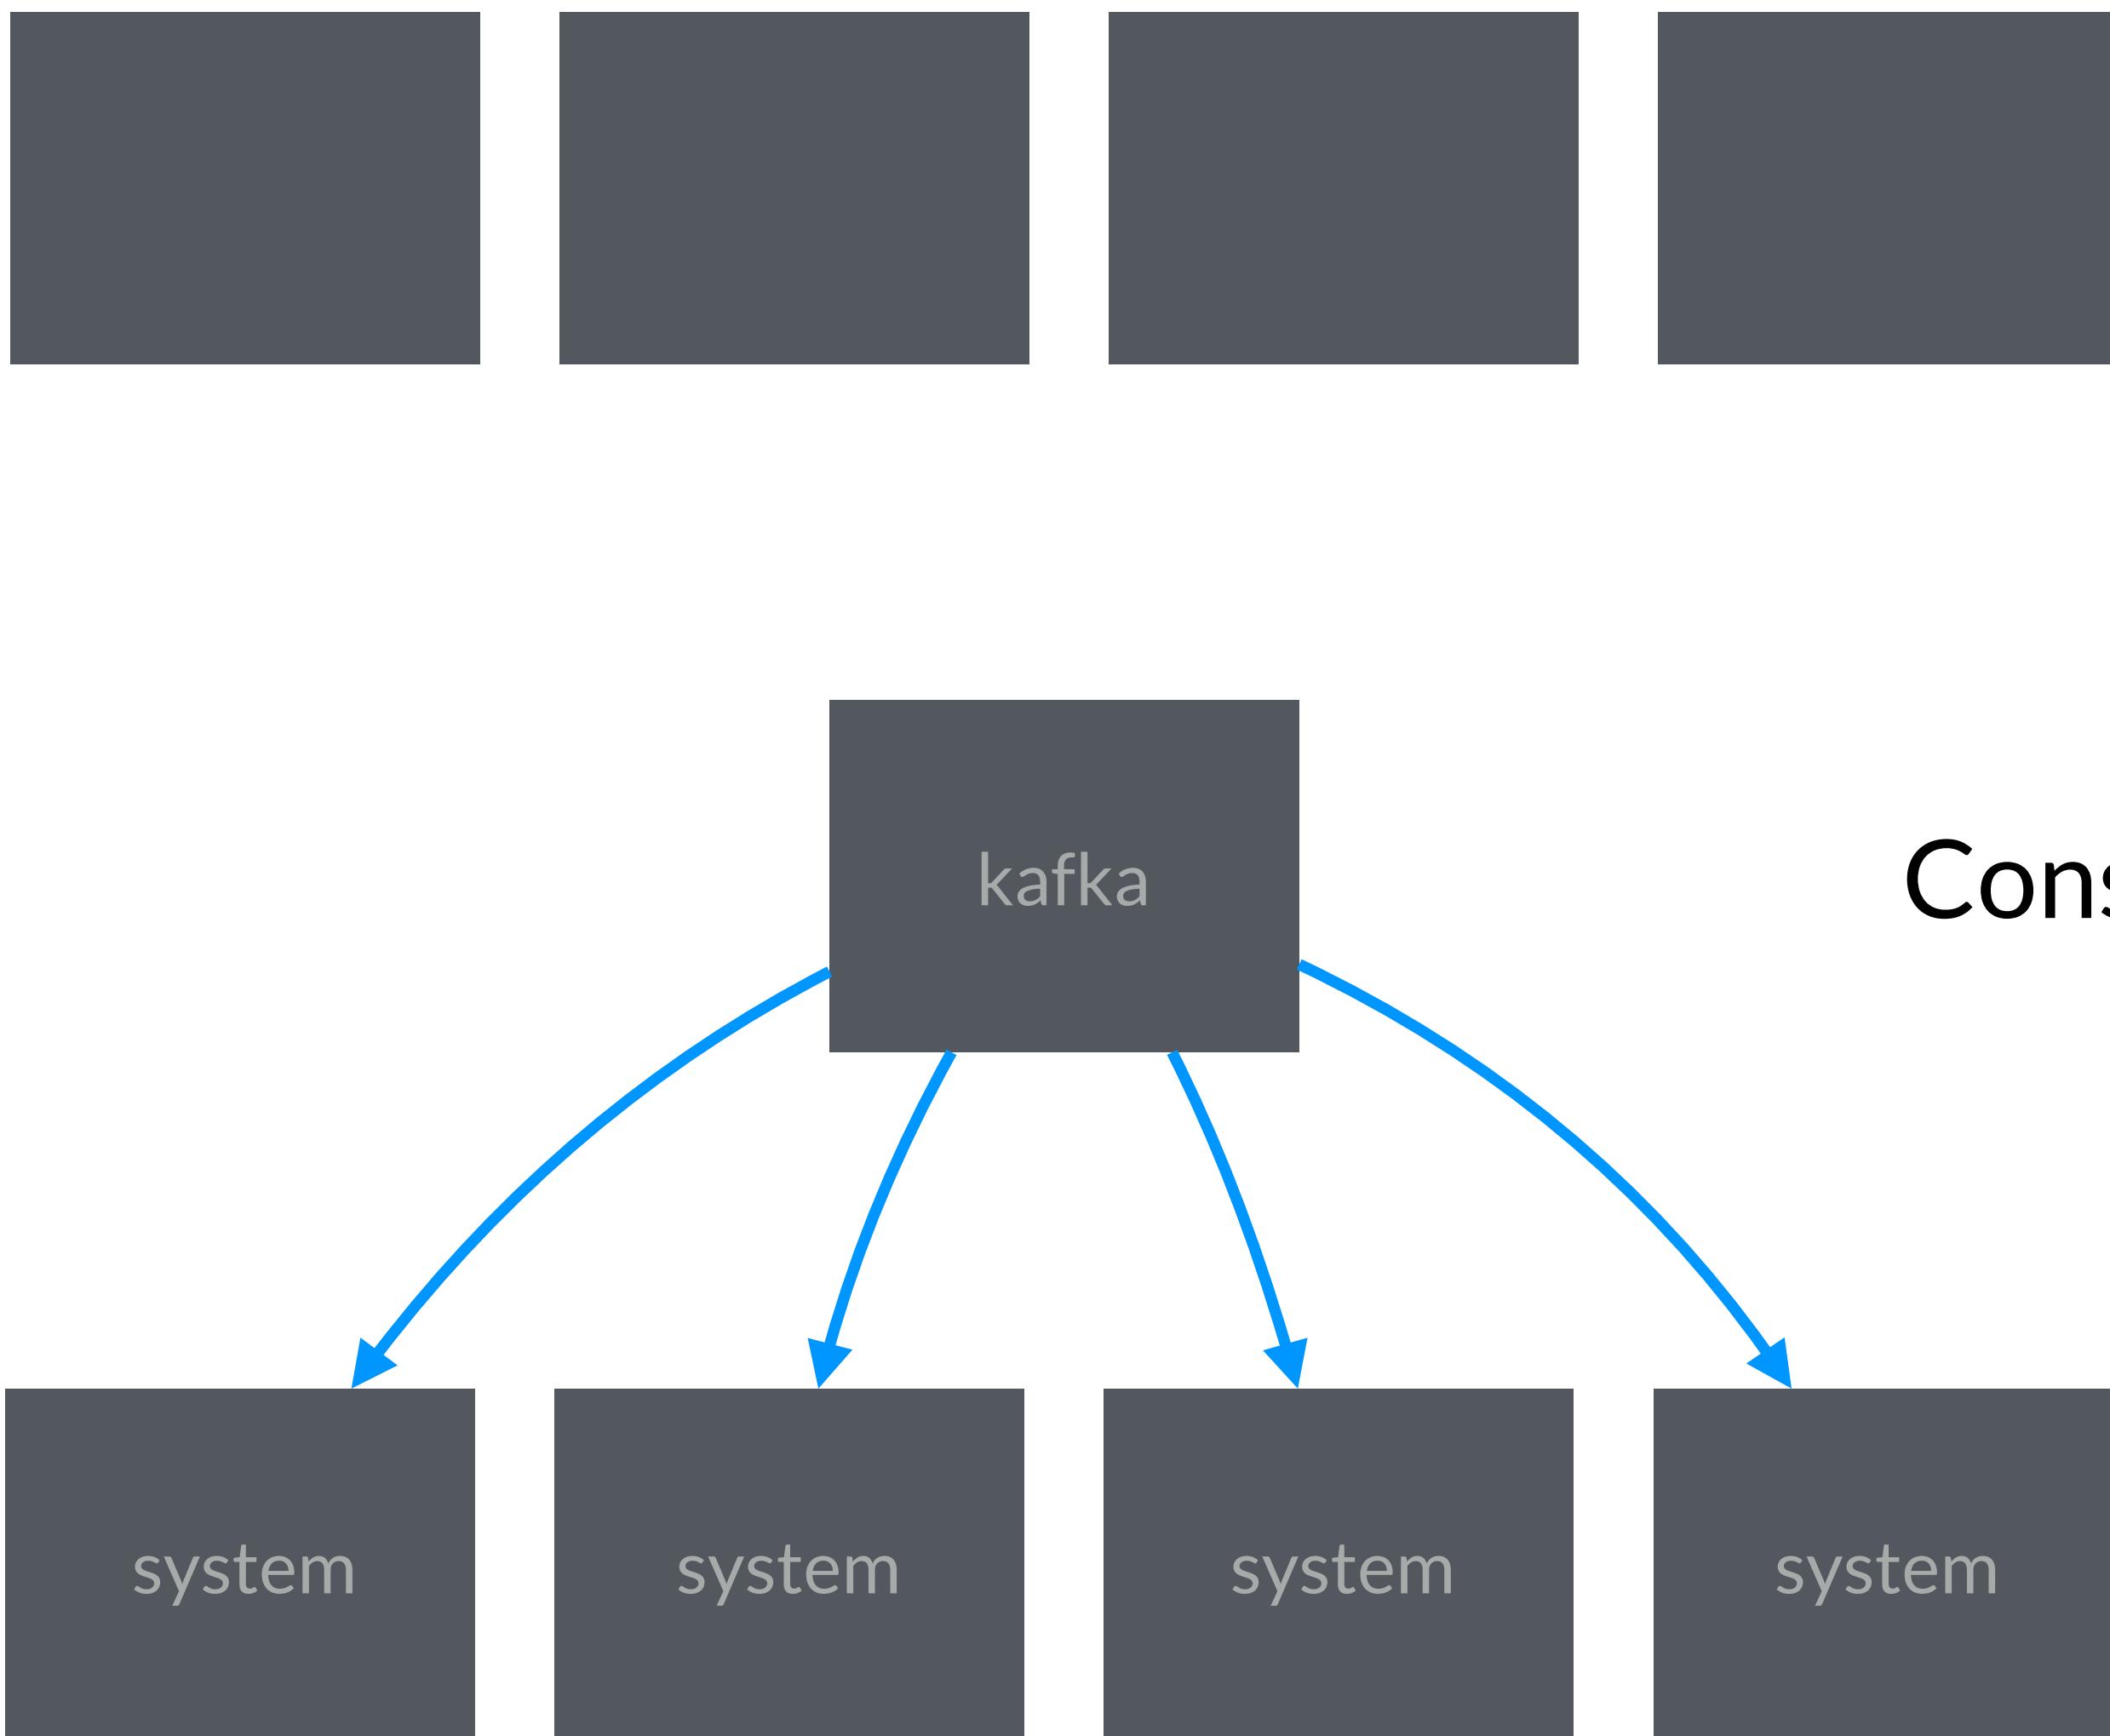
Agenda

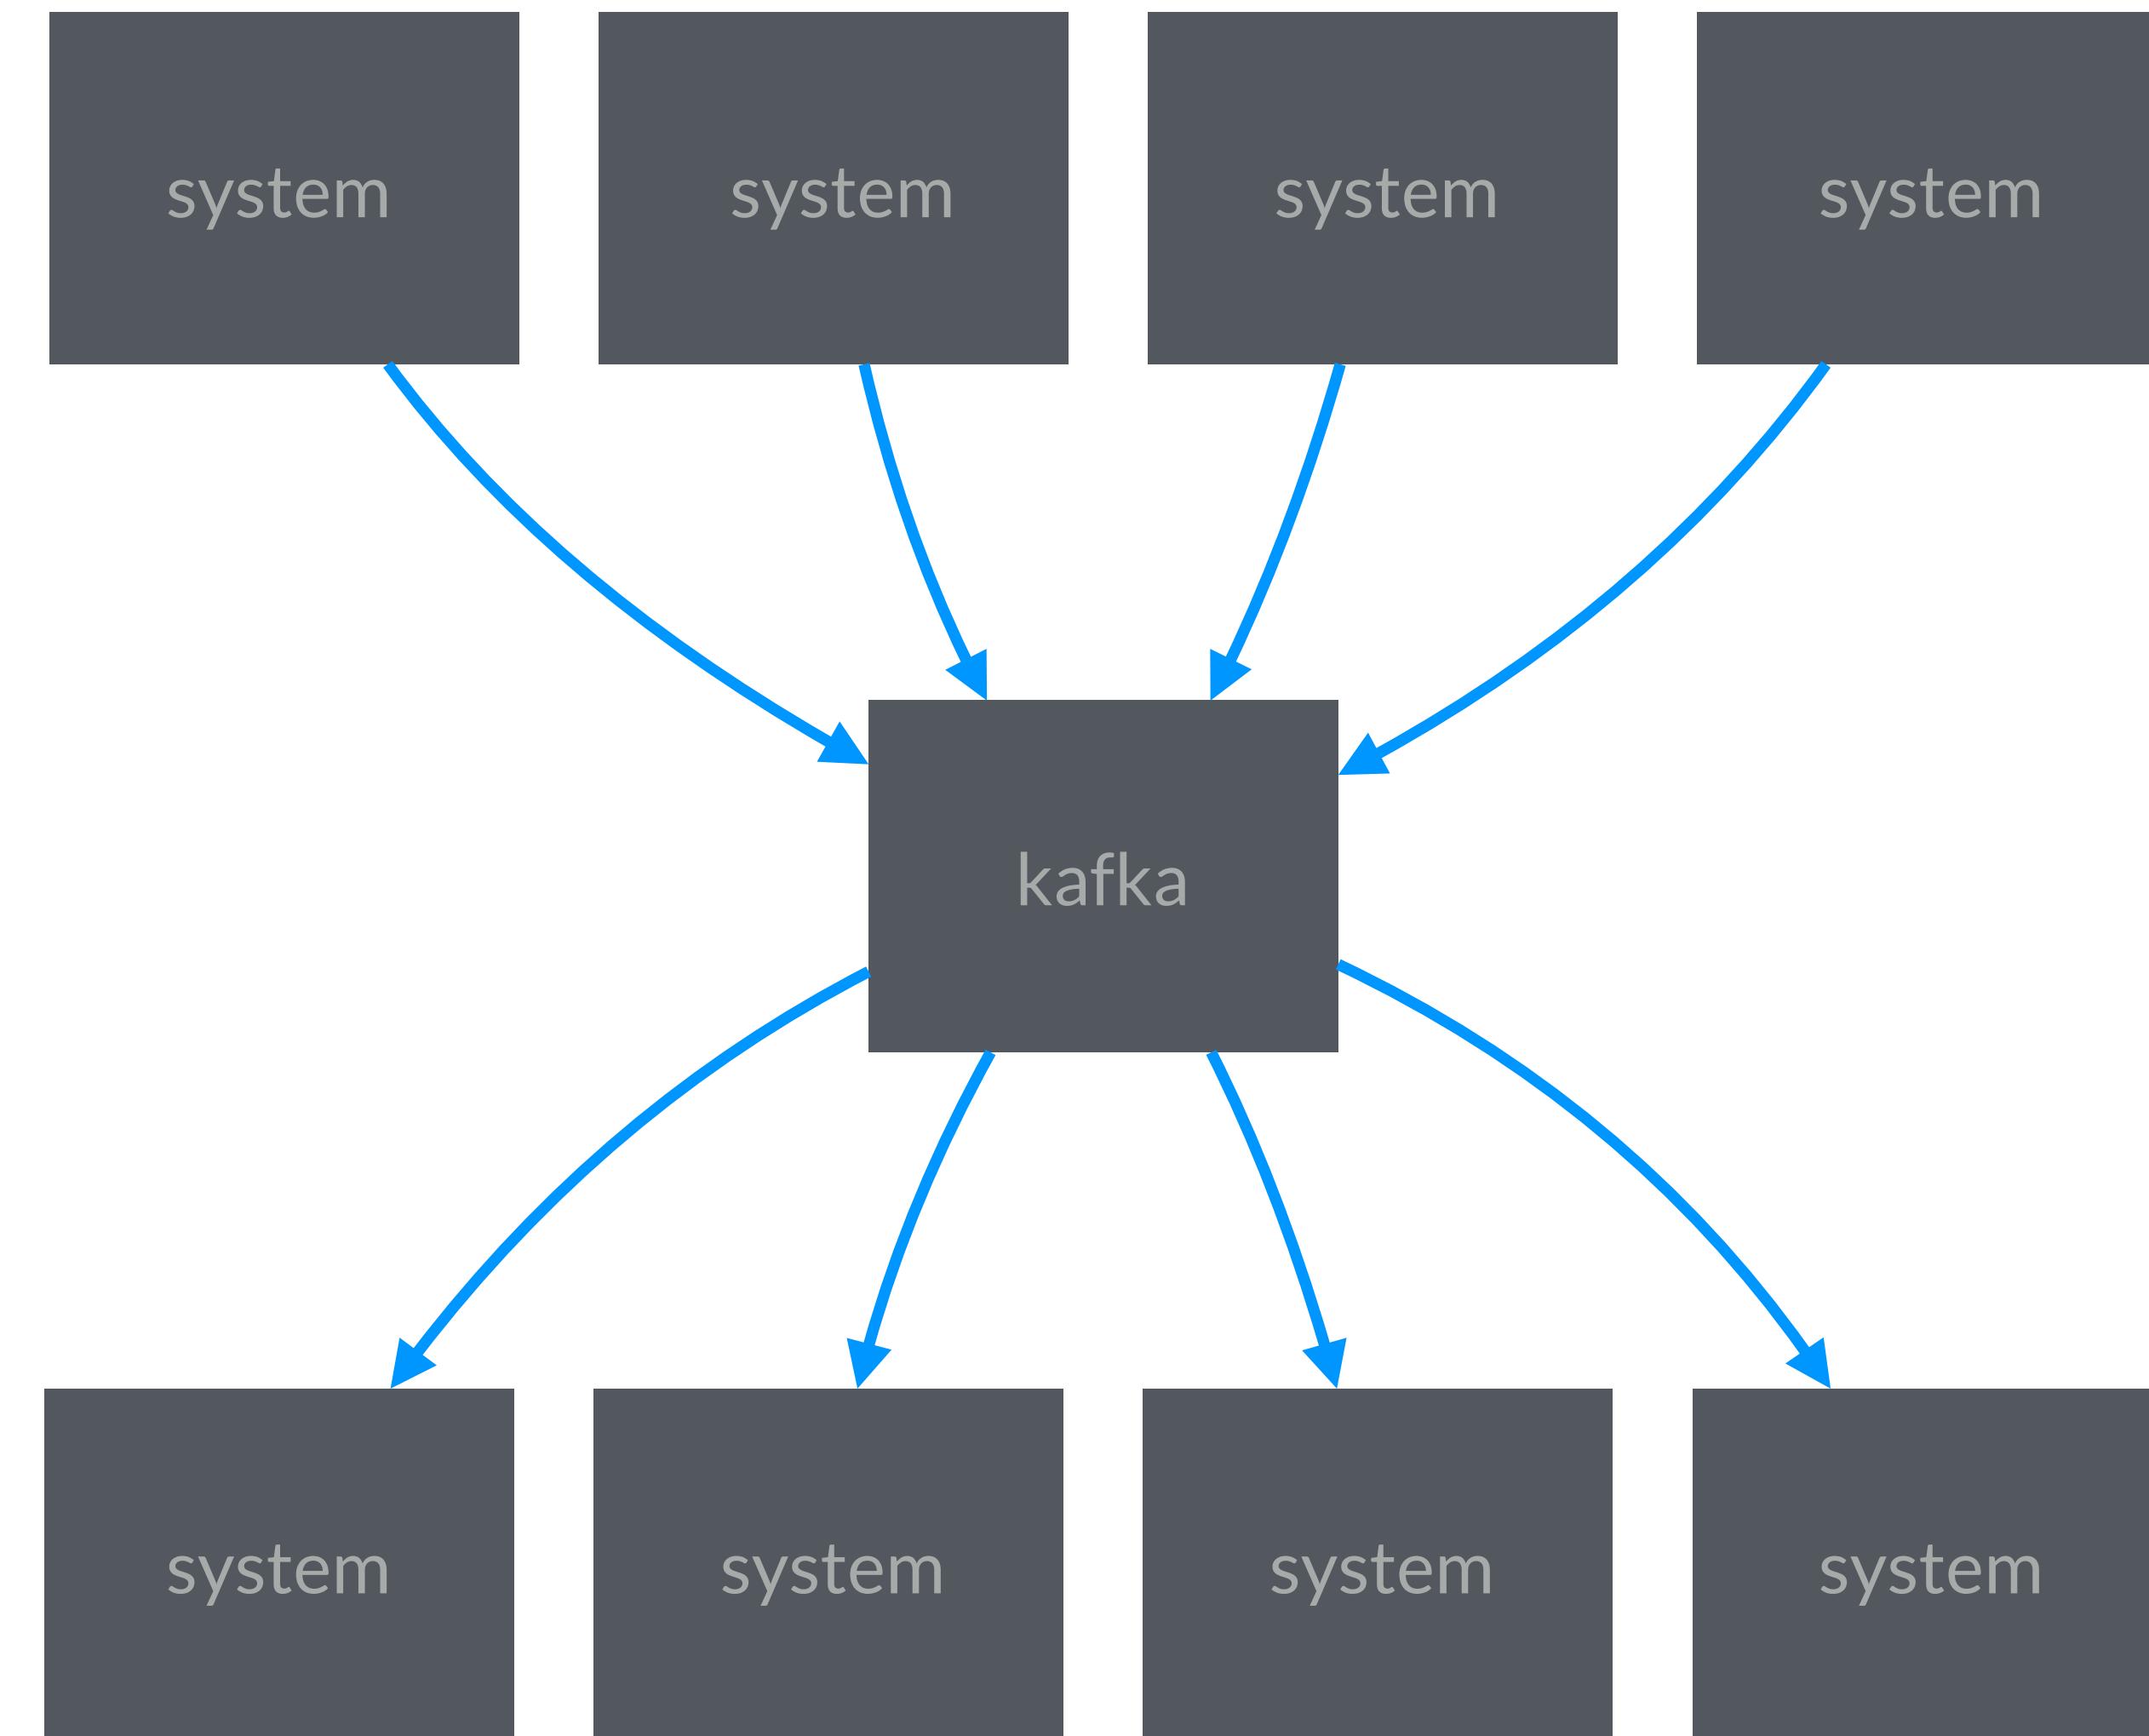
- ★ Understand Kafka
- ★ Understand Producer
- ★ Understand Consumer
- ★ Understand Rebalancing
- ★ Understand Streaming
- ★ Understand Tables
- ★ Understand KSQLDB

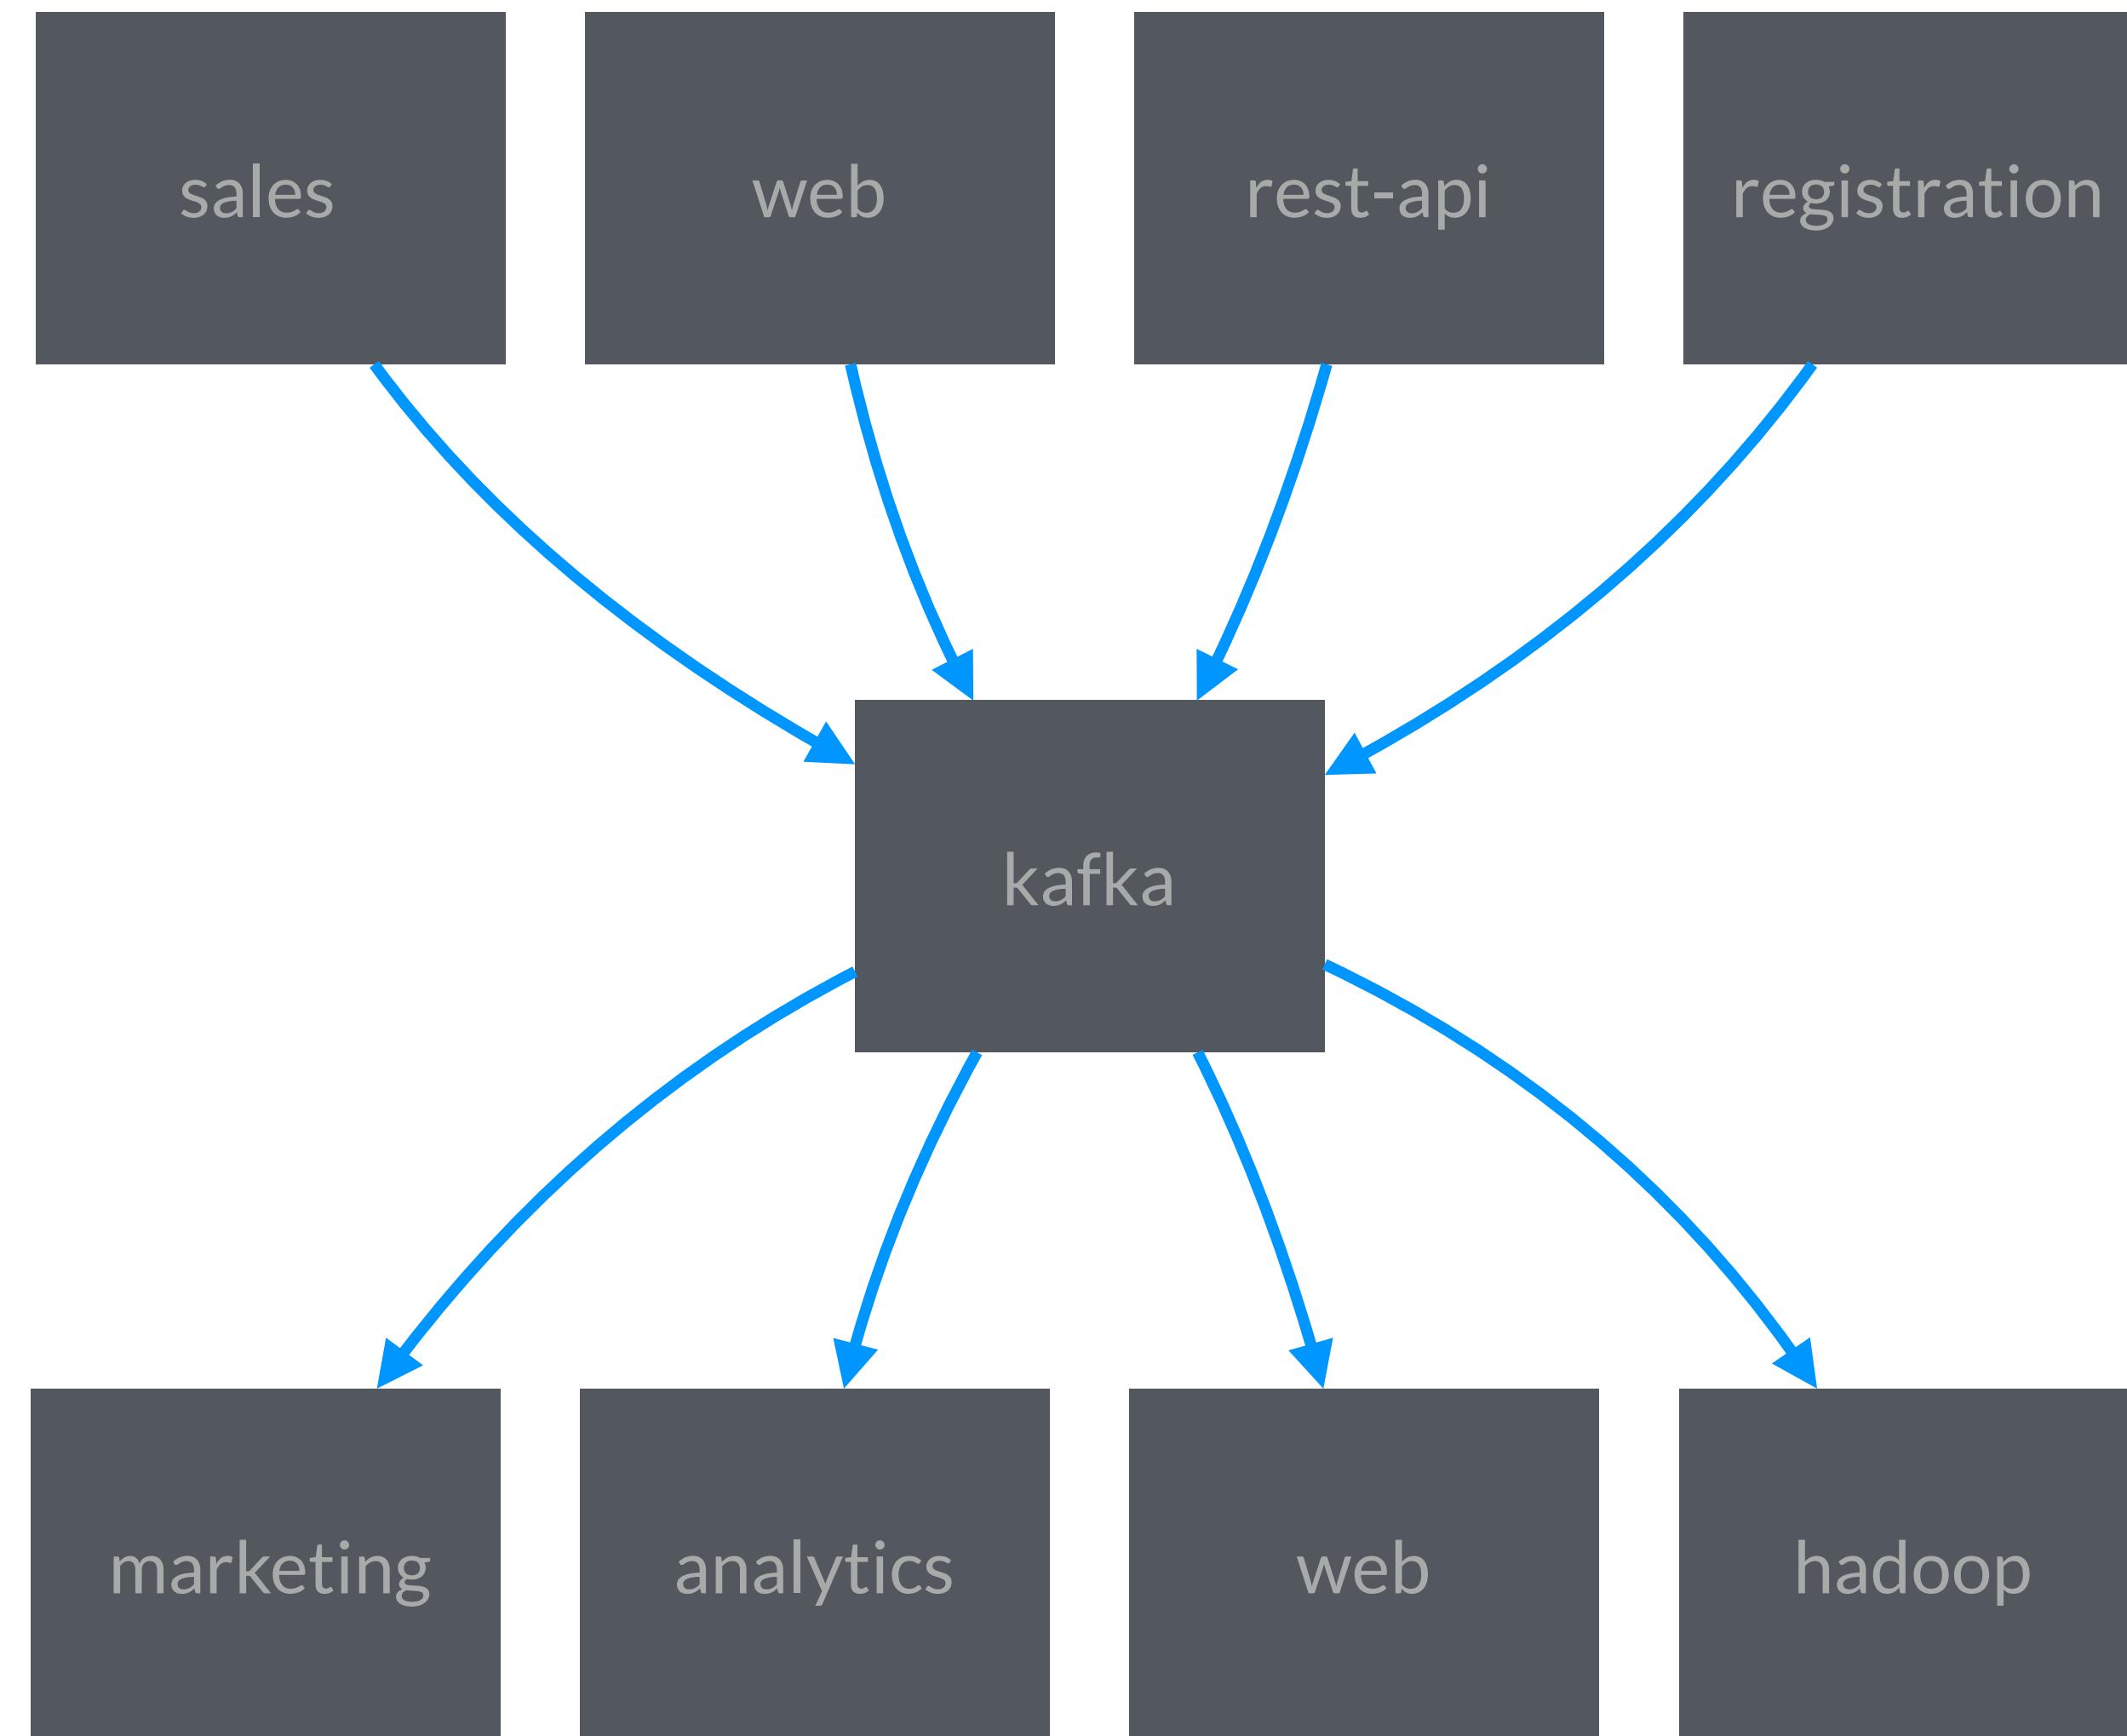
Kafka Introduction









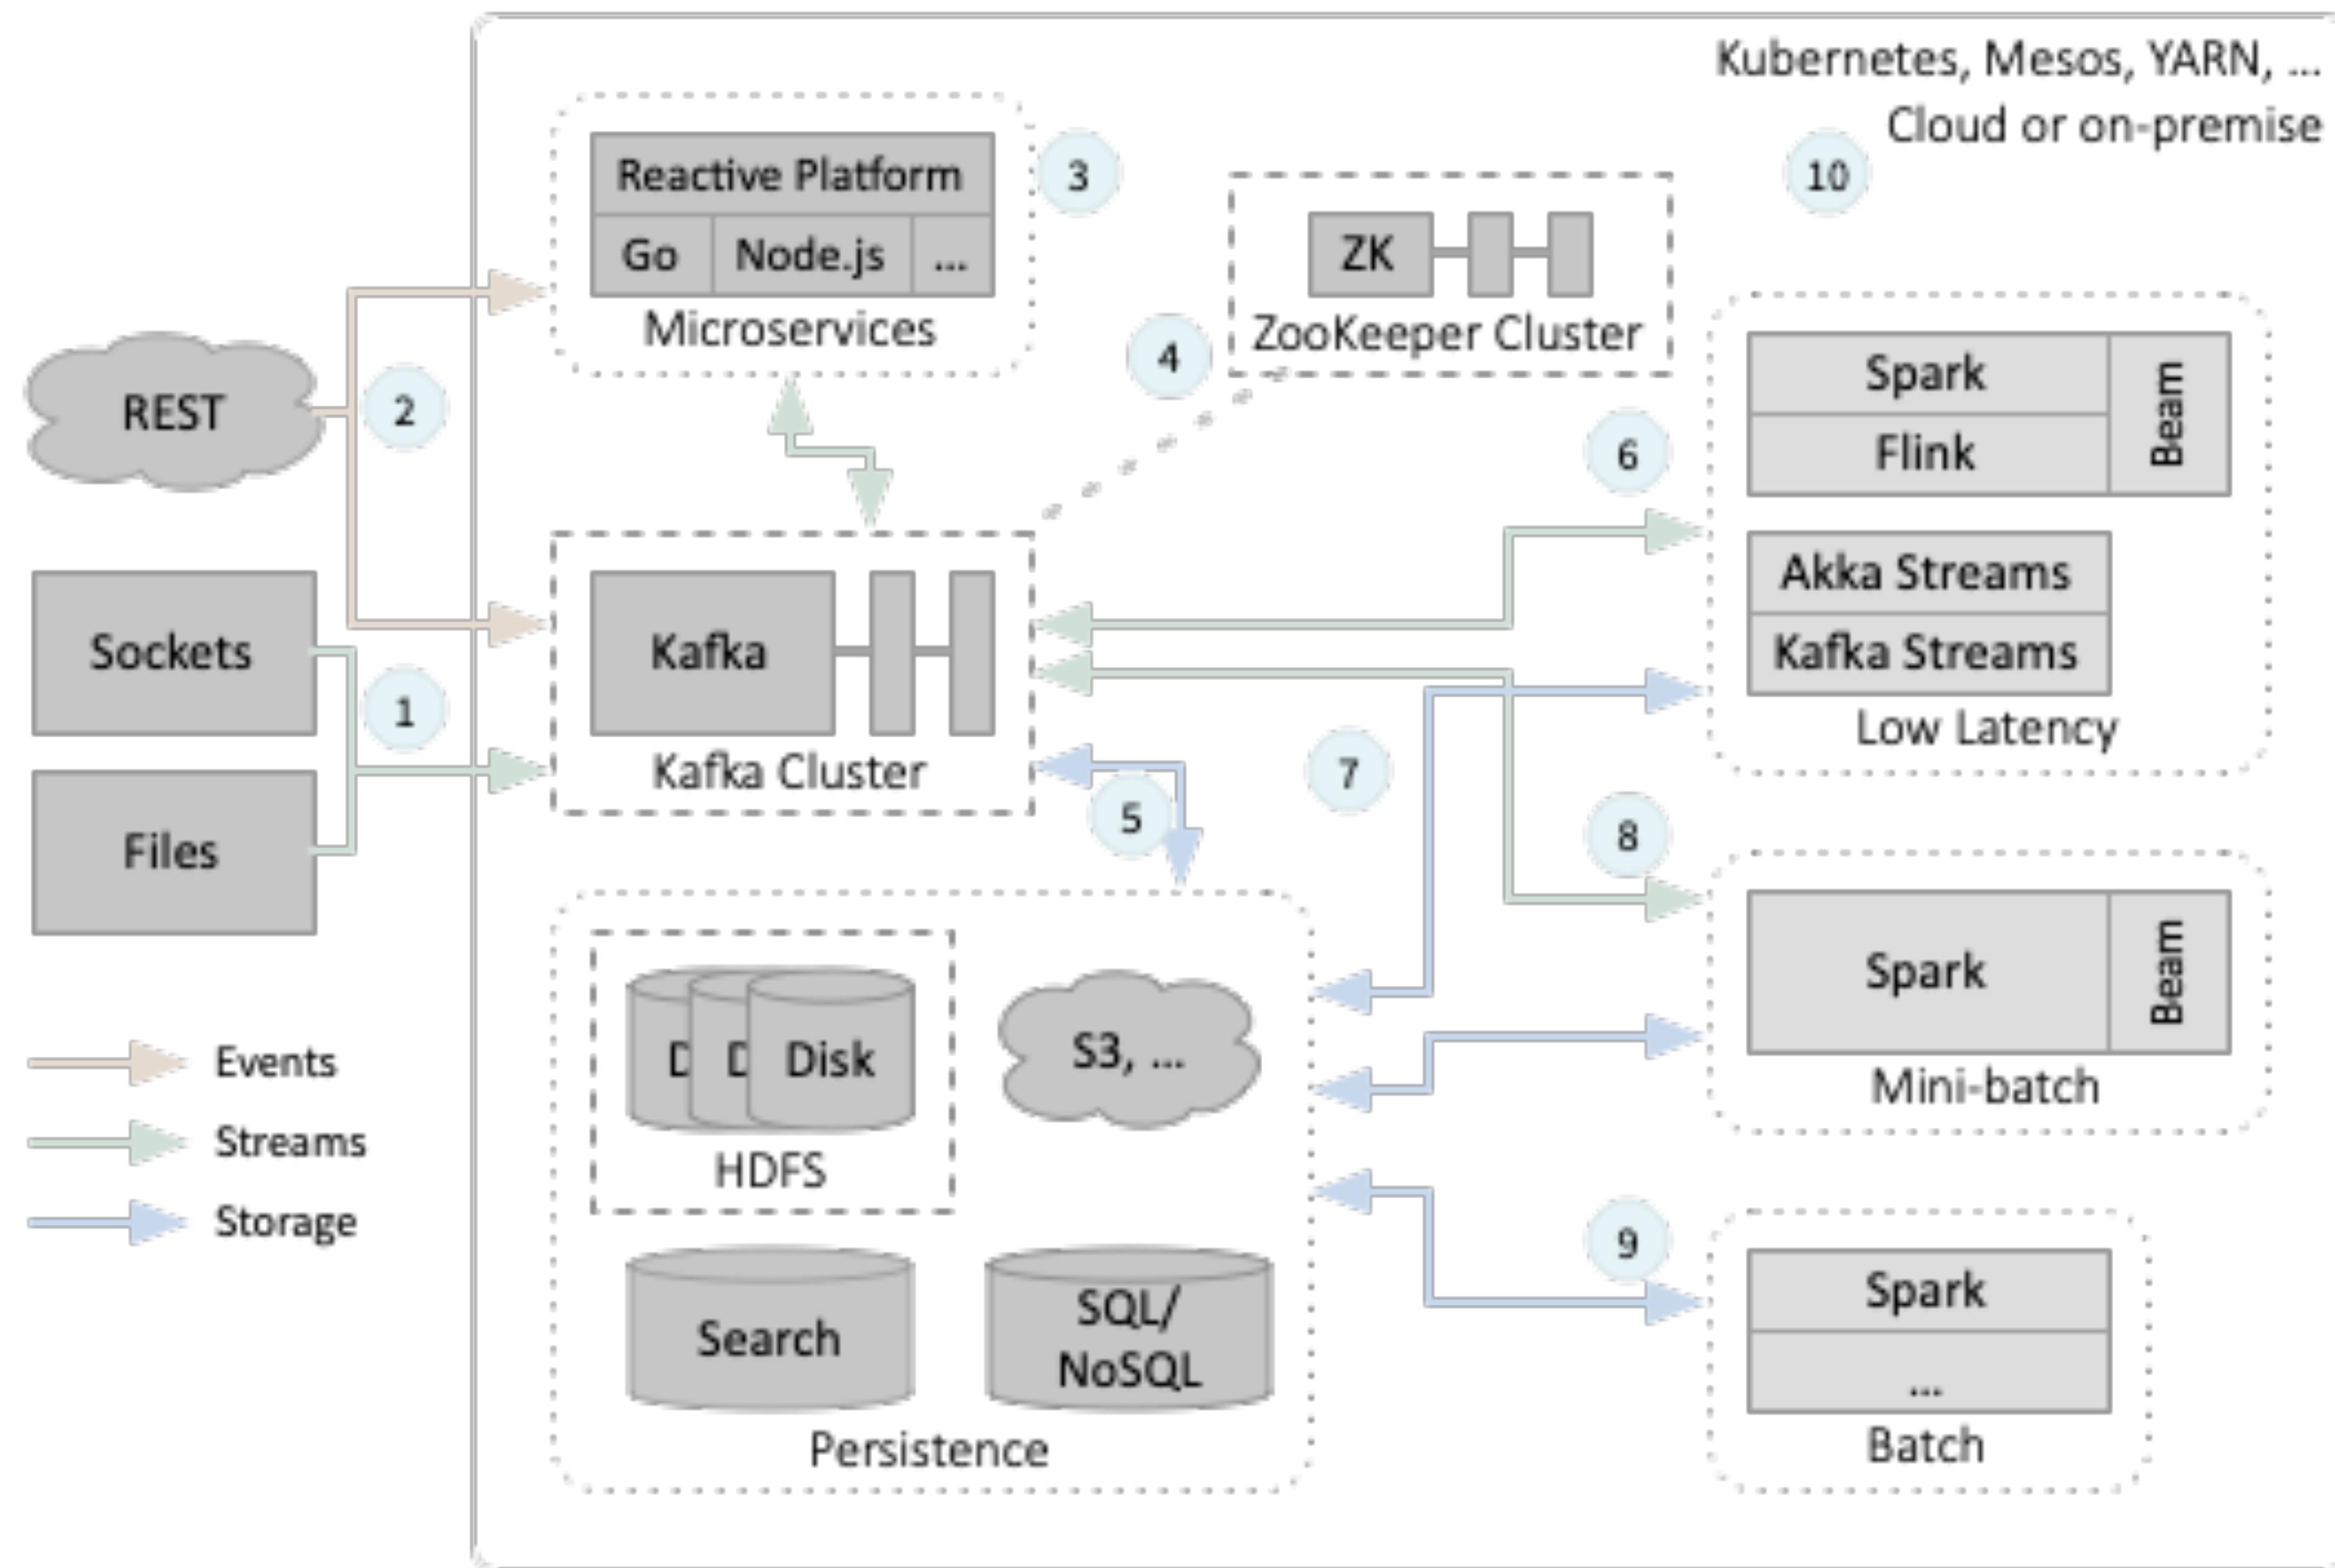


Note: A Producer can be a Consumer

About Kafka

- Handles millions of messages per second, high throughput, high volume
- Distributed and Replicated Commit-log
- Real Time Data Processing
- Stream processing


Dean Wampler Architecture



Kafka Messages

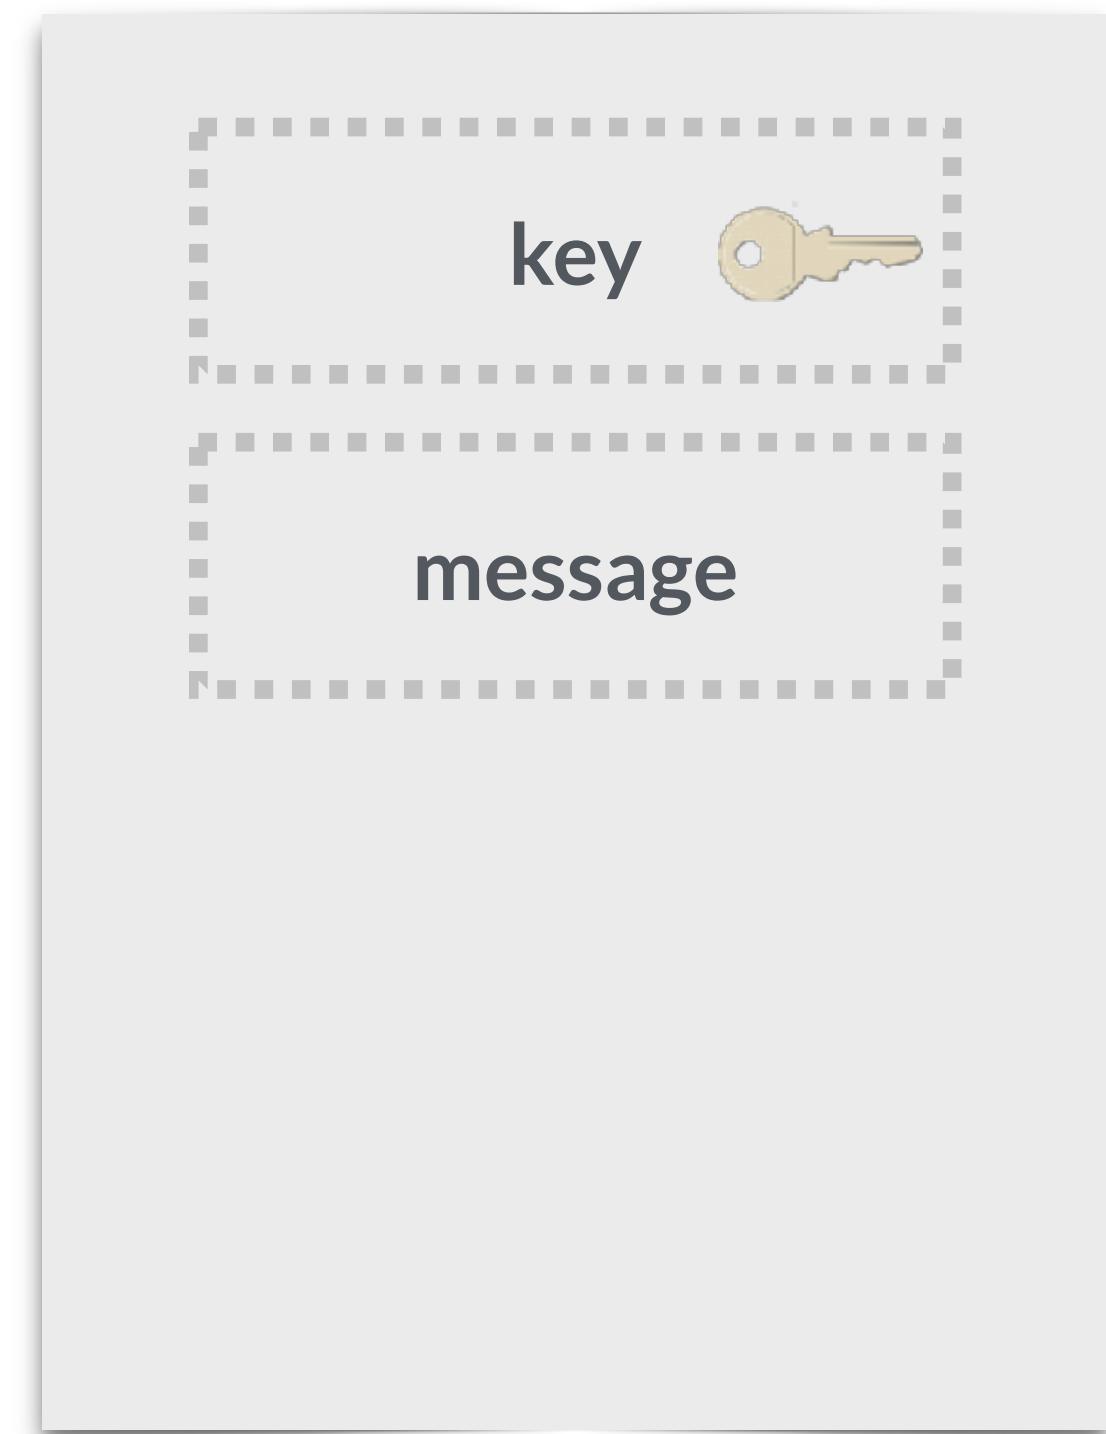
A Kafka Message

- Similar to a *row* or a *record*
- Message is an array of bytes
- No special serialization, that is done at the producer or consumer



A Kafka Message Key


- Message may contain a *key* for better distribution to partitions
- The *key* is also an array of bytes
- If a *key* is provided, a partitioner will hash the key and map it to a single partition
- Therefore it is the only time that something is guaranteed to be in order



Kafka Producers


kafka broker: 0


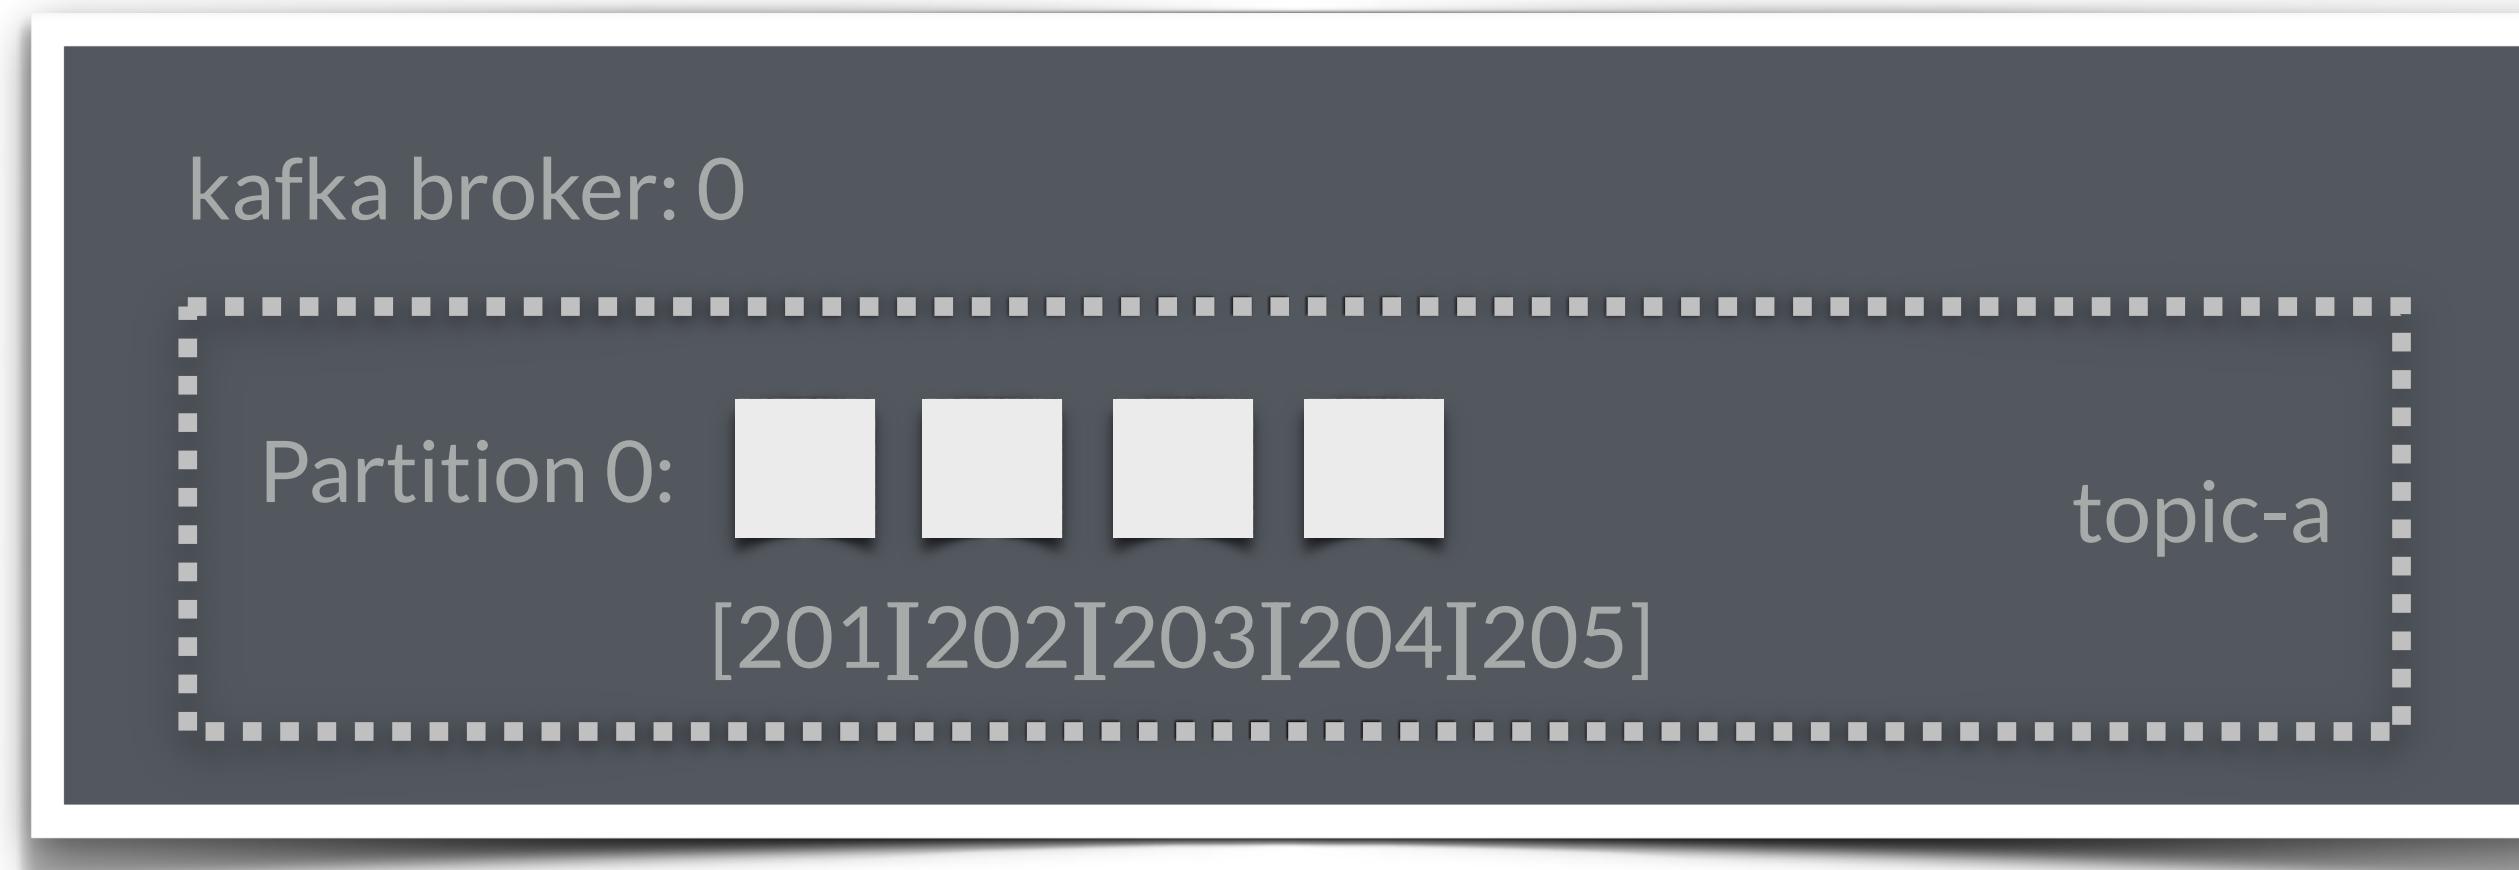
Partition 0:



[0] [1] [2] [3] [4]

topic-a







Retention: The data is temporary

Kafka Batch

- A collection of messages, that is sent, configured in *bytes*
- Sent to the same *topic* and the same *partition*
- *Avoids overhead of sending multiple message over the wire*





Every batch knows where
each partition is going

$\text{murmur2(bytes)} \% \text{ number partitions}$

A-E




F-K




L-S




T-Z

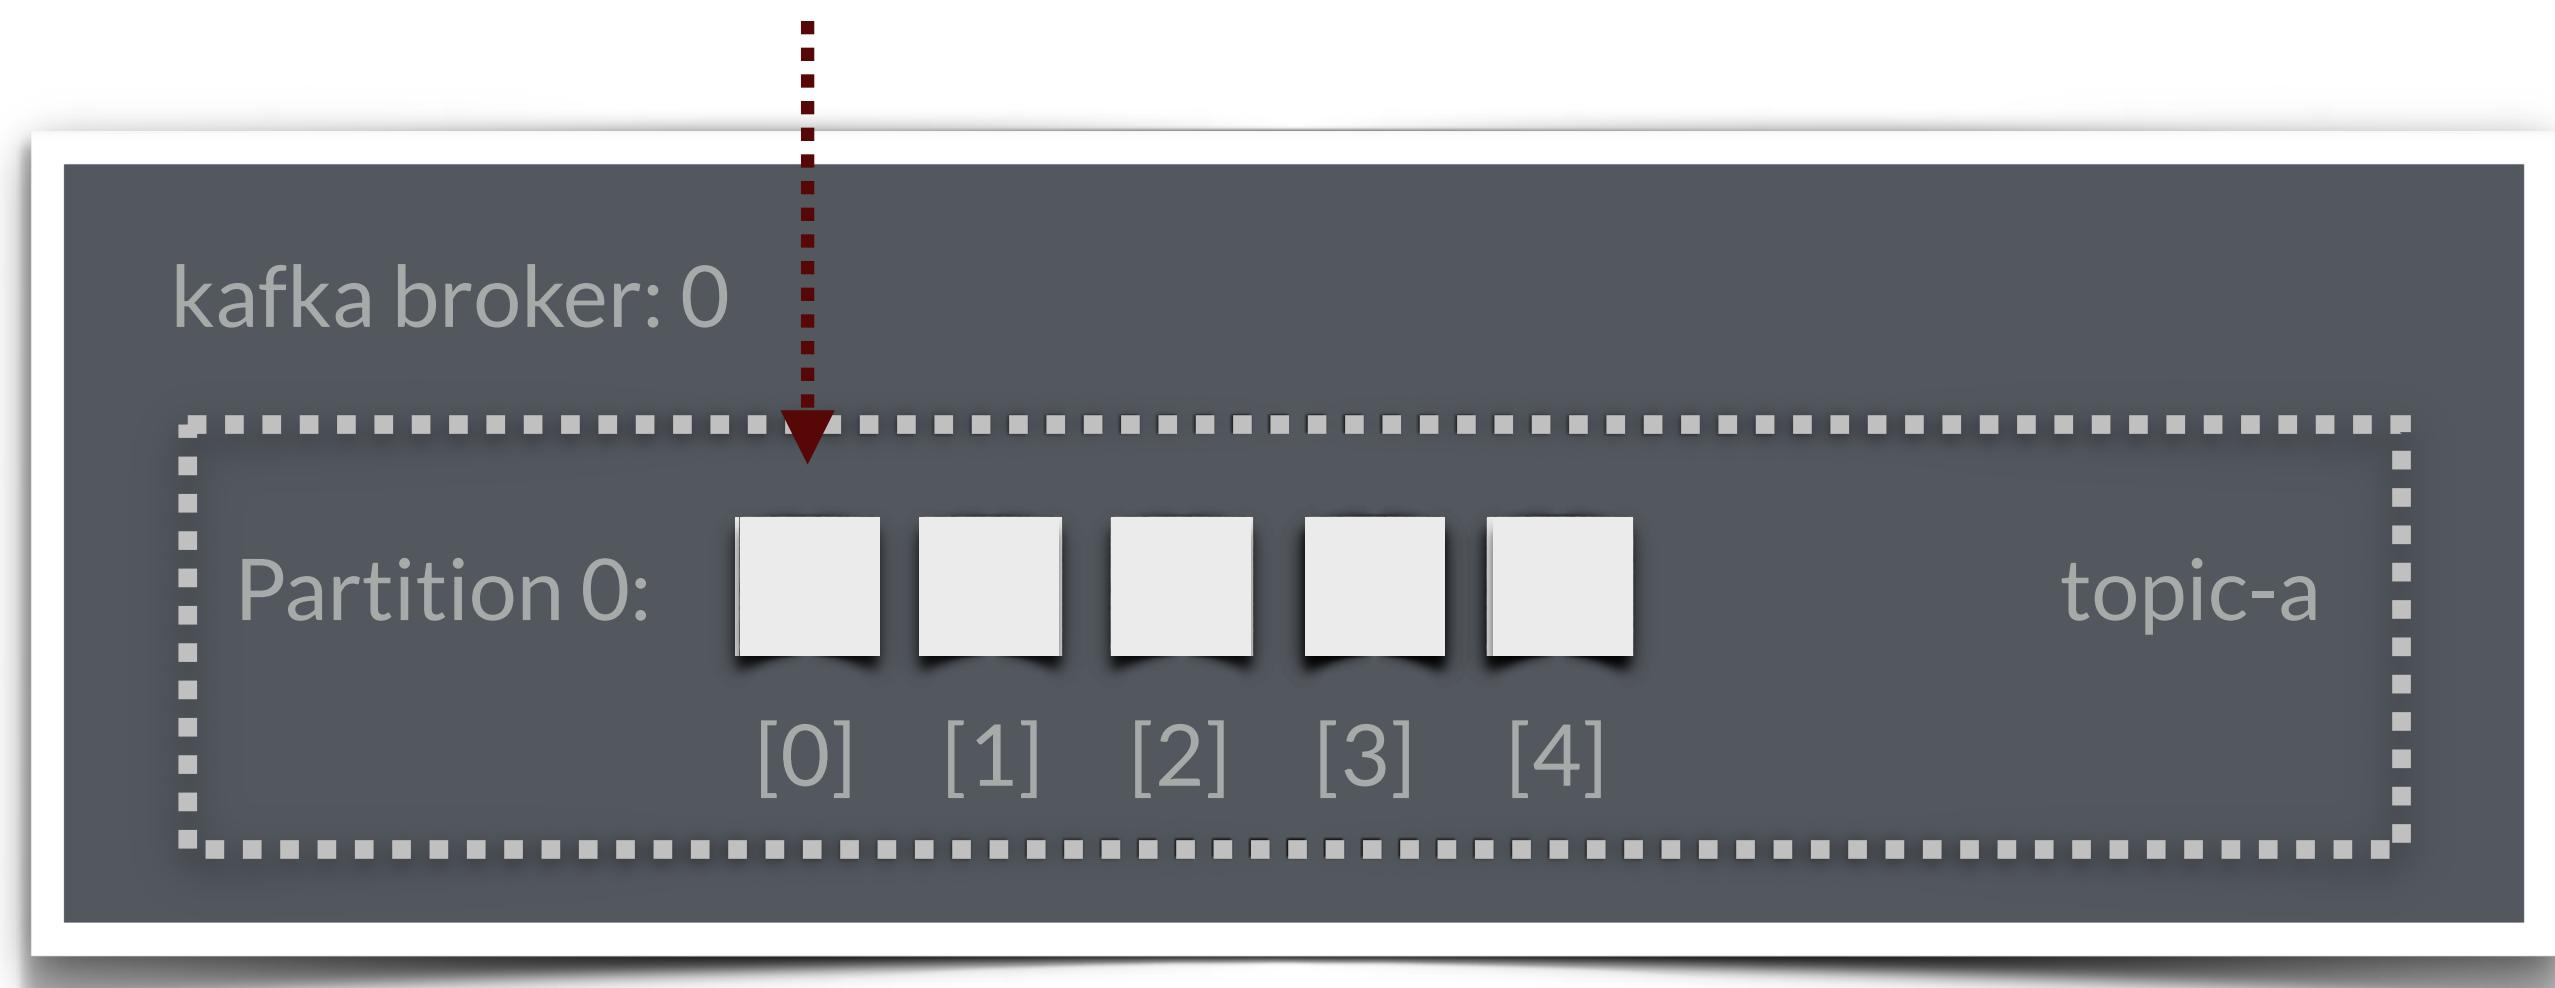



Kafka Consumers

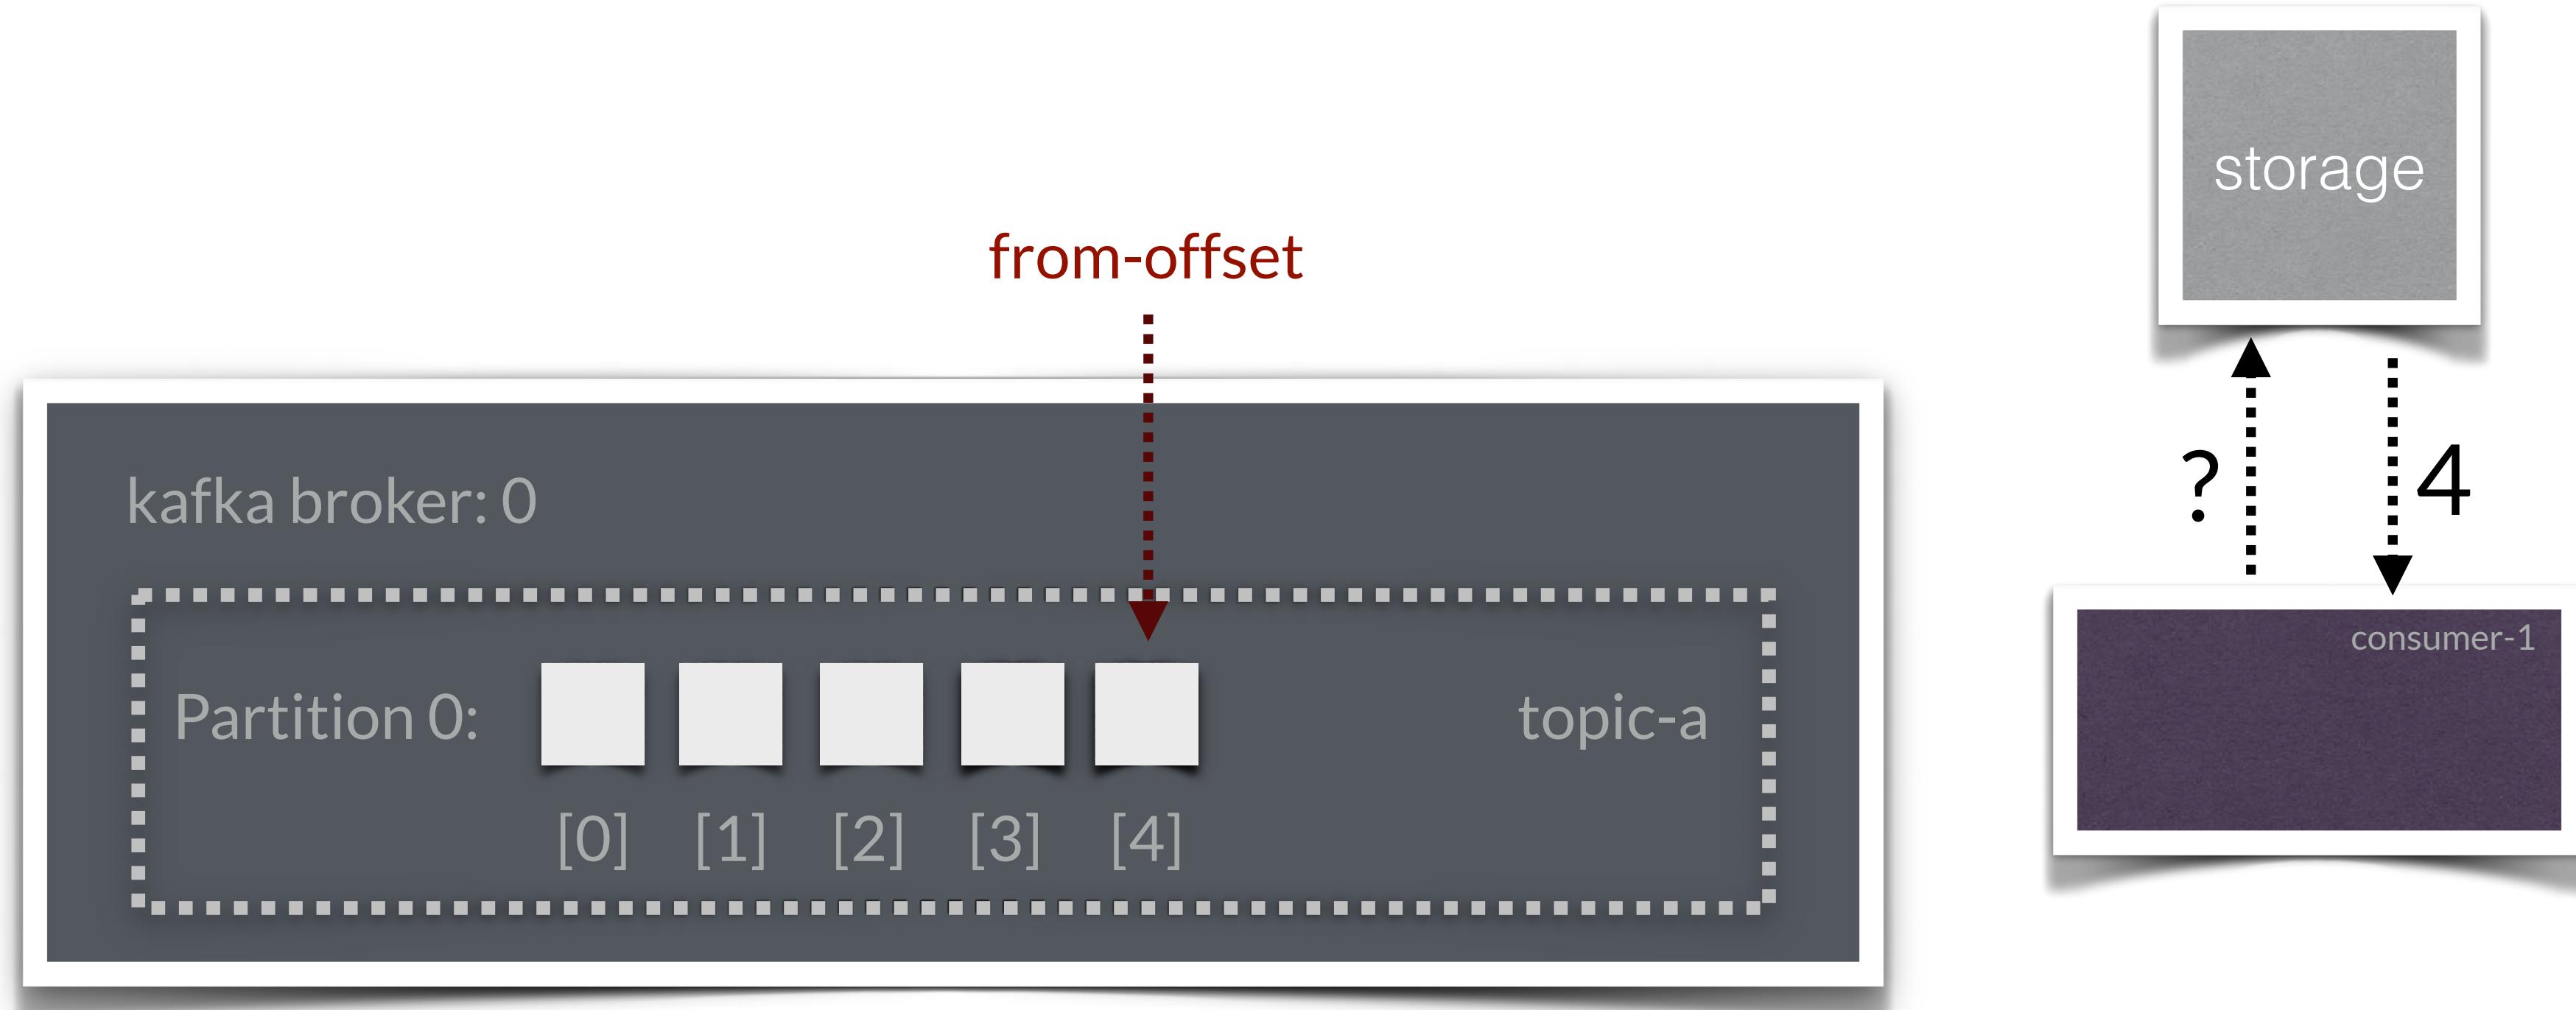




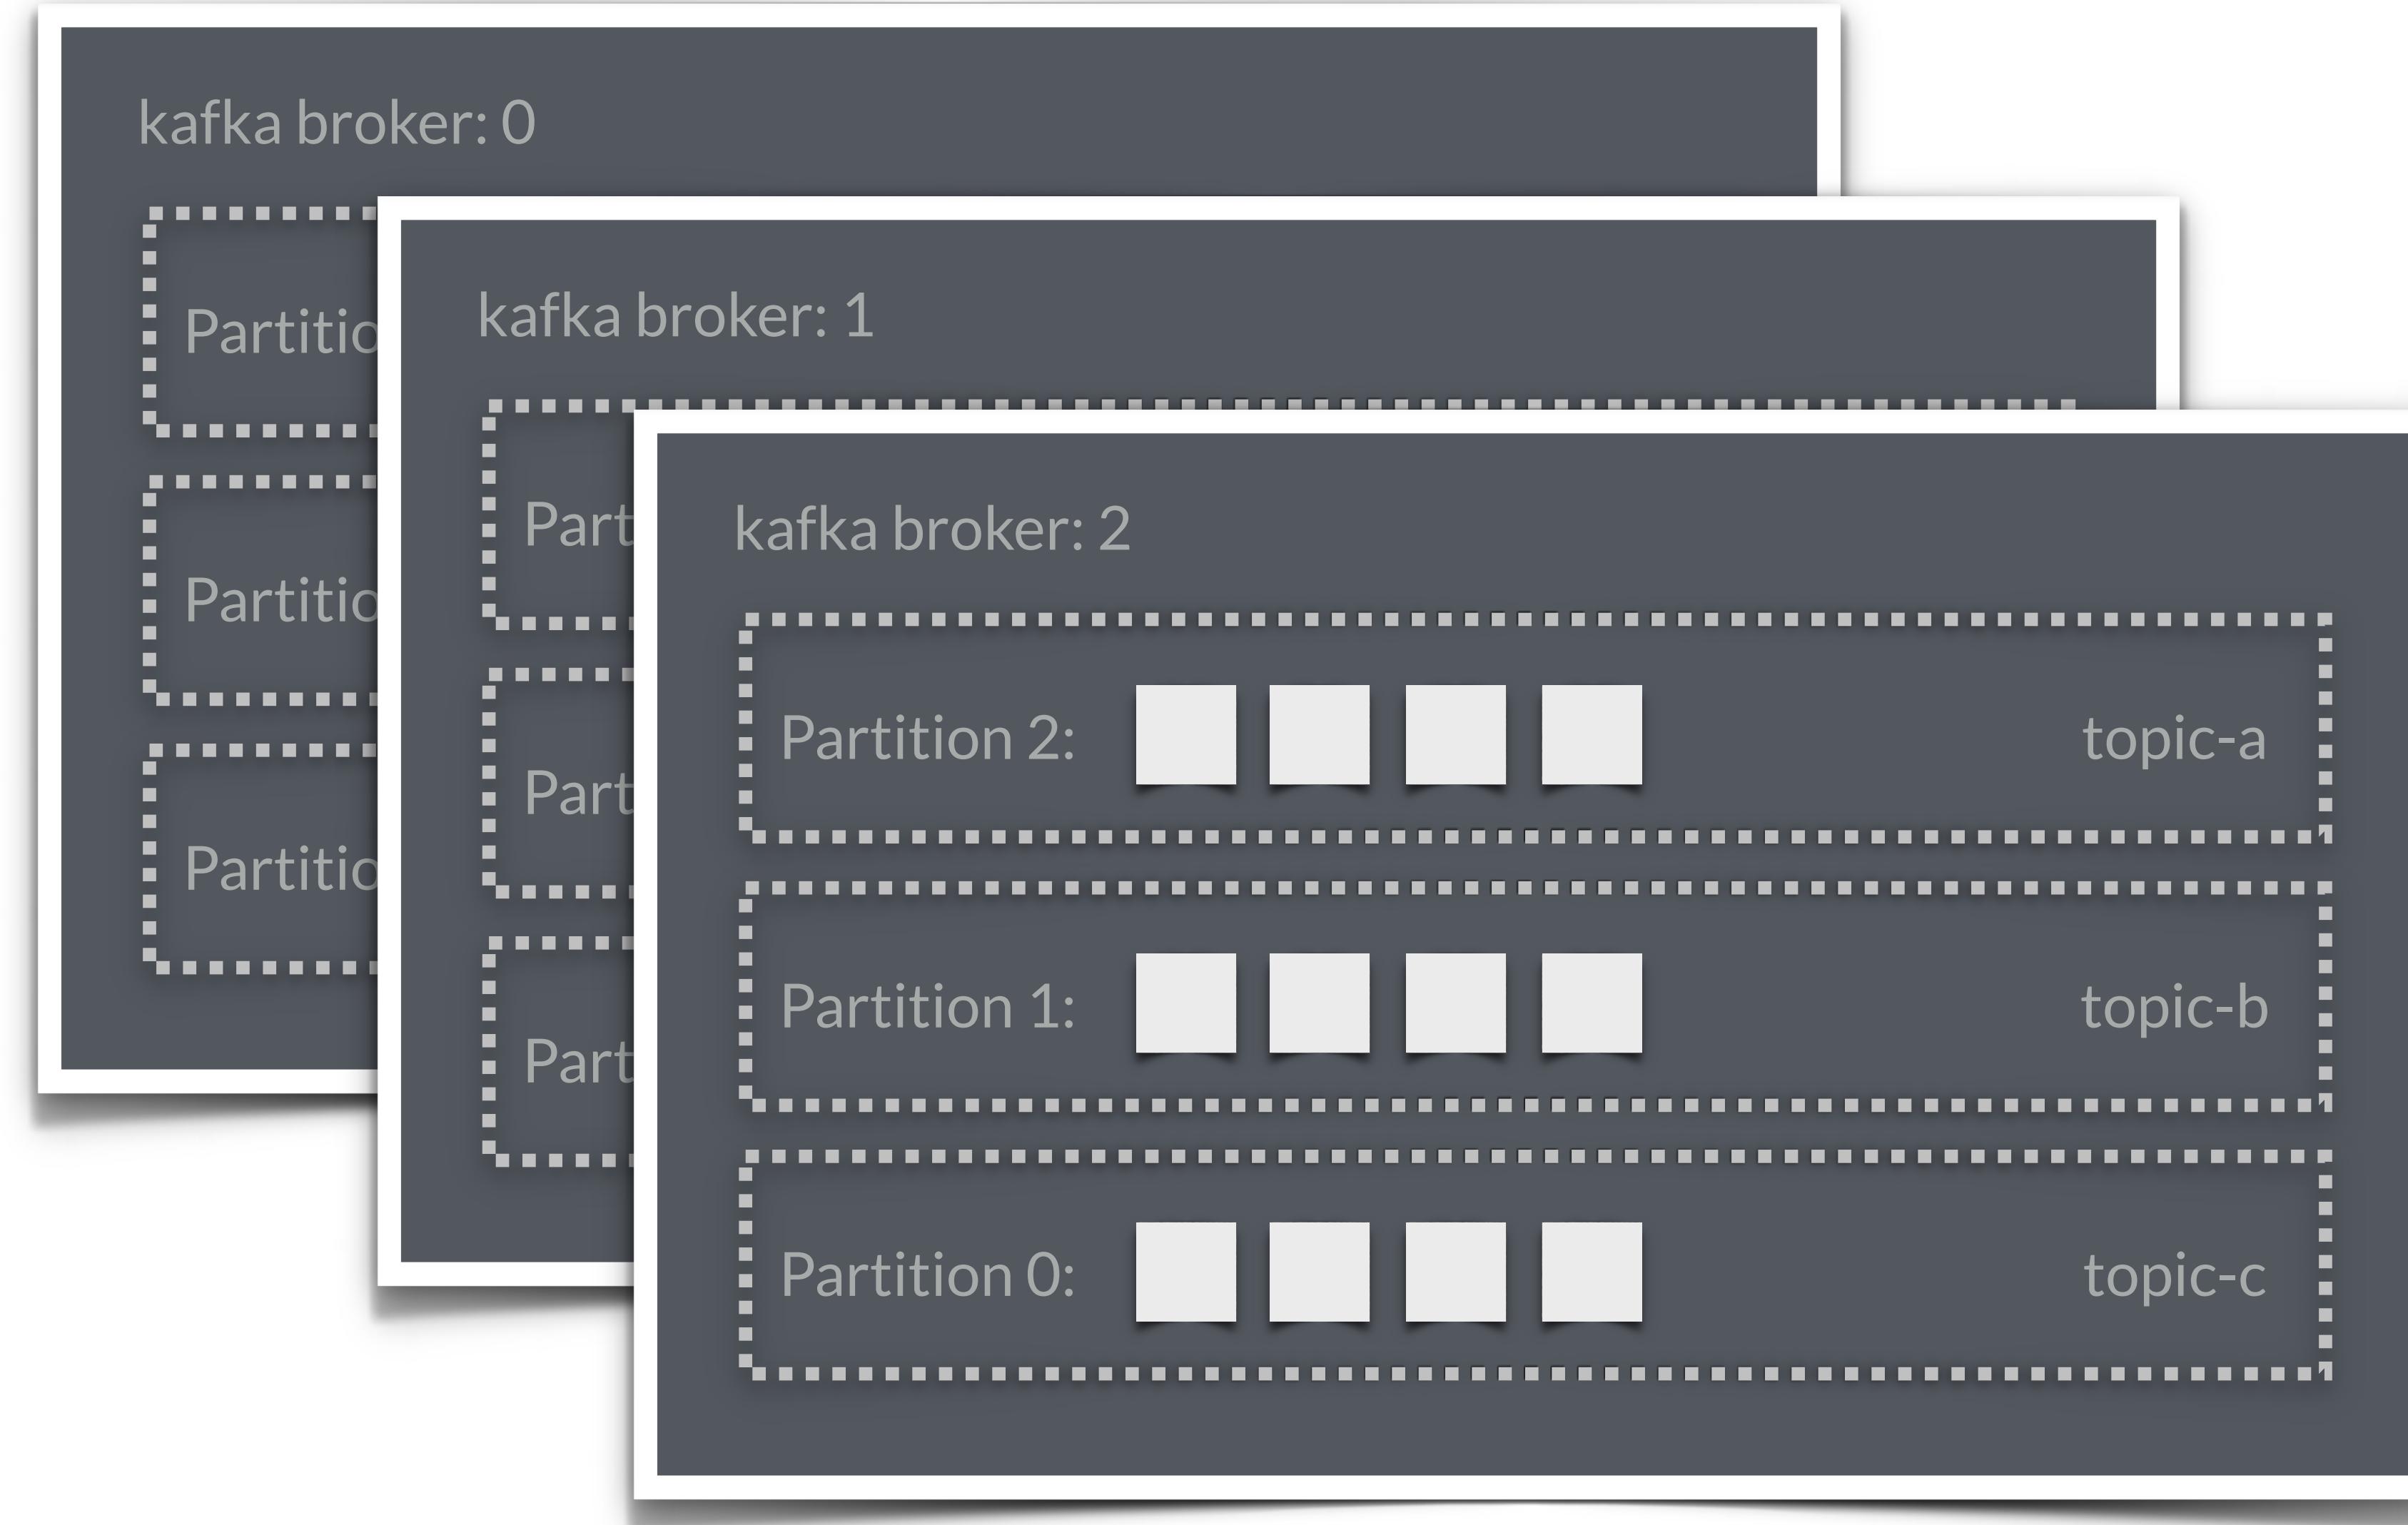
from-beginning







Kafka Partitions



Each partition is on a different broker,
therefore a single topic is scaled



A-E



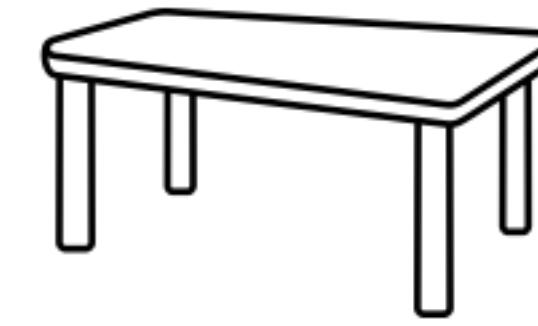
F-K



L-S



T-Z



Kafka Pub/Sub Comparisons

Comparison Pub/Sub

- Kafka as the distributed log
- Ability to be replayed in a consistent manner
- Kafka is also stored durable, in order, and deterministic, if key is available
- Data can also be distributed for resiliency
- Scalable and Elastic

Zookeeper

Zookeeper Essence

- Centralized Coordination Service
- Maintains:
 - Metadata
 - Naming
 - Configuration

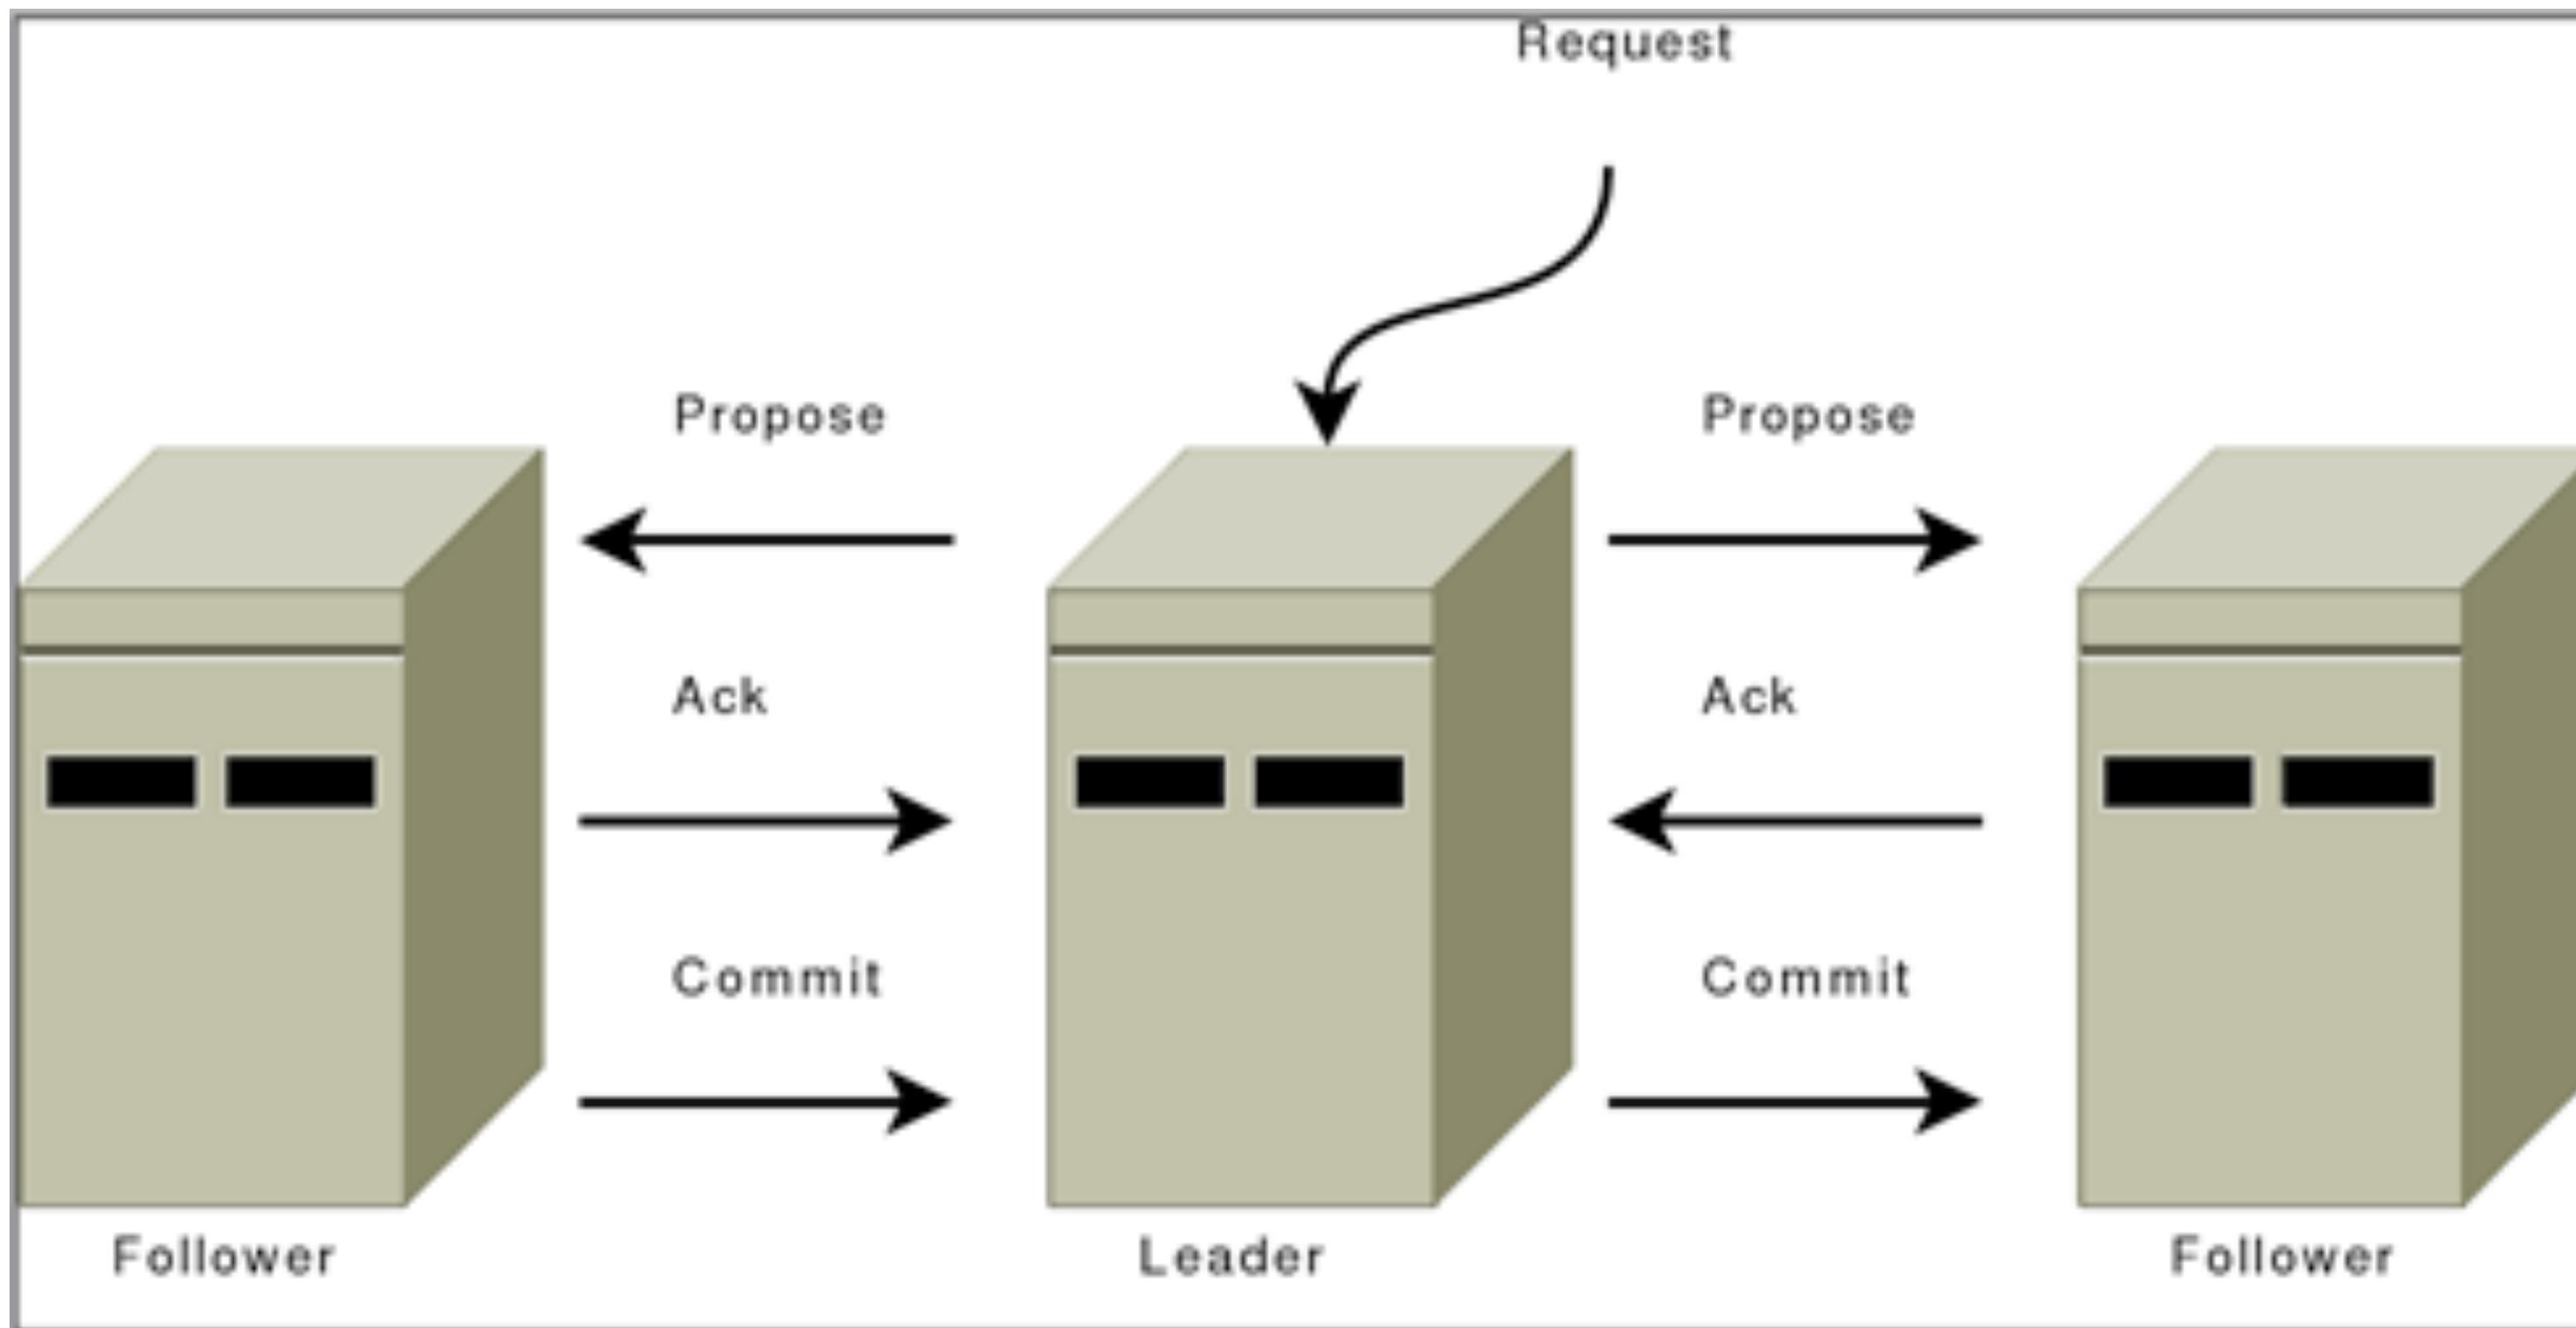


What is the purpose of Zookeeper?

- Zookeeper up to this point has been coordination service of Kafka
- Zookeeper is network coordinator in a distributed system
- Used extensively for discovery, access control lists, resource allocation, leader elements, and notifications.
- If you lost the Kafka data in ZooKeeper, the mapping of replicas to Broker and topic configurations would be lost as well, making your Kafka cluster no longer functional and potentially resulting in total data loss.


How does Zookeeper work?

- Zookeeper is a centralized service for configuration, naming, and provide synchronization.
- Doing such things are difficult and can run afoul of bugs and race conditions
- Zookeeper was created to handle this difficult problem
- Zookeeper uses the ZooKeeper Atomic Broadcast Protocol to communicate settings to others




https://zookeeper.apache.org/doc/r3.1.2/zookeeperInternals.html#sc_atomicBroadcast

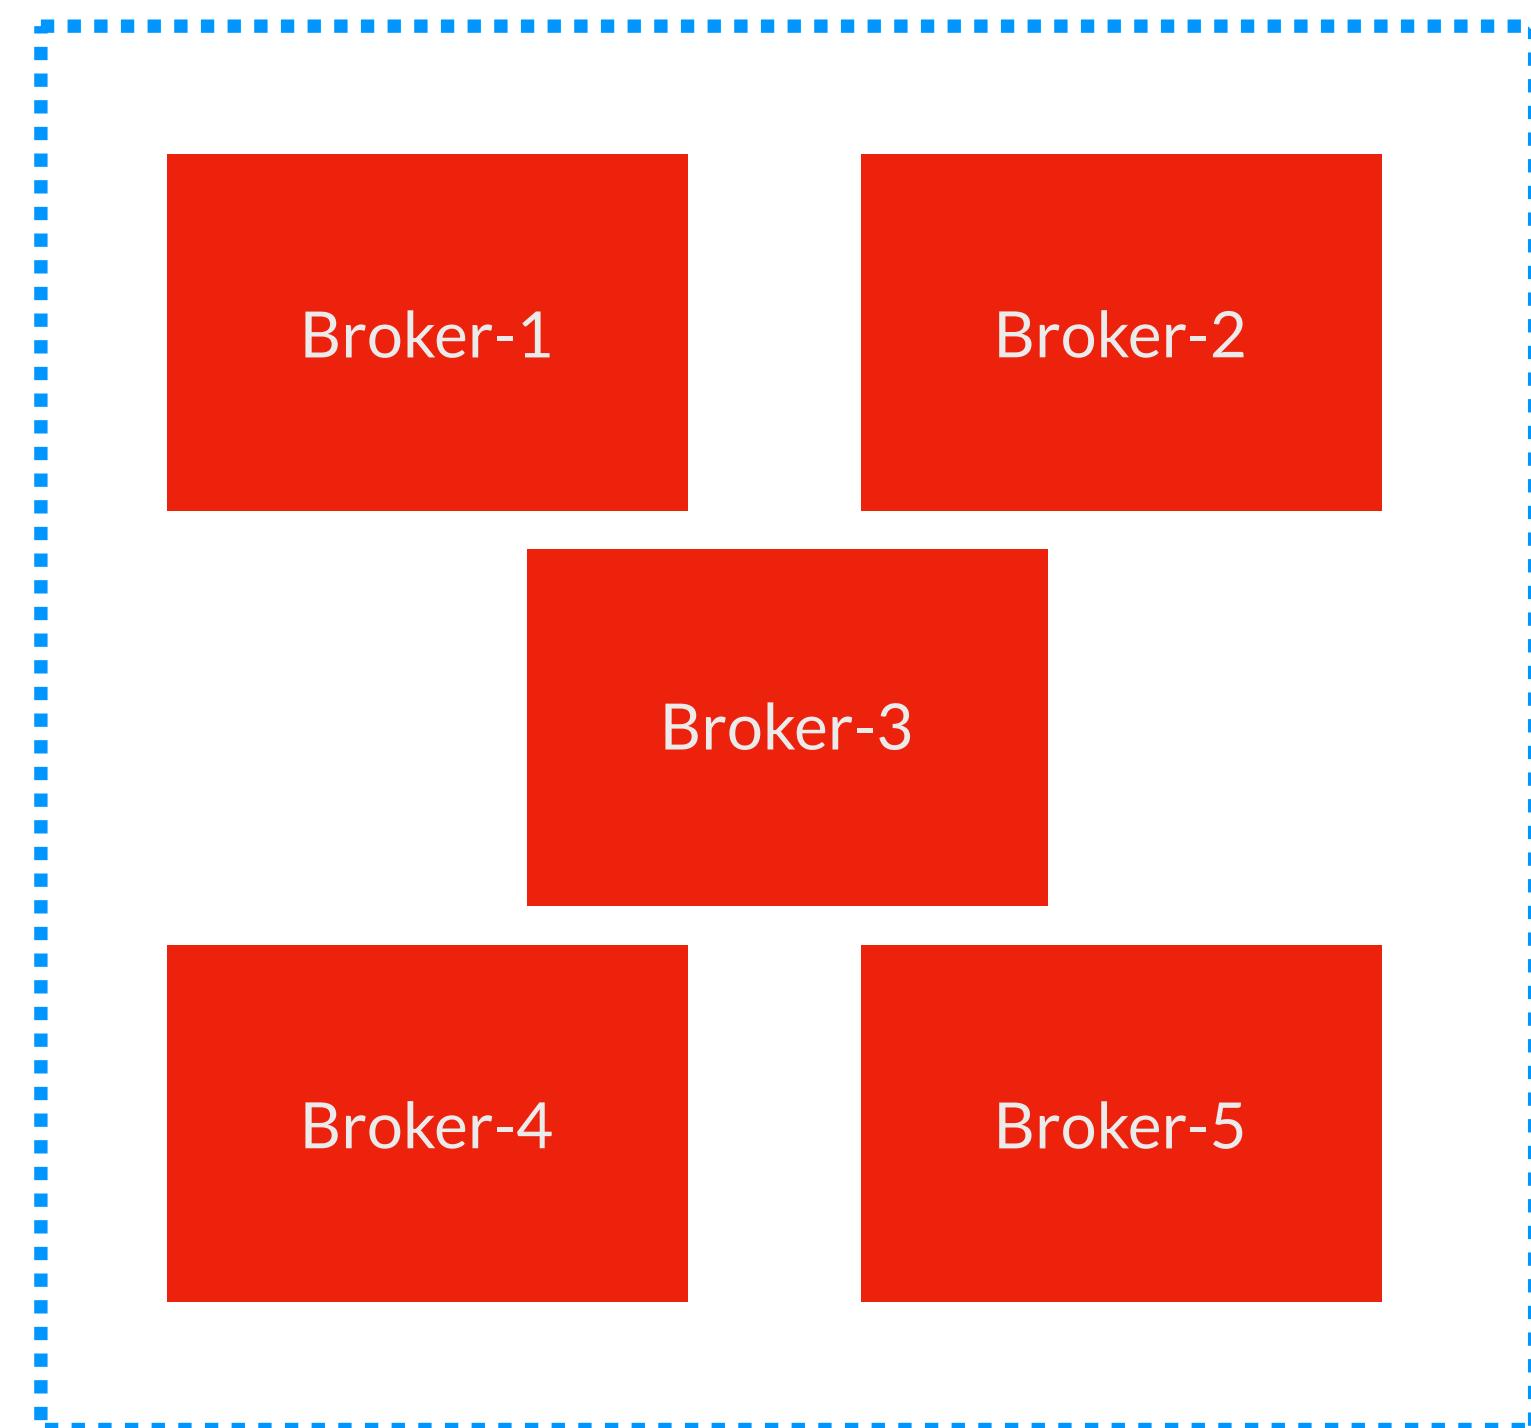
Kafka Cluster




Zookeeper Ensemble



Kafka Cluster



Zookeeper Ensemble



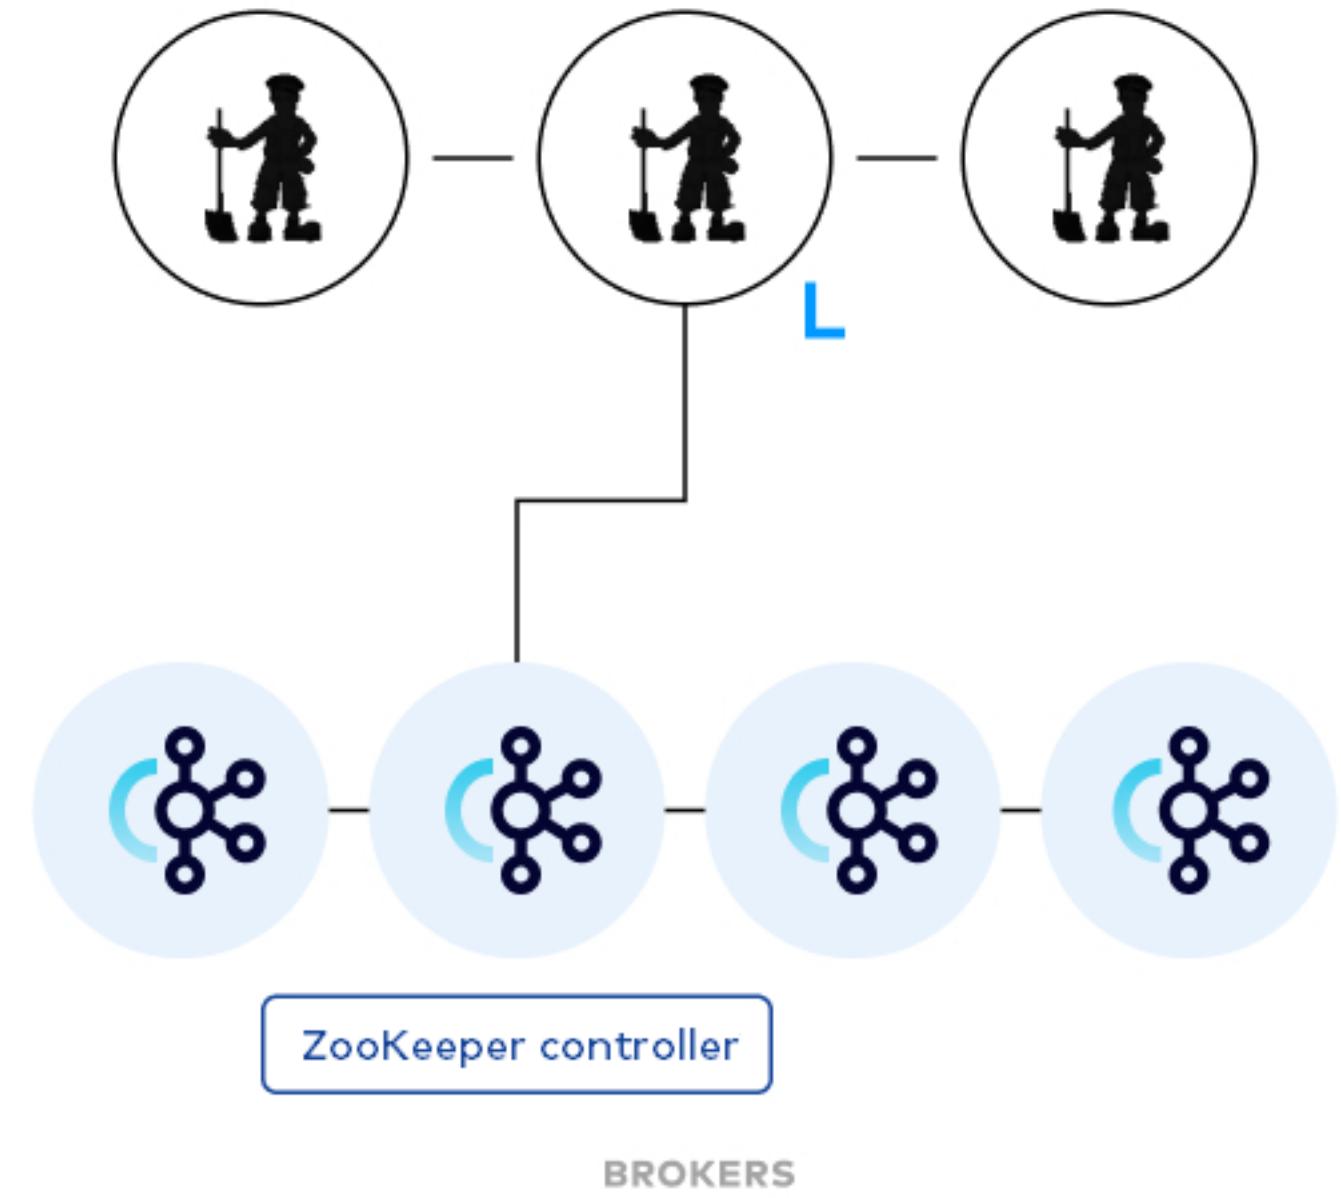
KRaft



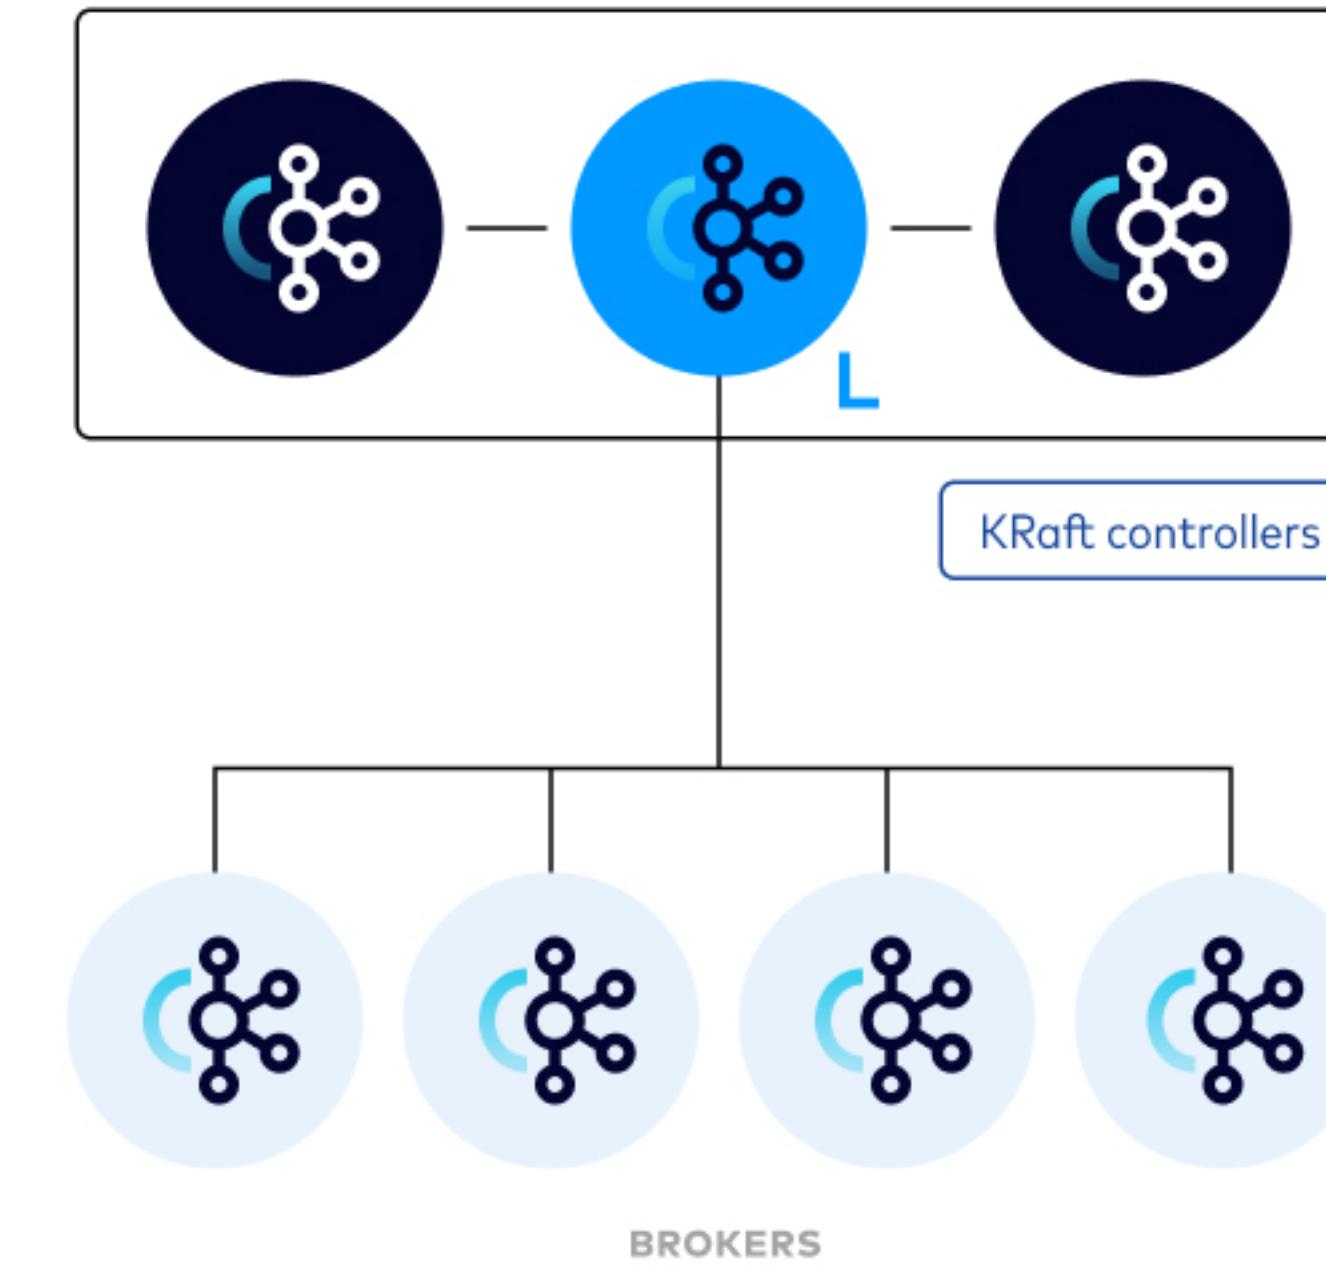
KRaft Protocol

- Introduced in [KIP-500](#), and pronounced as “Craft”
- Apache Kafka® Raft (KRaft) is the consensus protocol that was introduced to remove Kafka’s dependency on ZooKeeper for metadata management.
- Based on the Raft Protocol, but running under a Kafka architecture
- Consolidates responsibility for metadata into Kafka itself, rather than requiring and configuring two different systems: ZooKeeper and Kafka.

ZooKeeper



KRaft



L denotes cluster metadata leader

Note: This shows KRaft running in Isolated Mode

What was the motivation?

- Moving to KRaft for metadata storage simplifies Kafka.
- ZooKeeper is a separate system, with its own configuration file syntax, management tools, and deployment patterns, which makes deploying Kafka more complicated for system administrators.
- KRaft enables right-sized clusters, meaning clusters that are sized with the appropriate number of brokers and compute that satisfies a use-case's throughput and latency requirements, to scale up to millions of partitions.
- Metadata failover is near-instantaneous with KRaft.
- KRaft provides a single system to manage since KRaft is part of Kafka itself. Single management model for authentication, configuration, protocols, and networking.

Controller Quorum

- The KRaft controller nodes comprise a *Raft quorum* which manages the Kafka metadata log.
- This log contains information about each change to the cluster metadata.
- Everything that is/was currently stored in ZooKeeper, such as *topics*, *partitions*, *ISRs*, *configurations*, and so on, is stored in this log.

How does Raft work?

Keep in Mind...


- This protocol is to ensure that data is equivalent on all nodes
- In the case of Kafka, think ACL lists, ISRs, Topics, and other state changes
- This is *metadata management*

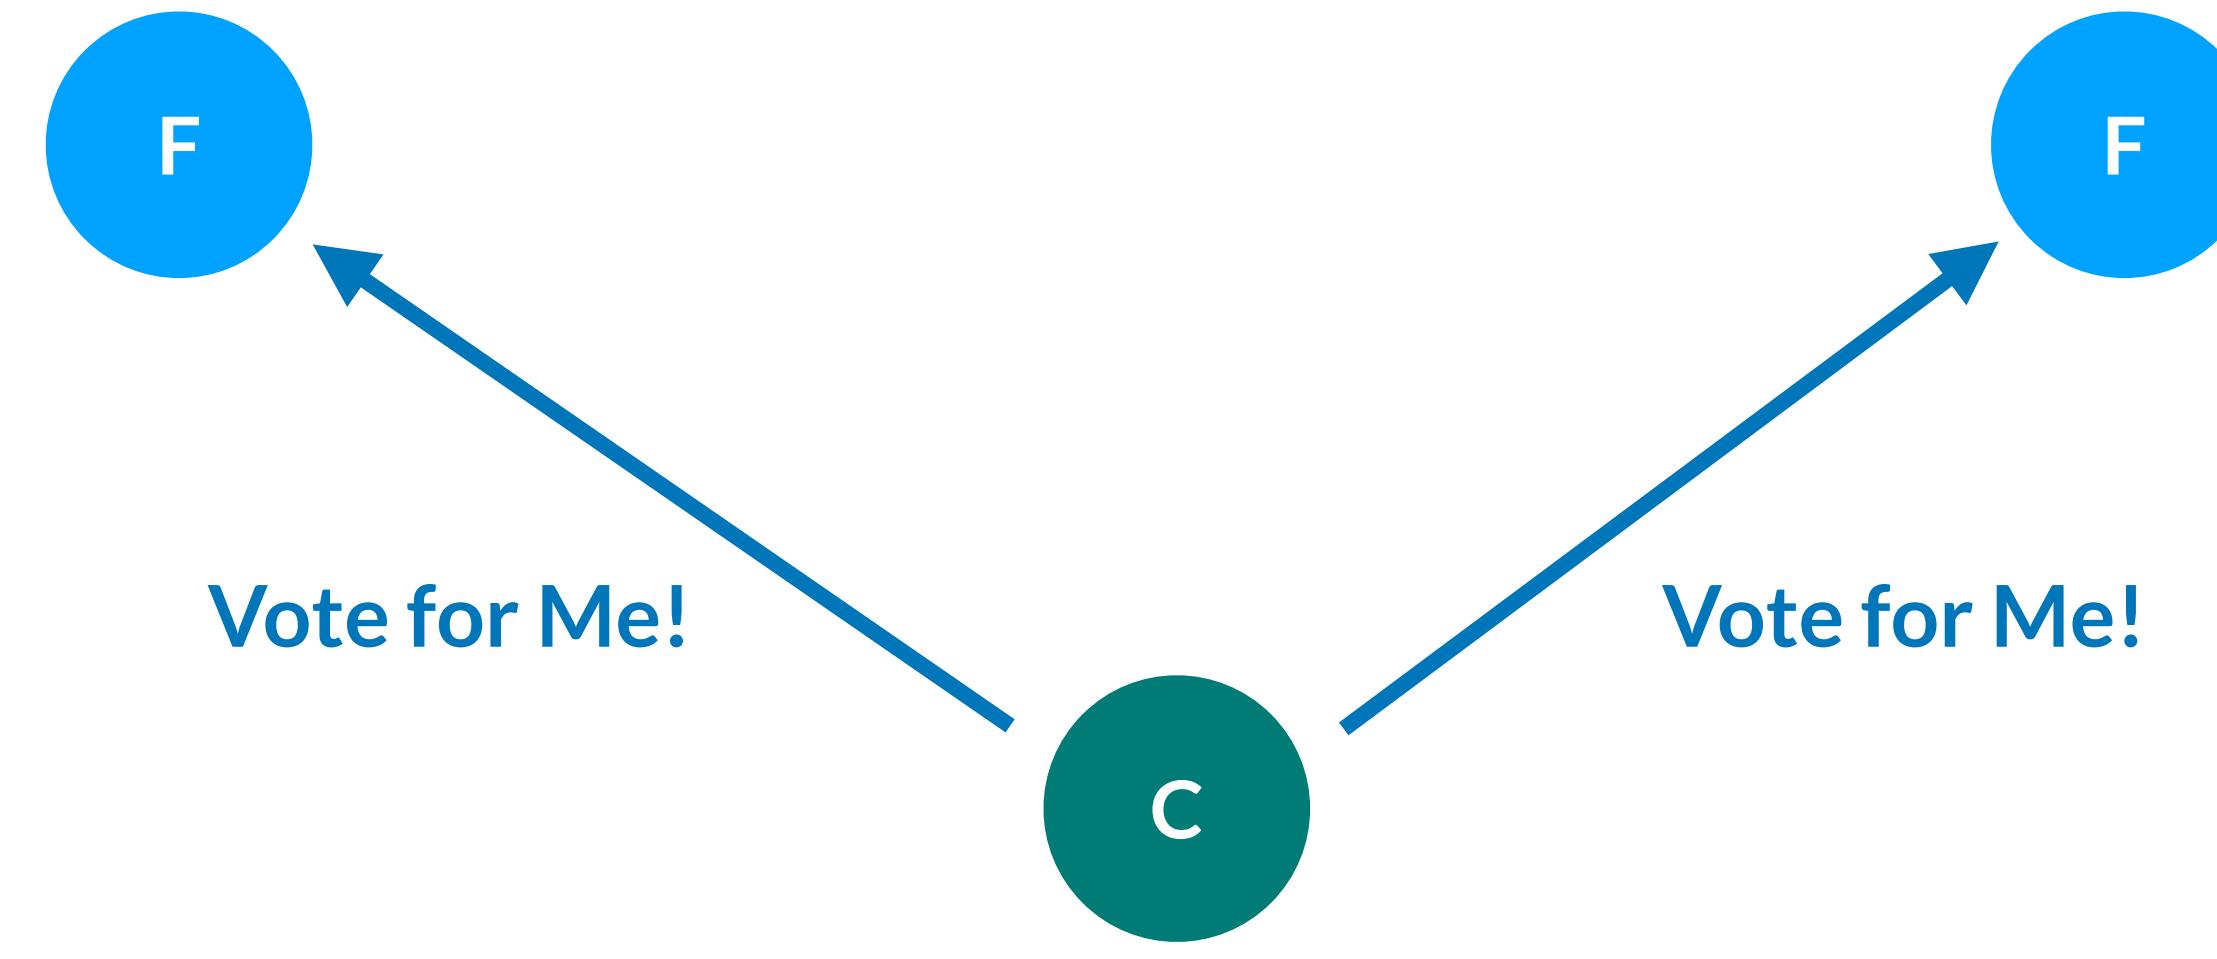
Roles

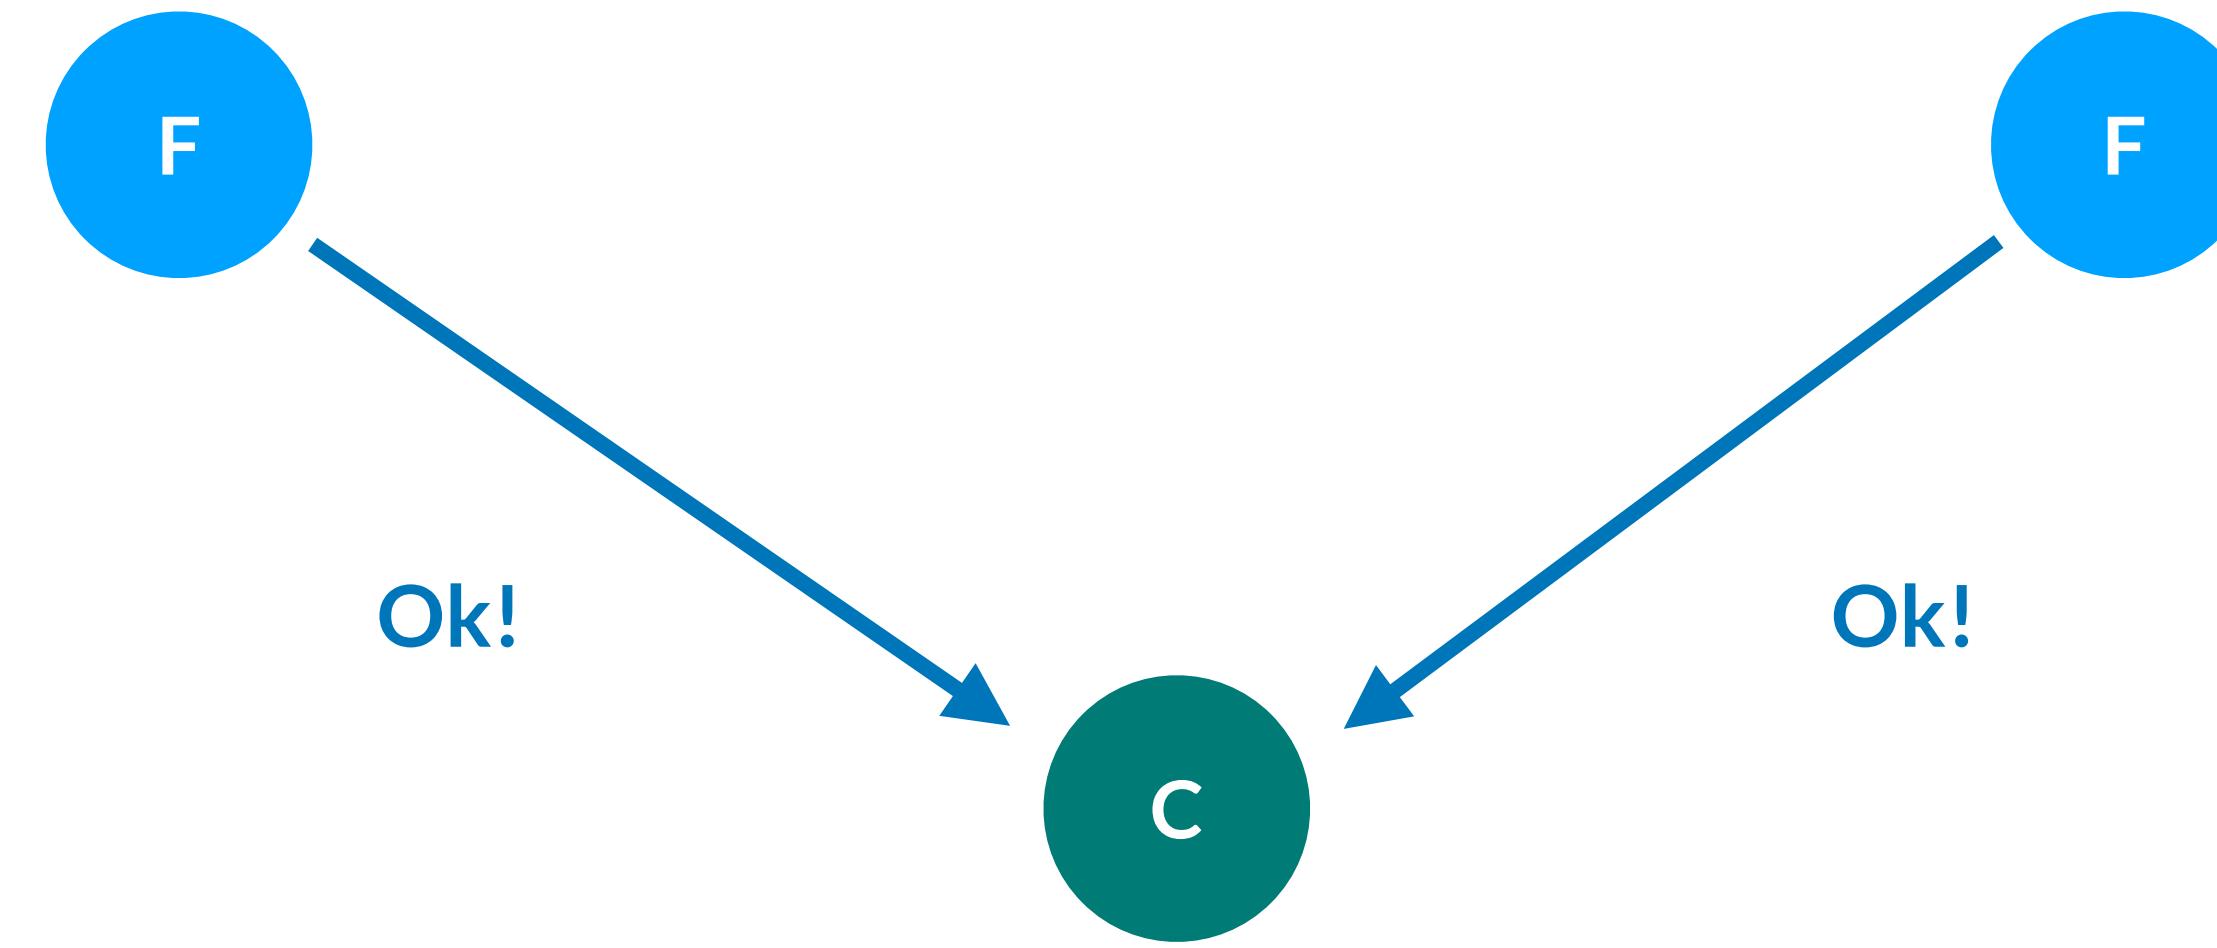
- **Leader** - Manages the replication to the followers
- **Follower** - Takes whatever the leader provides
- **Candidate** - If a follower doesn't hear from the leader, they become candidates

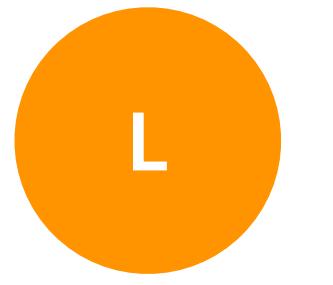
Happy Path

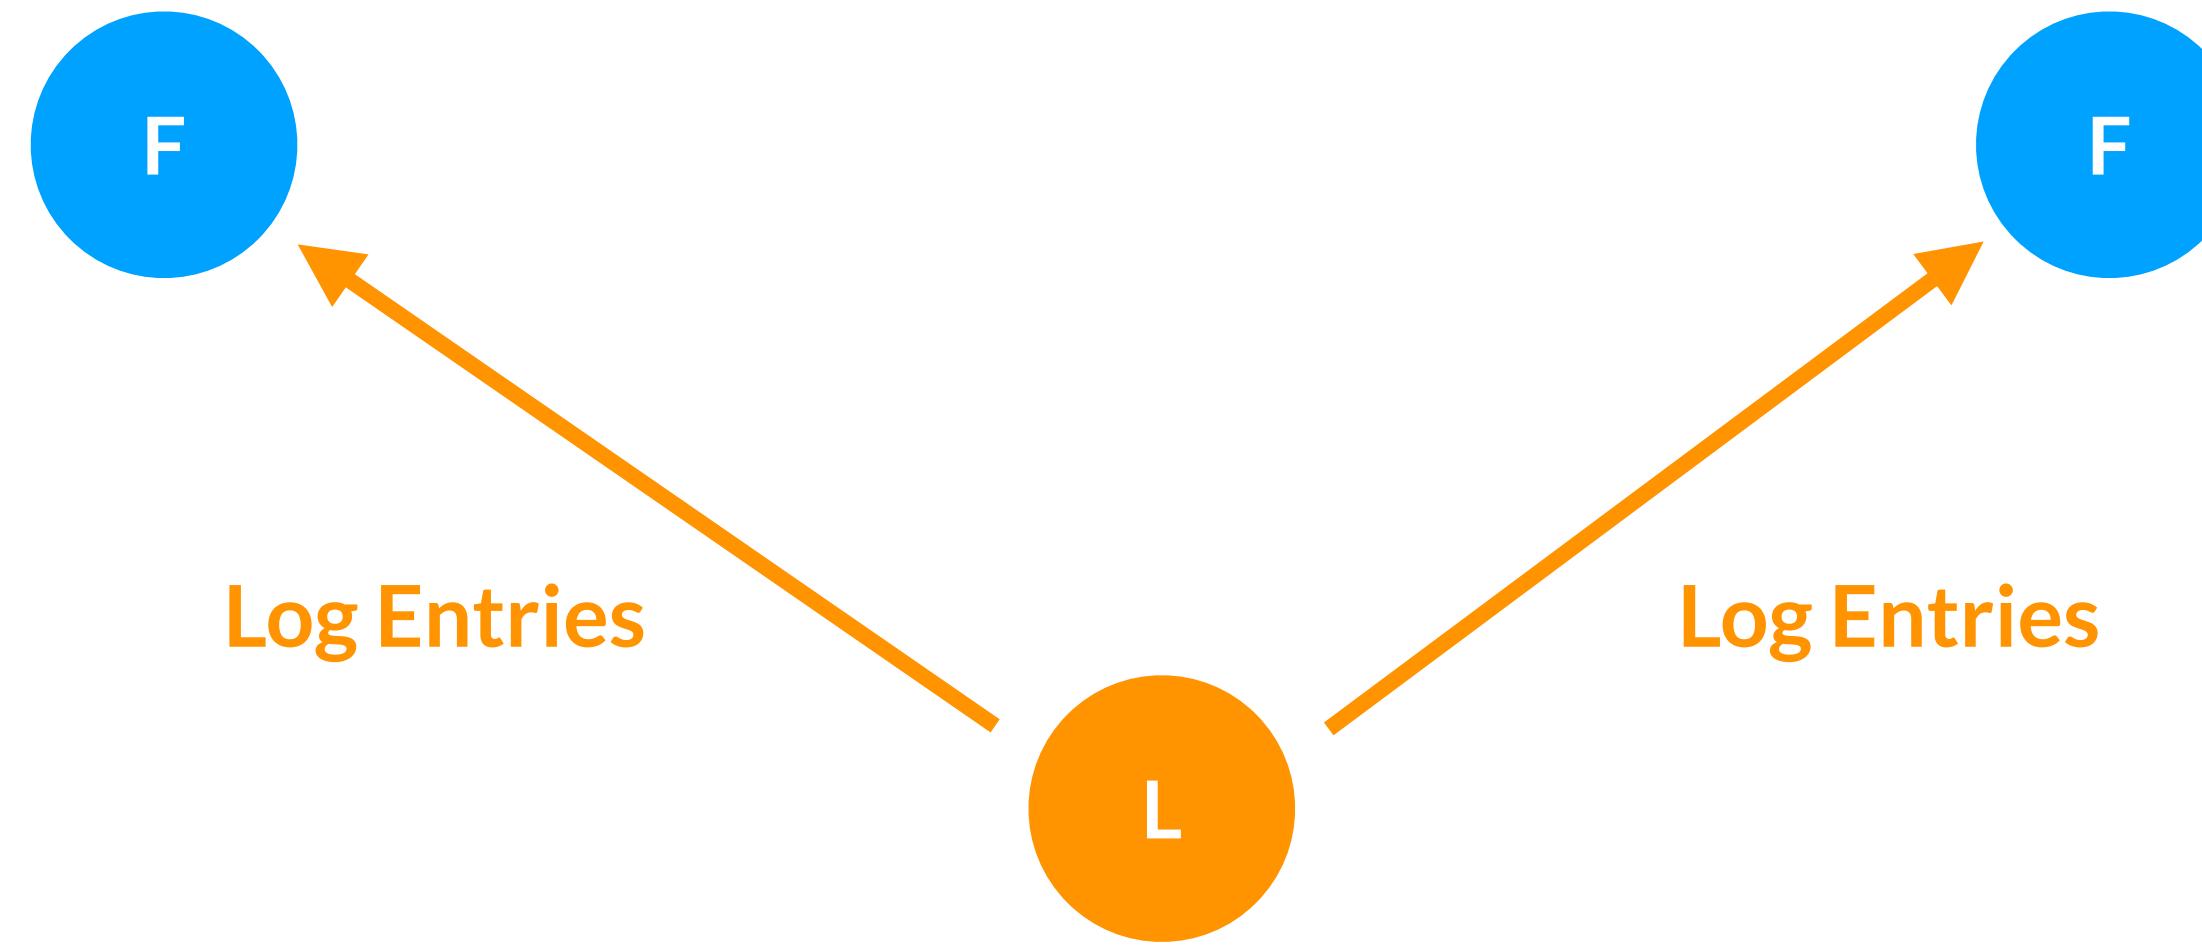




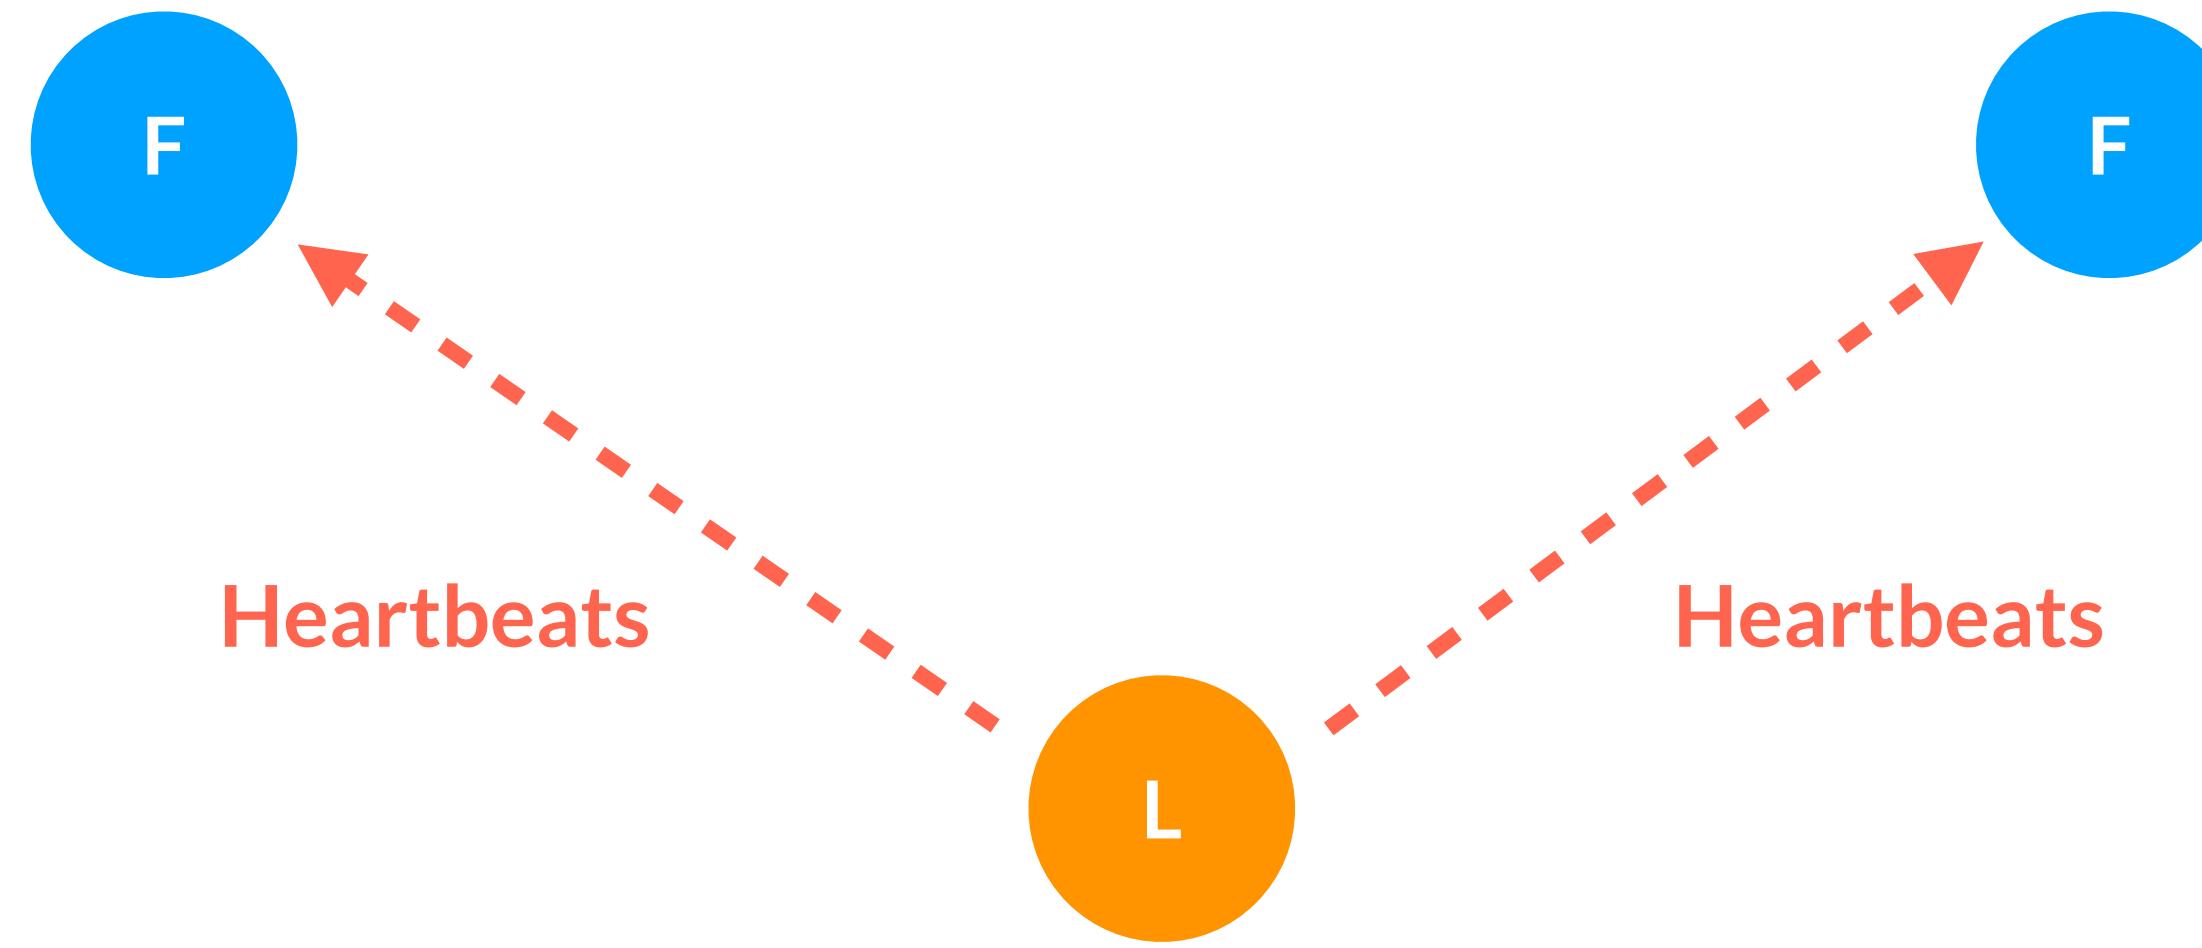









State changes are actively communicated to the followers

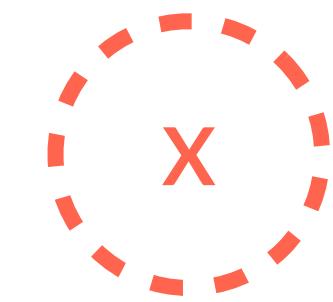
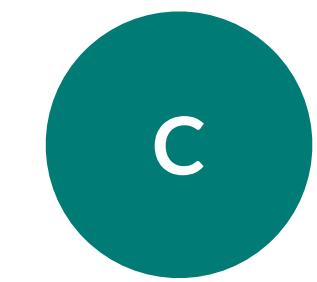


If no active changes to state are being done, then heartbeats are sent


Leader Failure



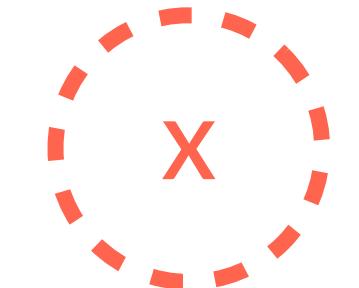
Leader has failed



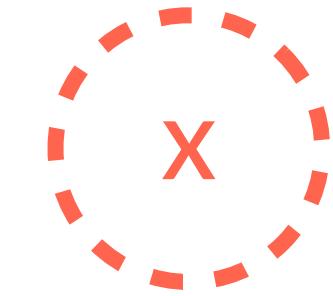
Since each clock is randomized, one will trip. It becomes a candidate



New candidate requests that the followers vote for it



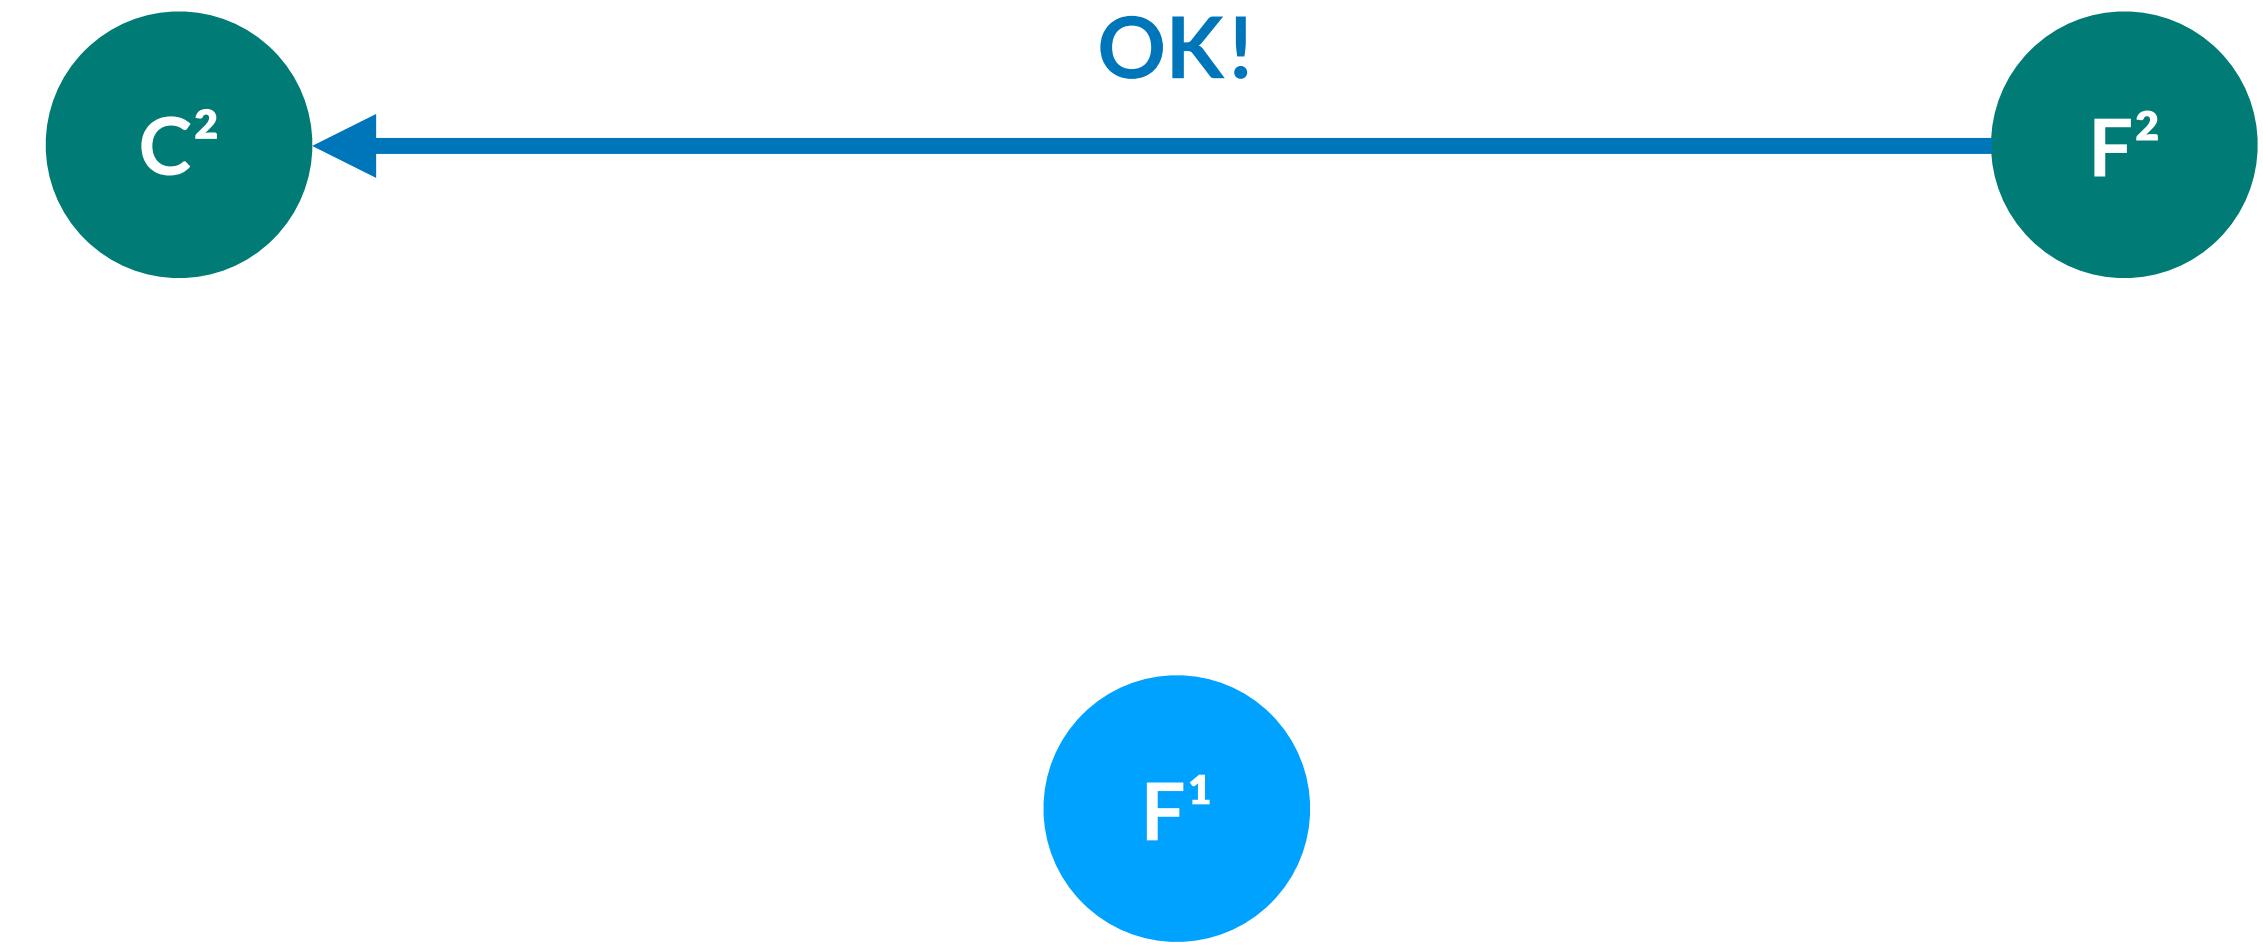
Follower votes for the candidate

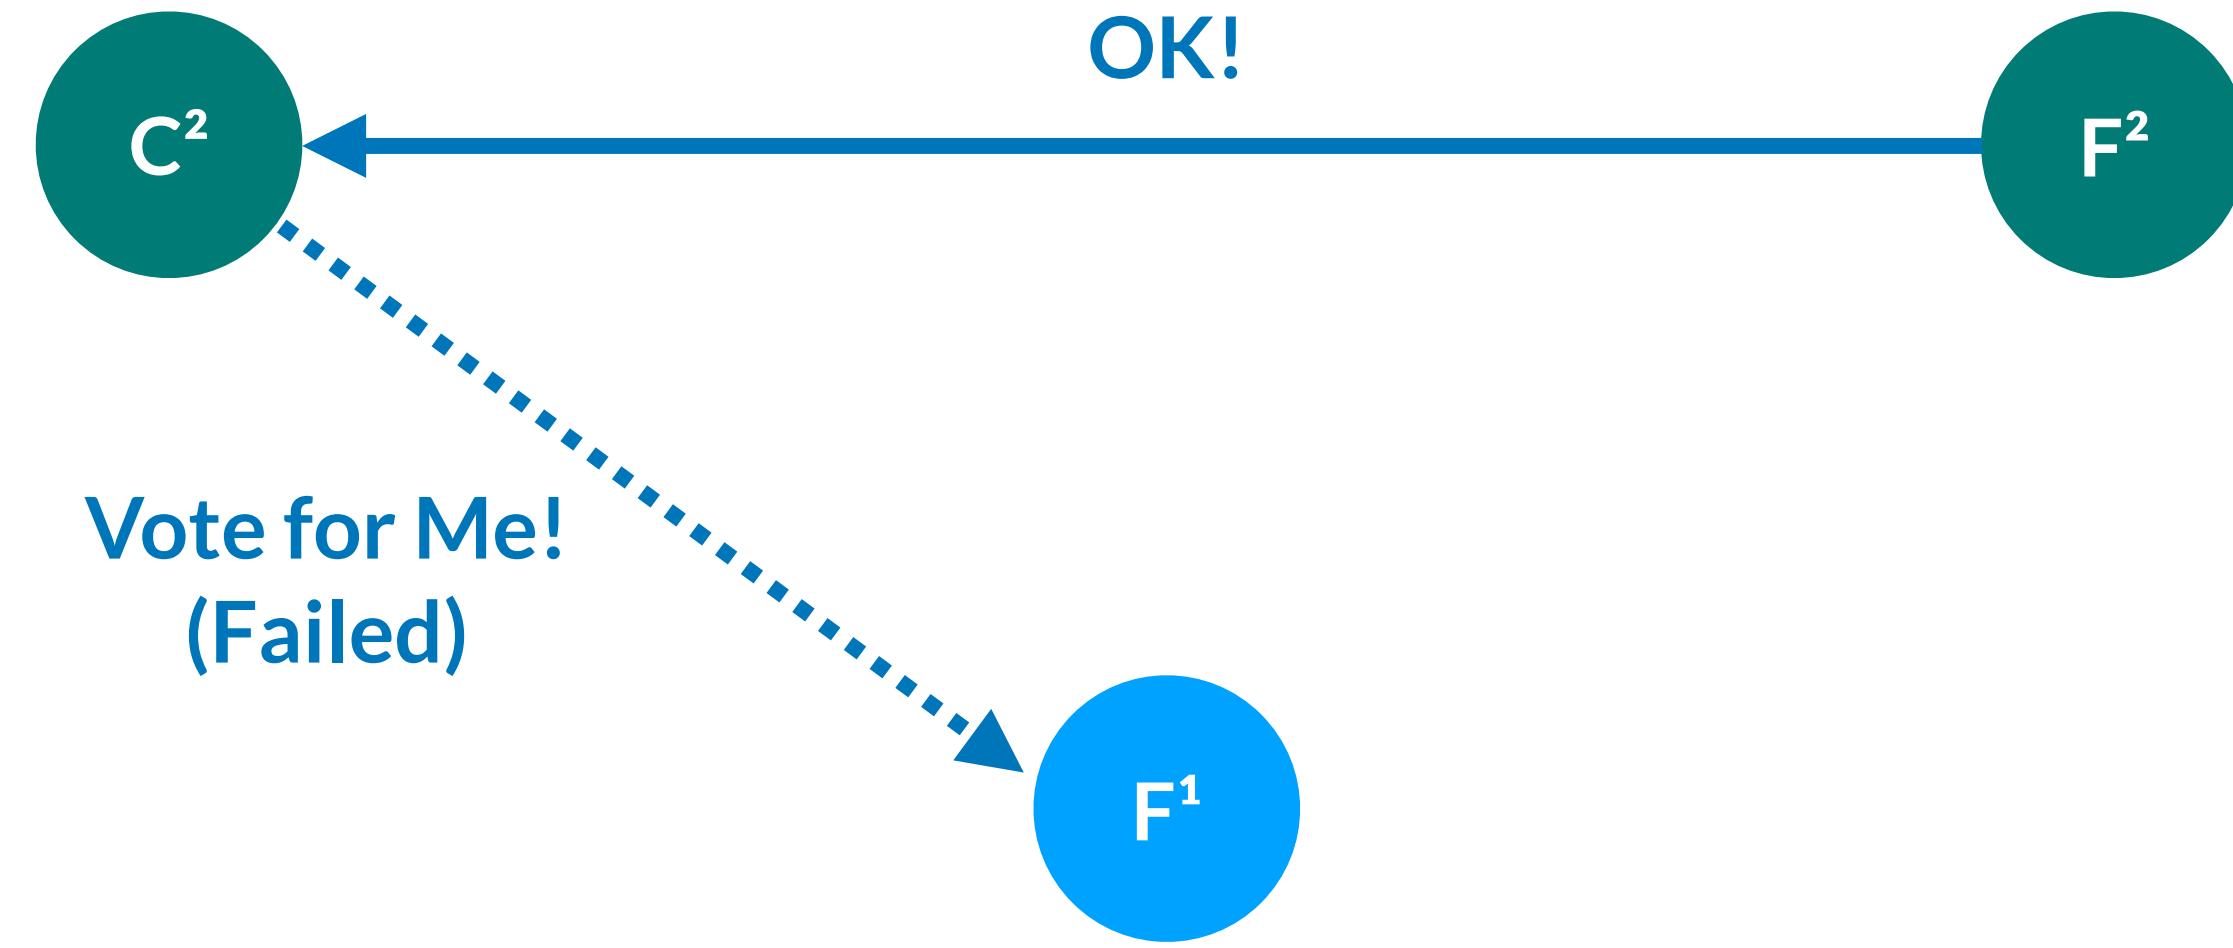


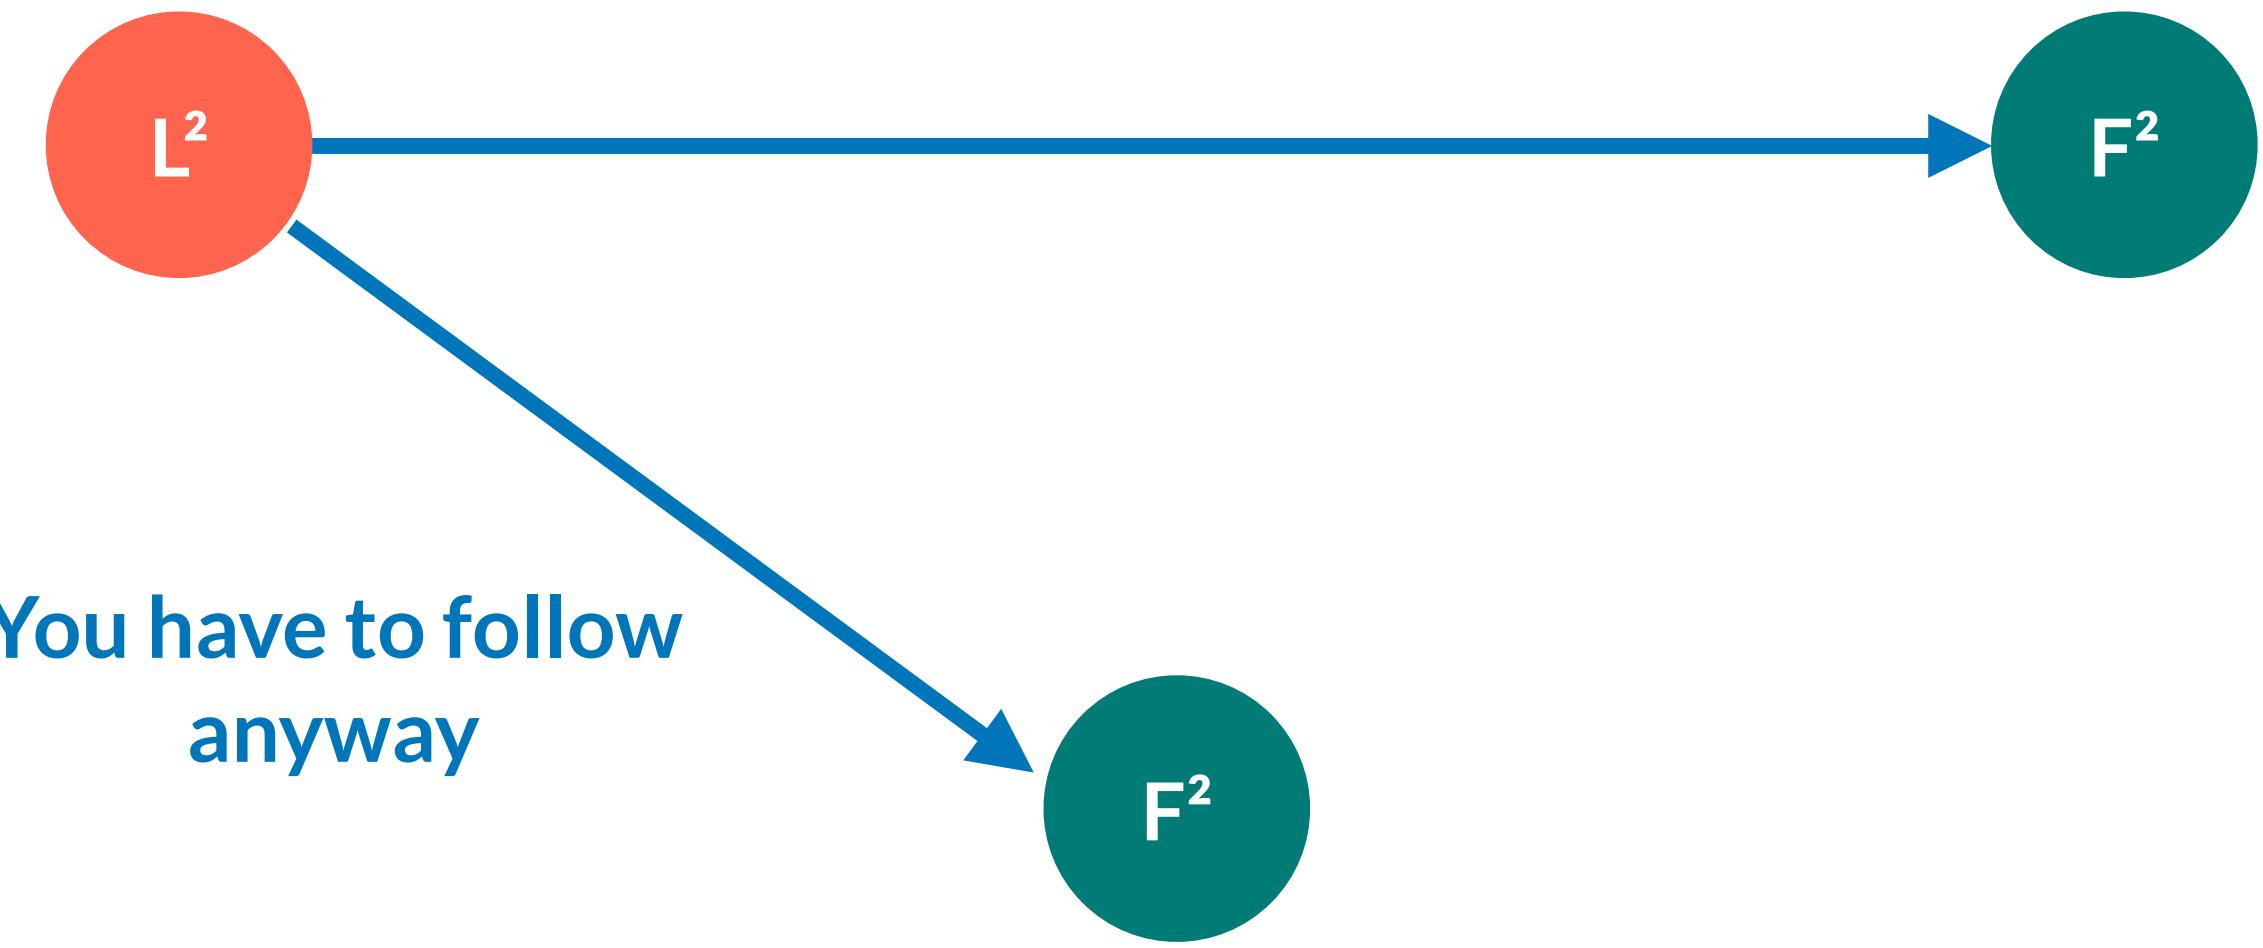
New Leader takes control. It had two votes, itself and the follower,
it sends state changes and heartbeats

Term Elections

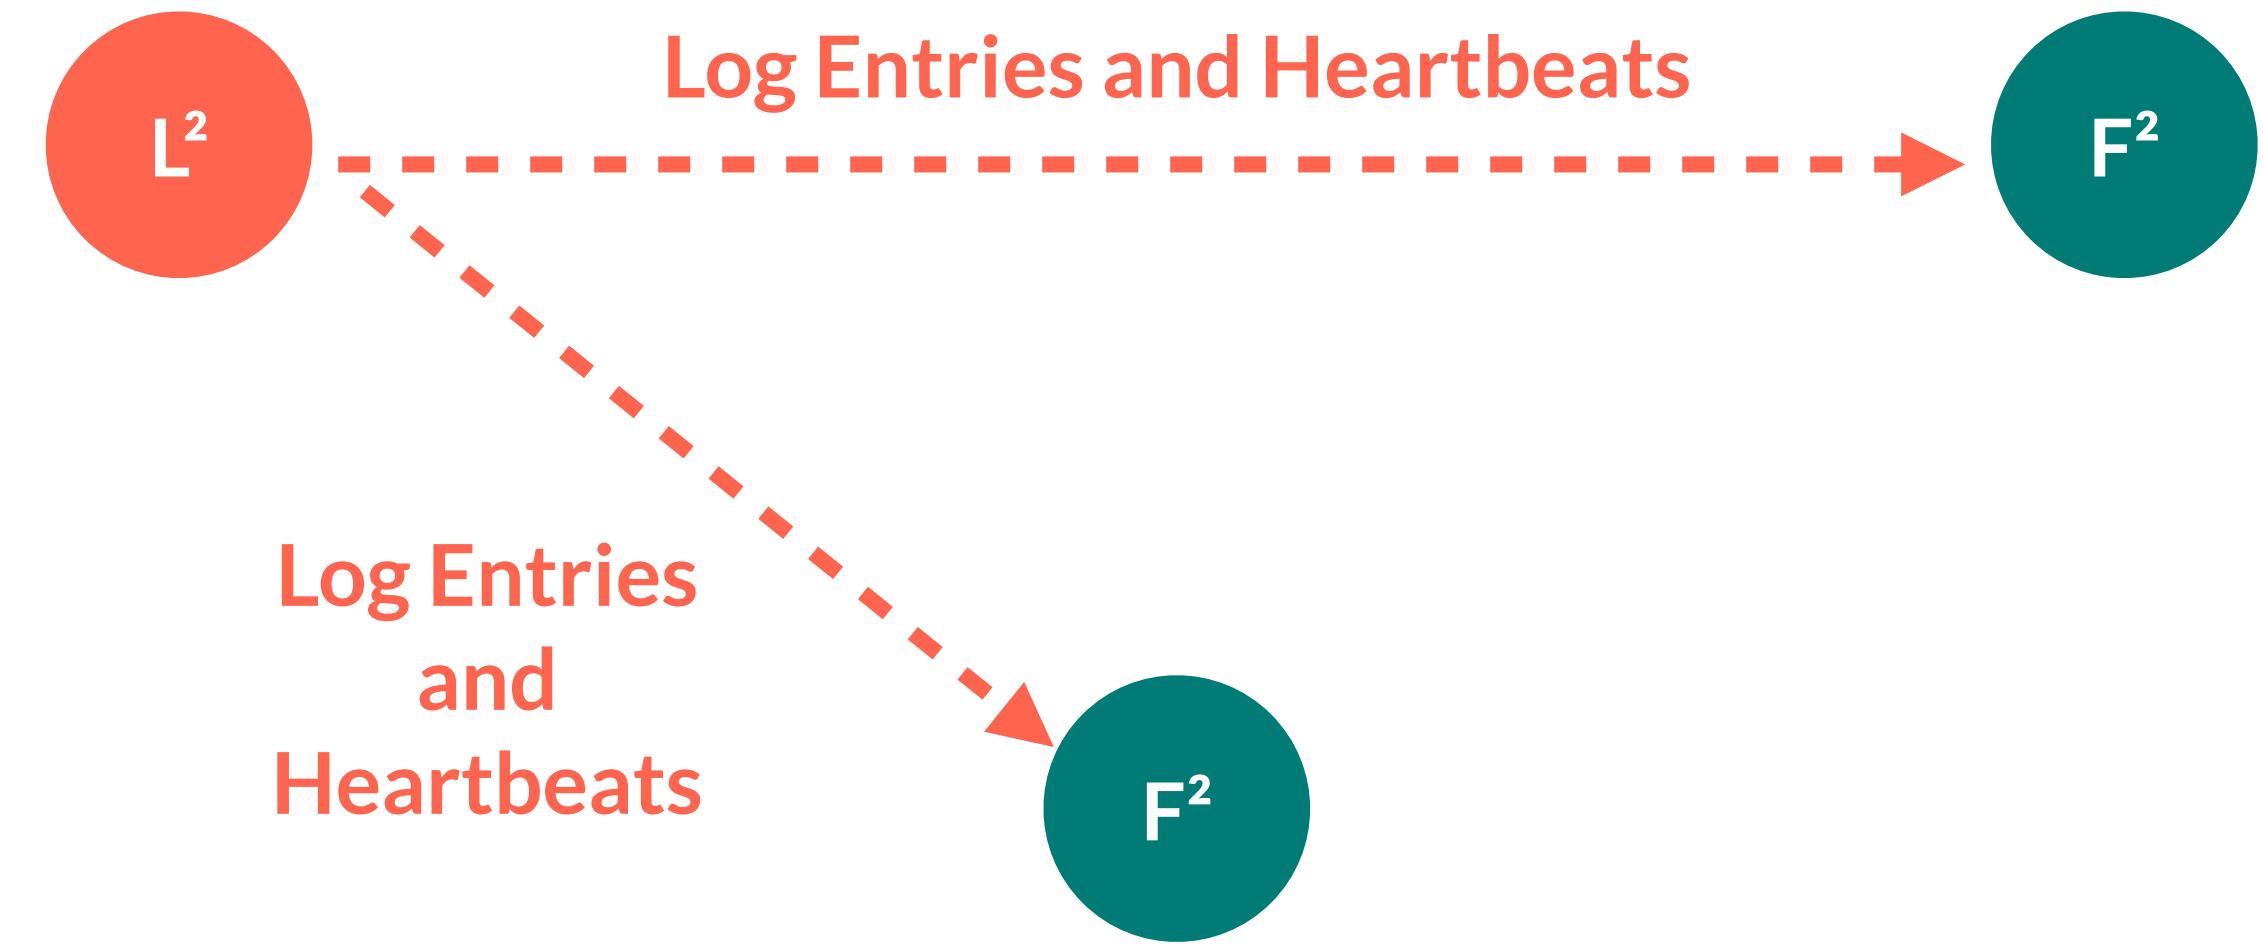




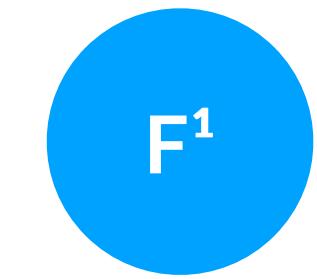




Both are now at Term 2



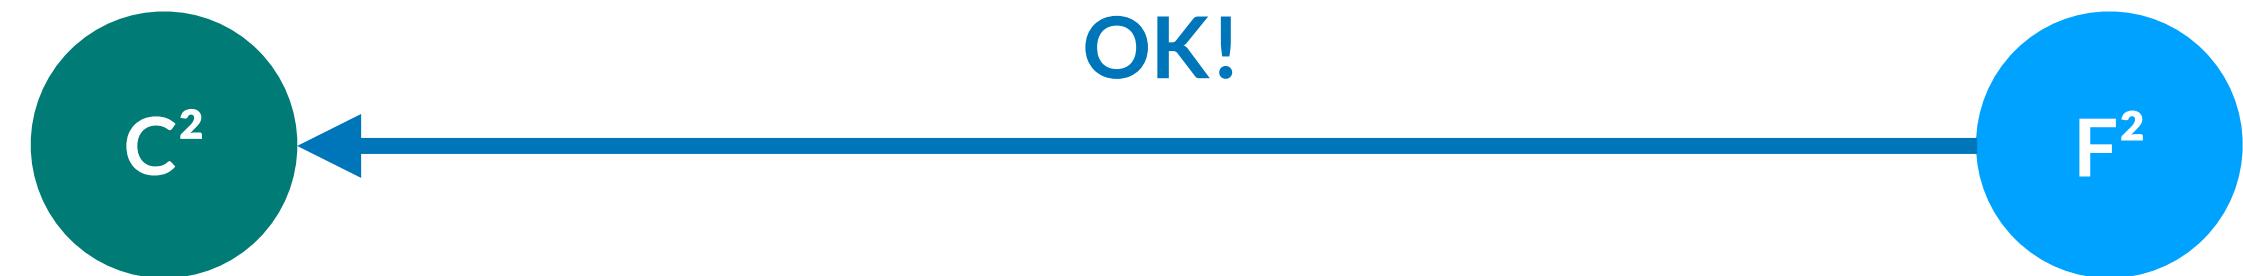
Split Votes



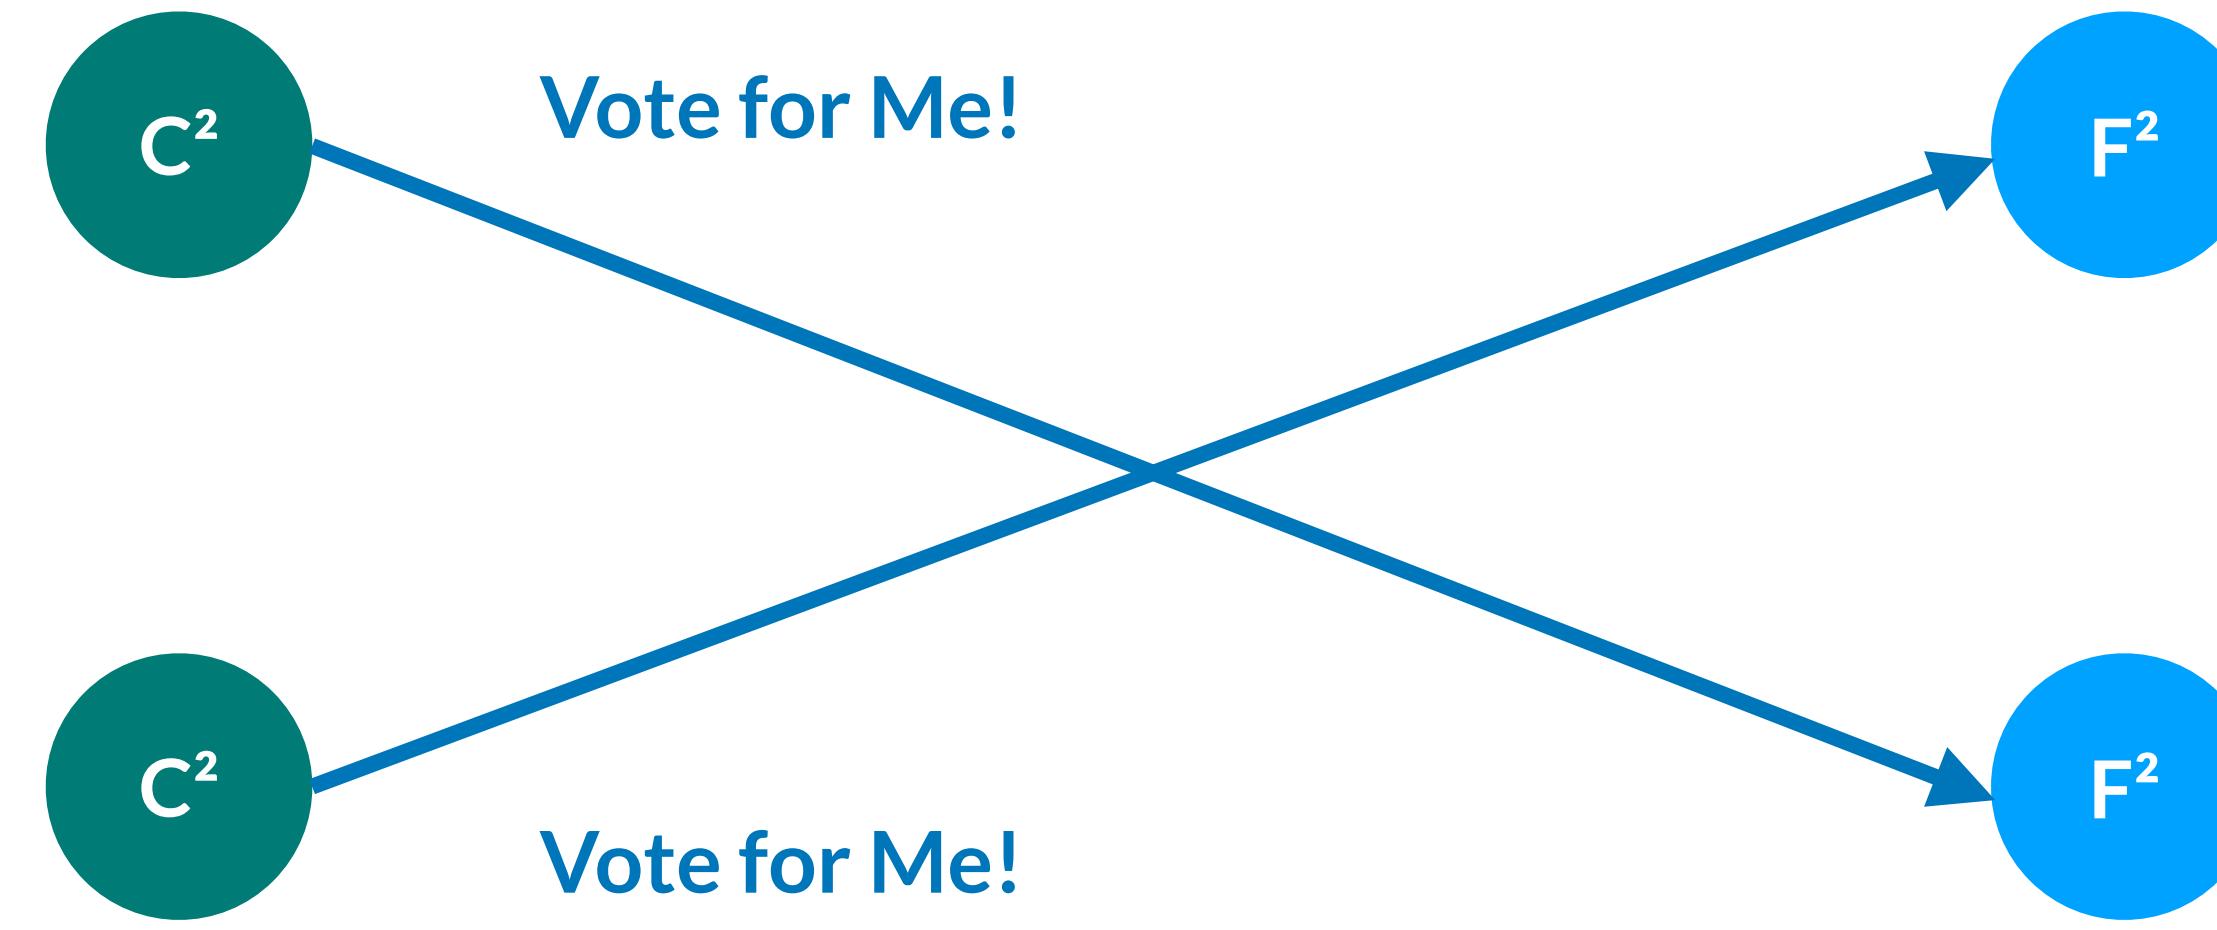
**Two followers become candidates
and being requesting votes**



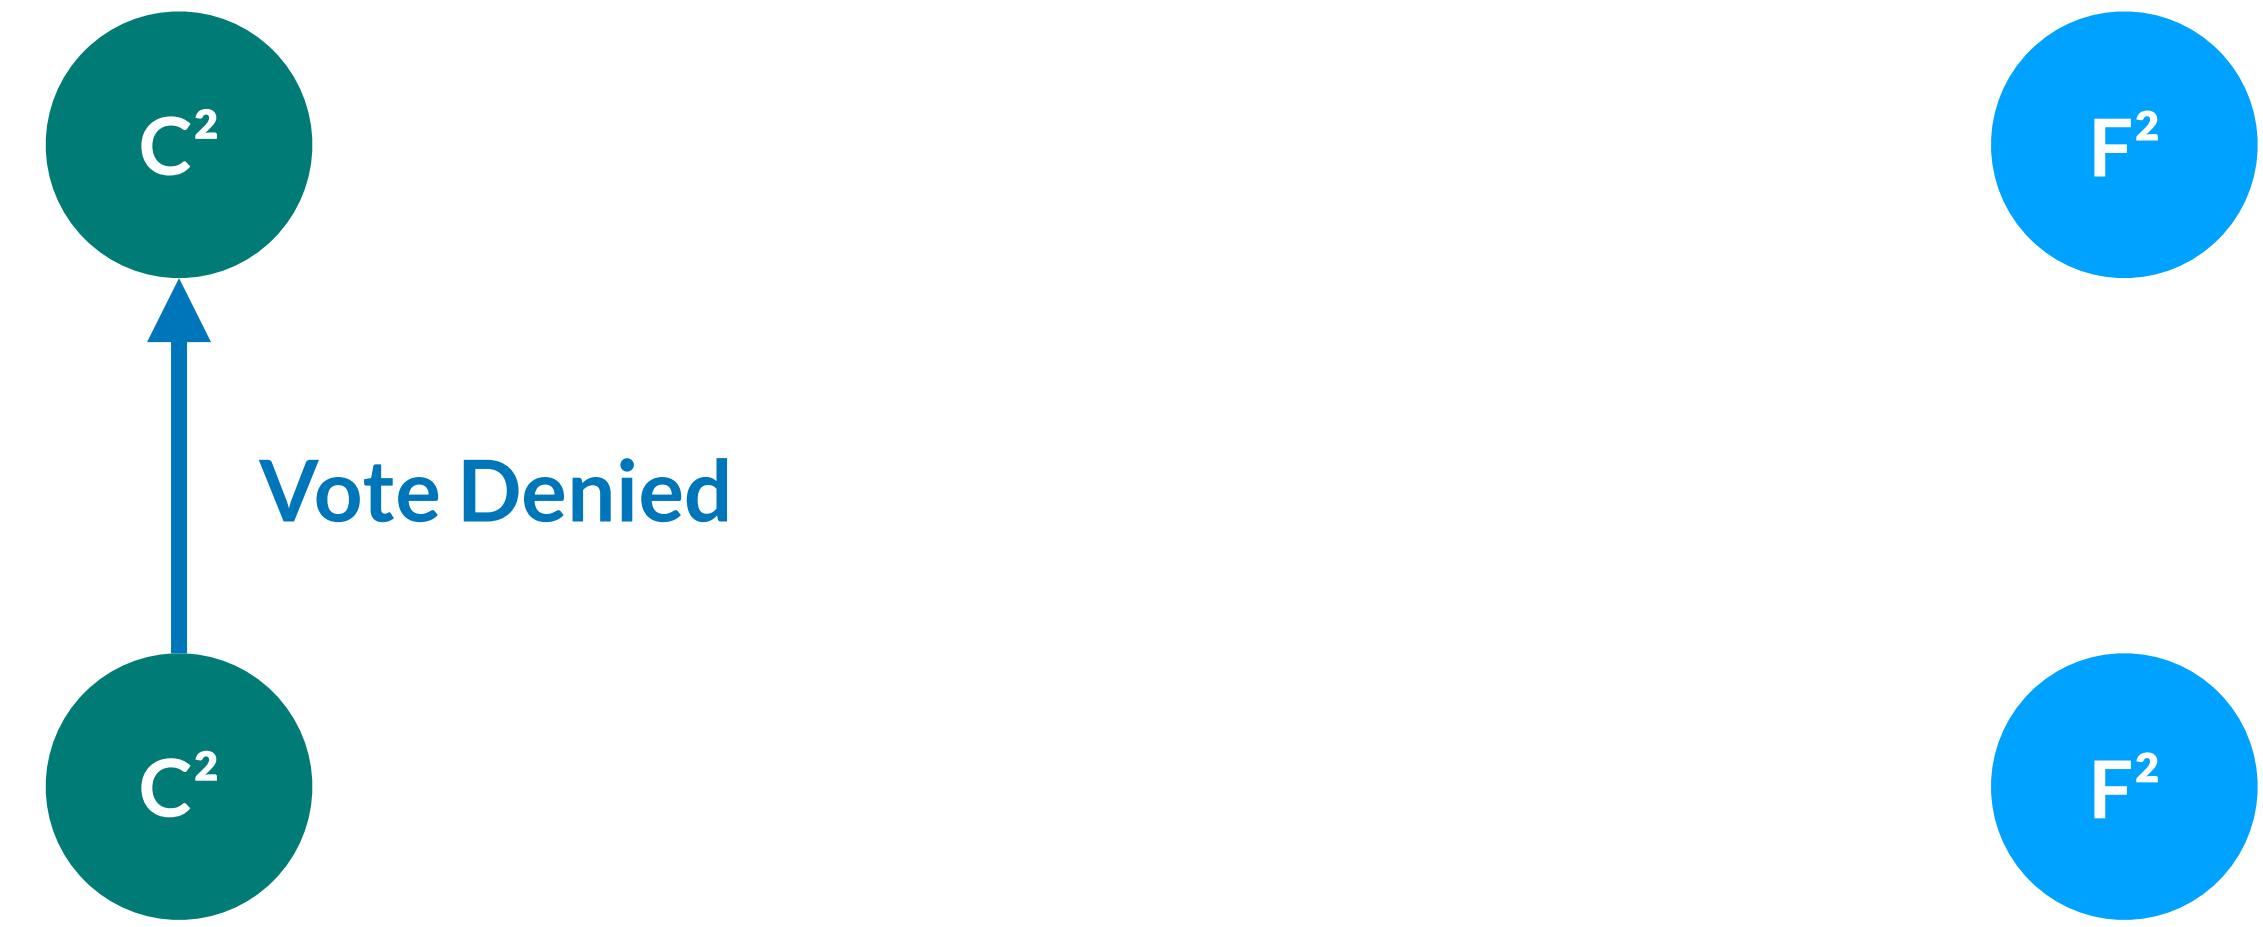
**Two followers become candidates
and being requesting votes**



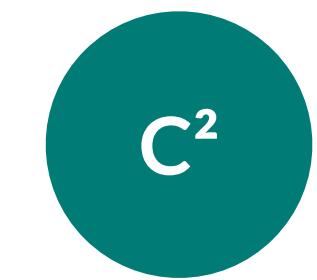
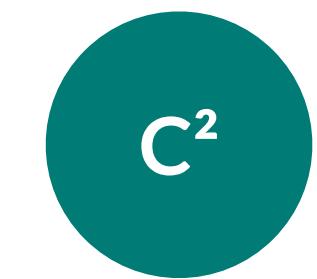
**Each candidate receives a vote from themselves
and one follower peer, but that is not enough**



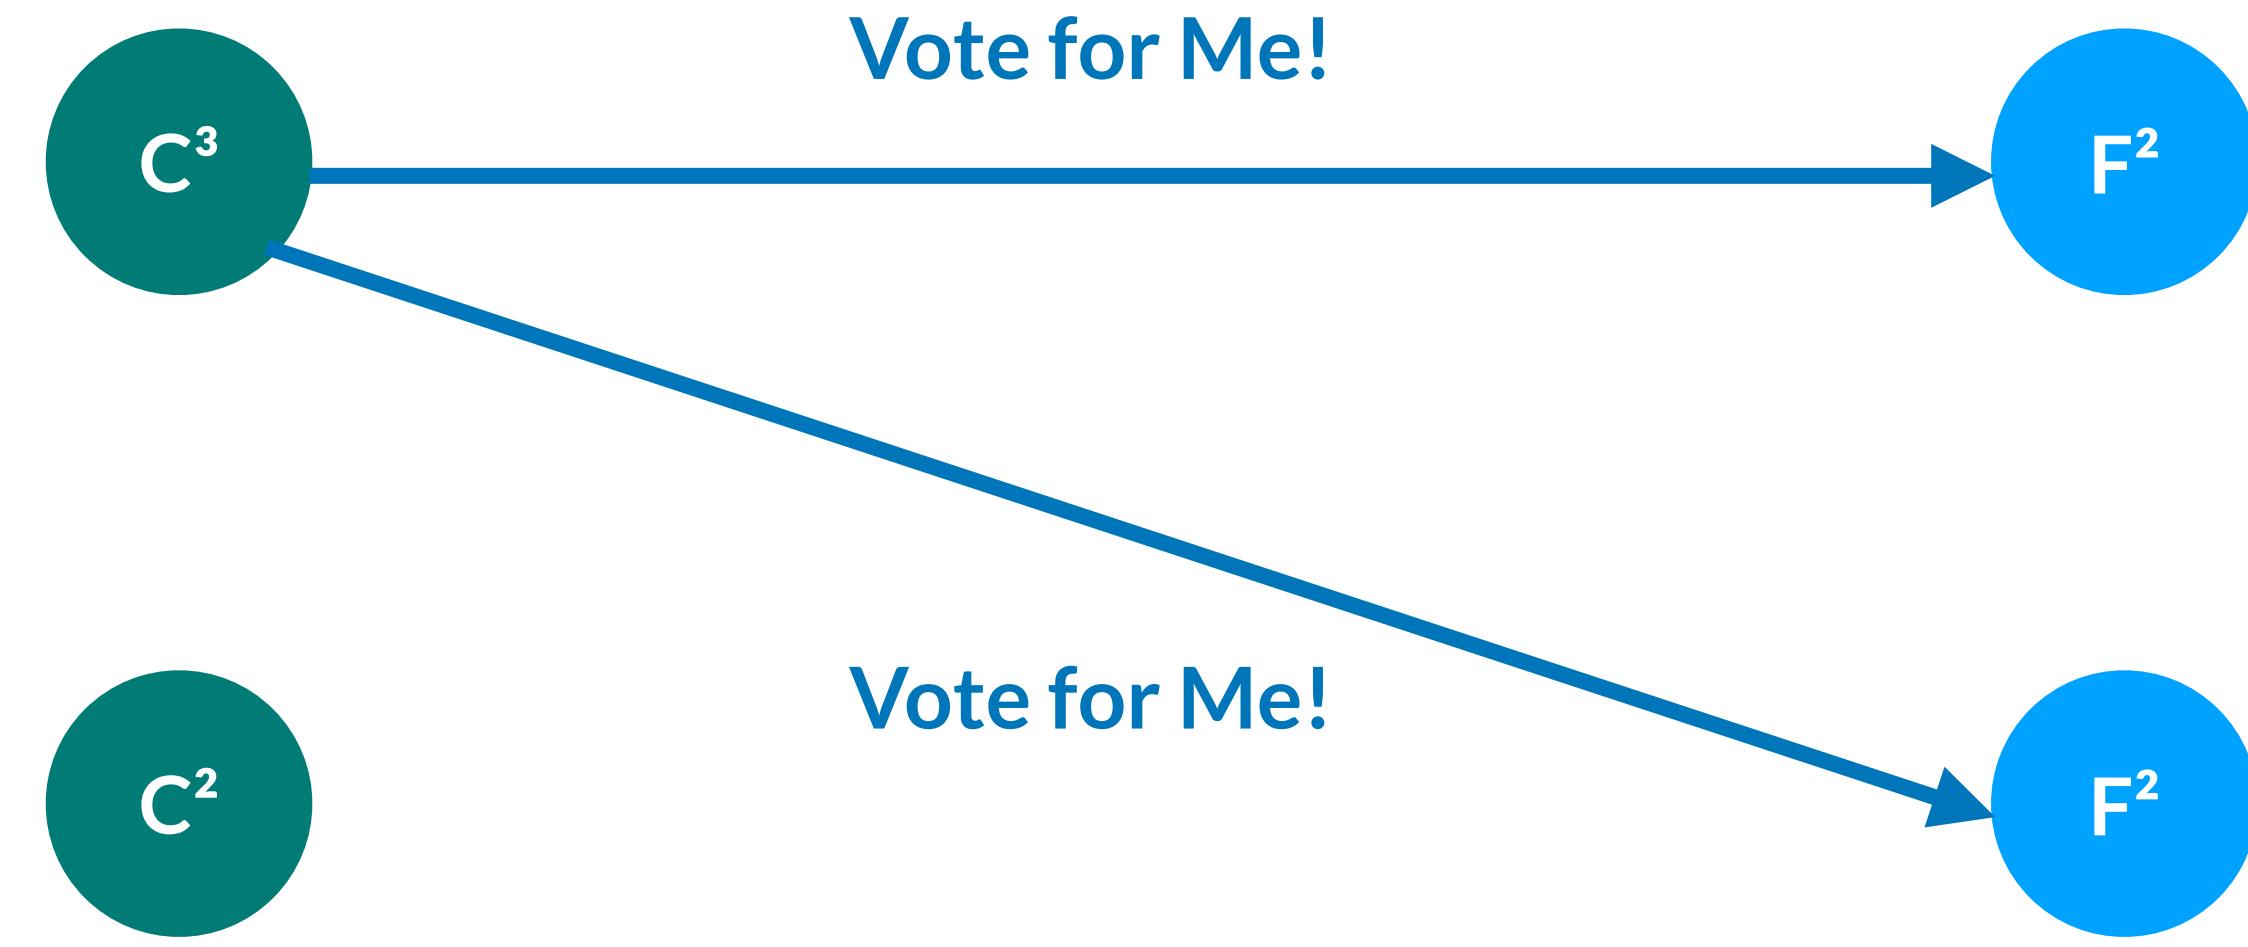
Each candidate now tries to ask for votes from a follower
who has already voted



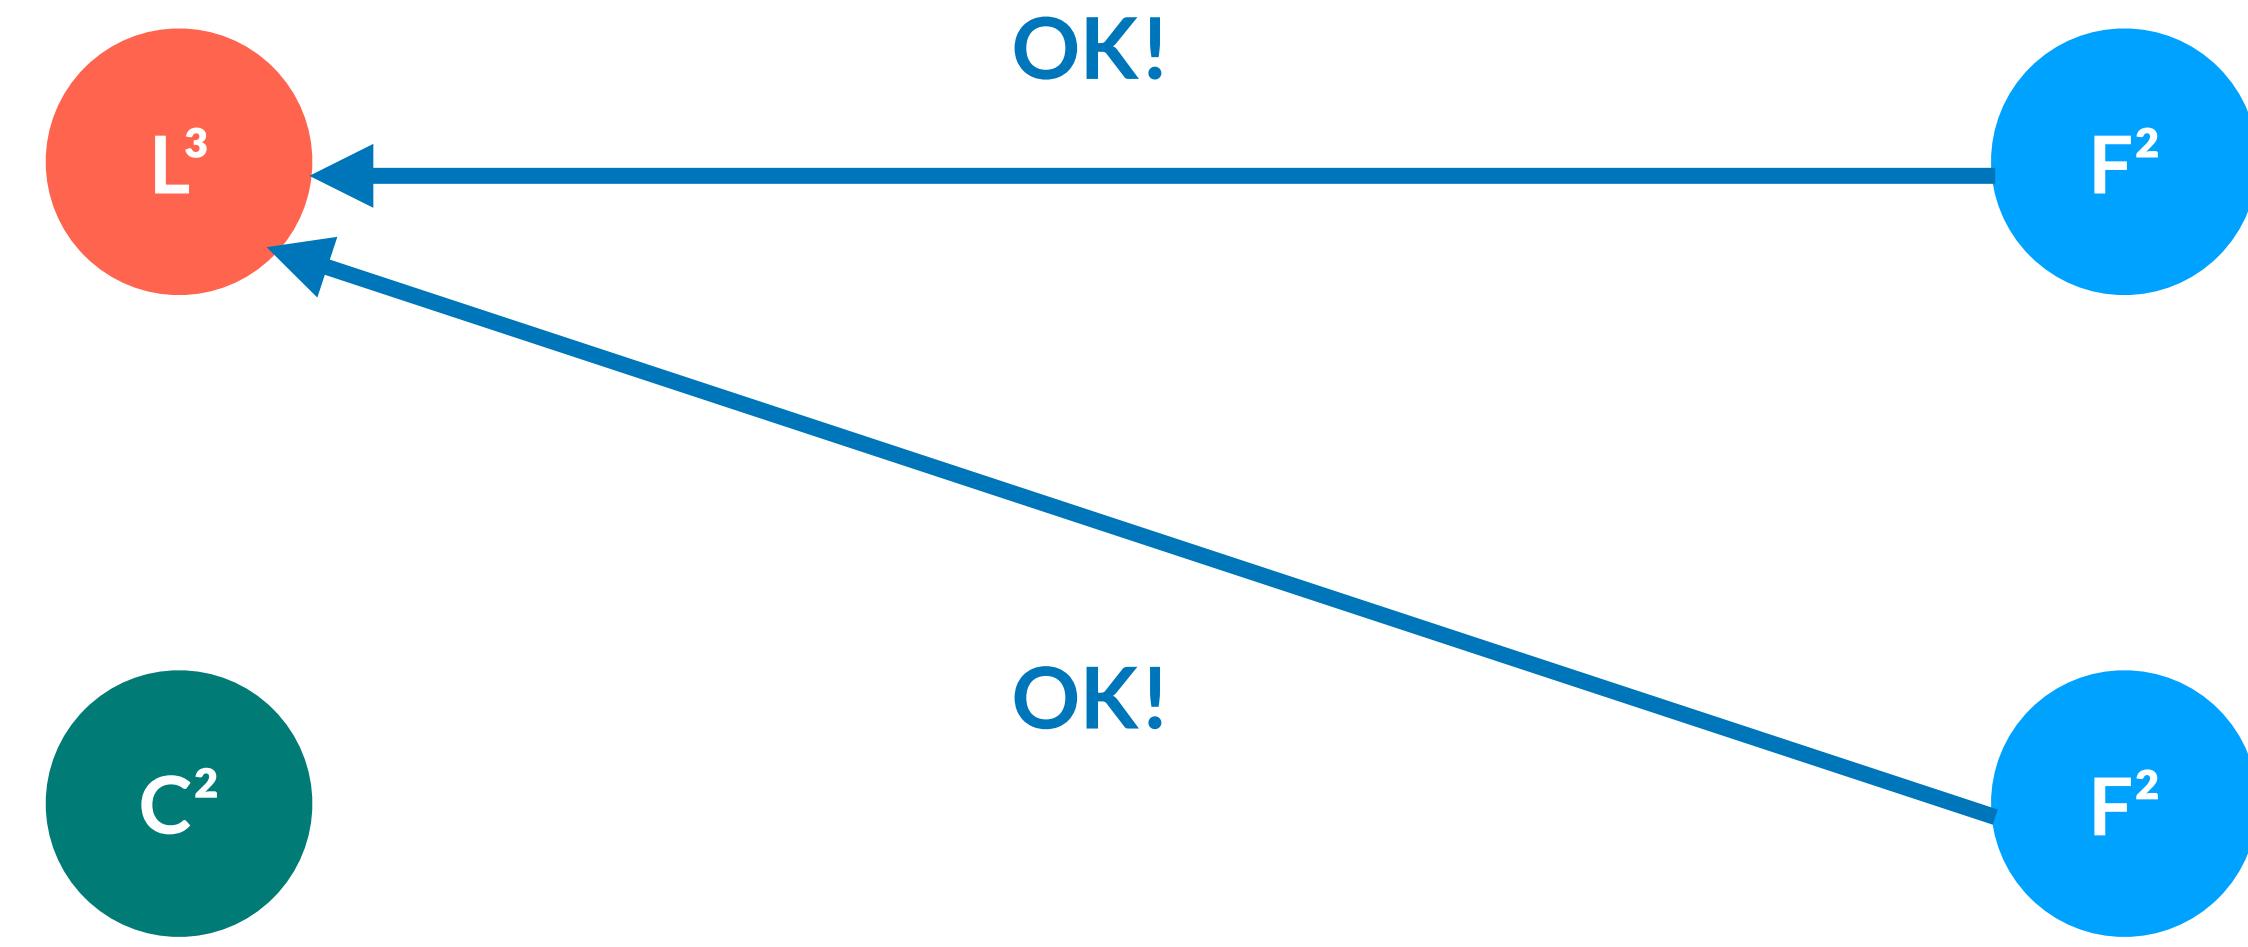
**Vote Requests are denied since they
already voted for themselves**



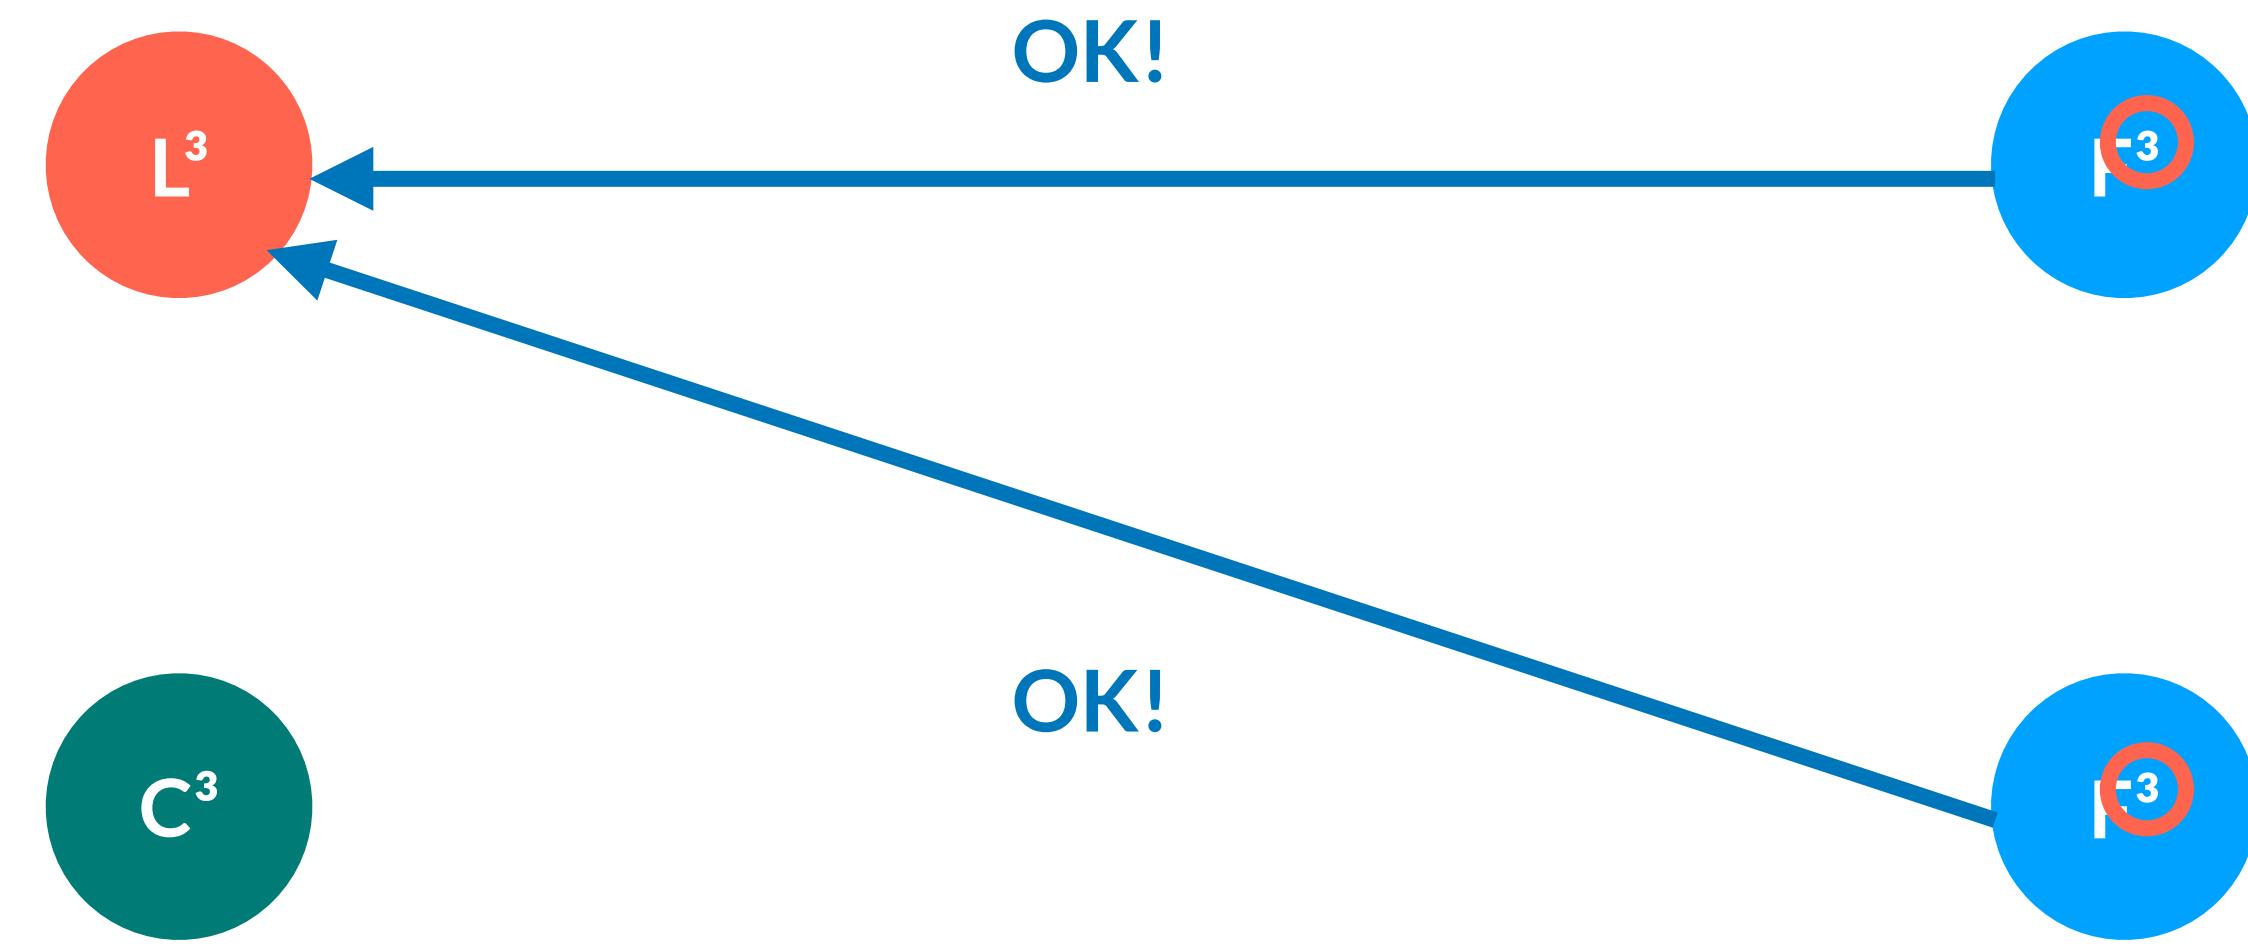
Stalemate. Candidates wait for another random election timeout to occur (150ms - 300ms)



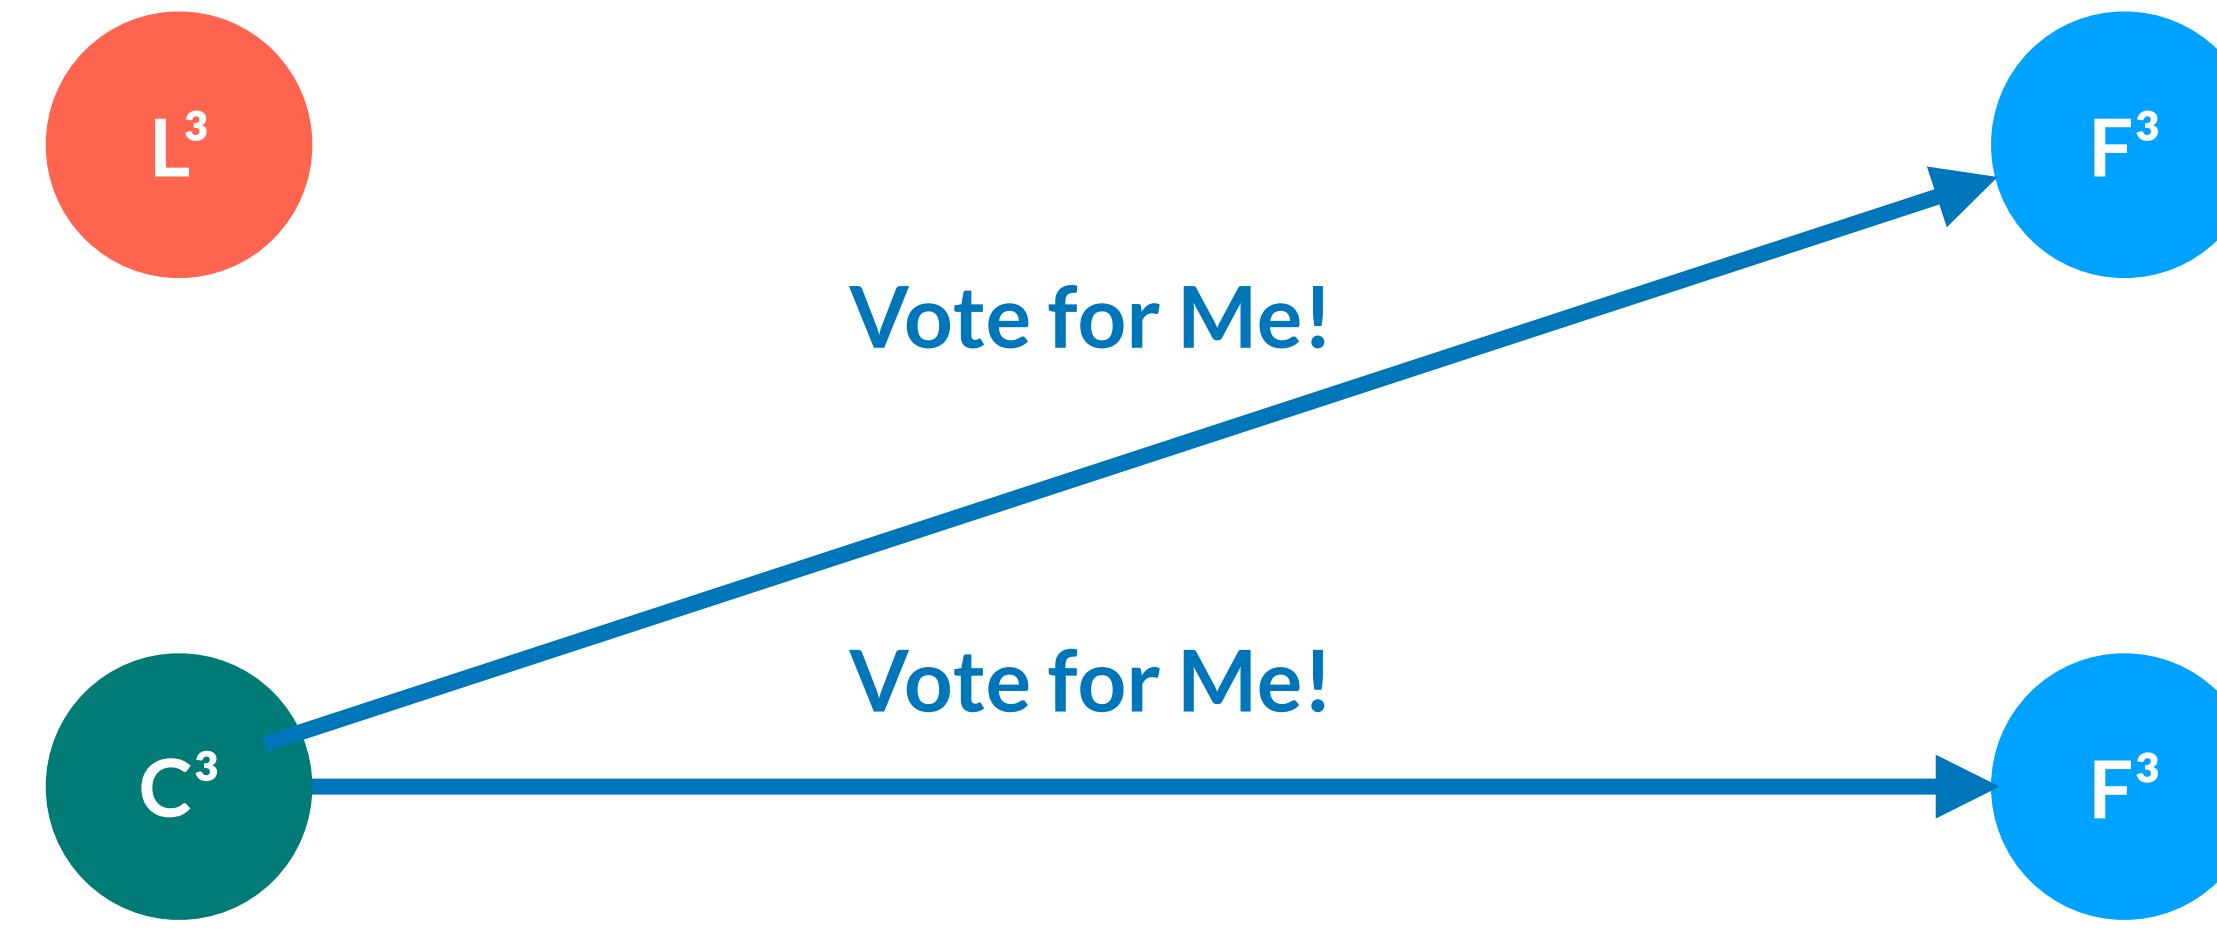
Let's try again. We are now in Term 3



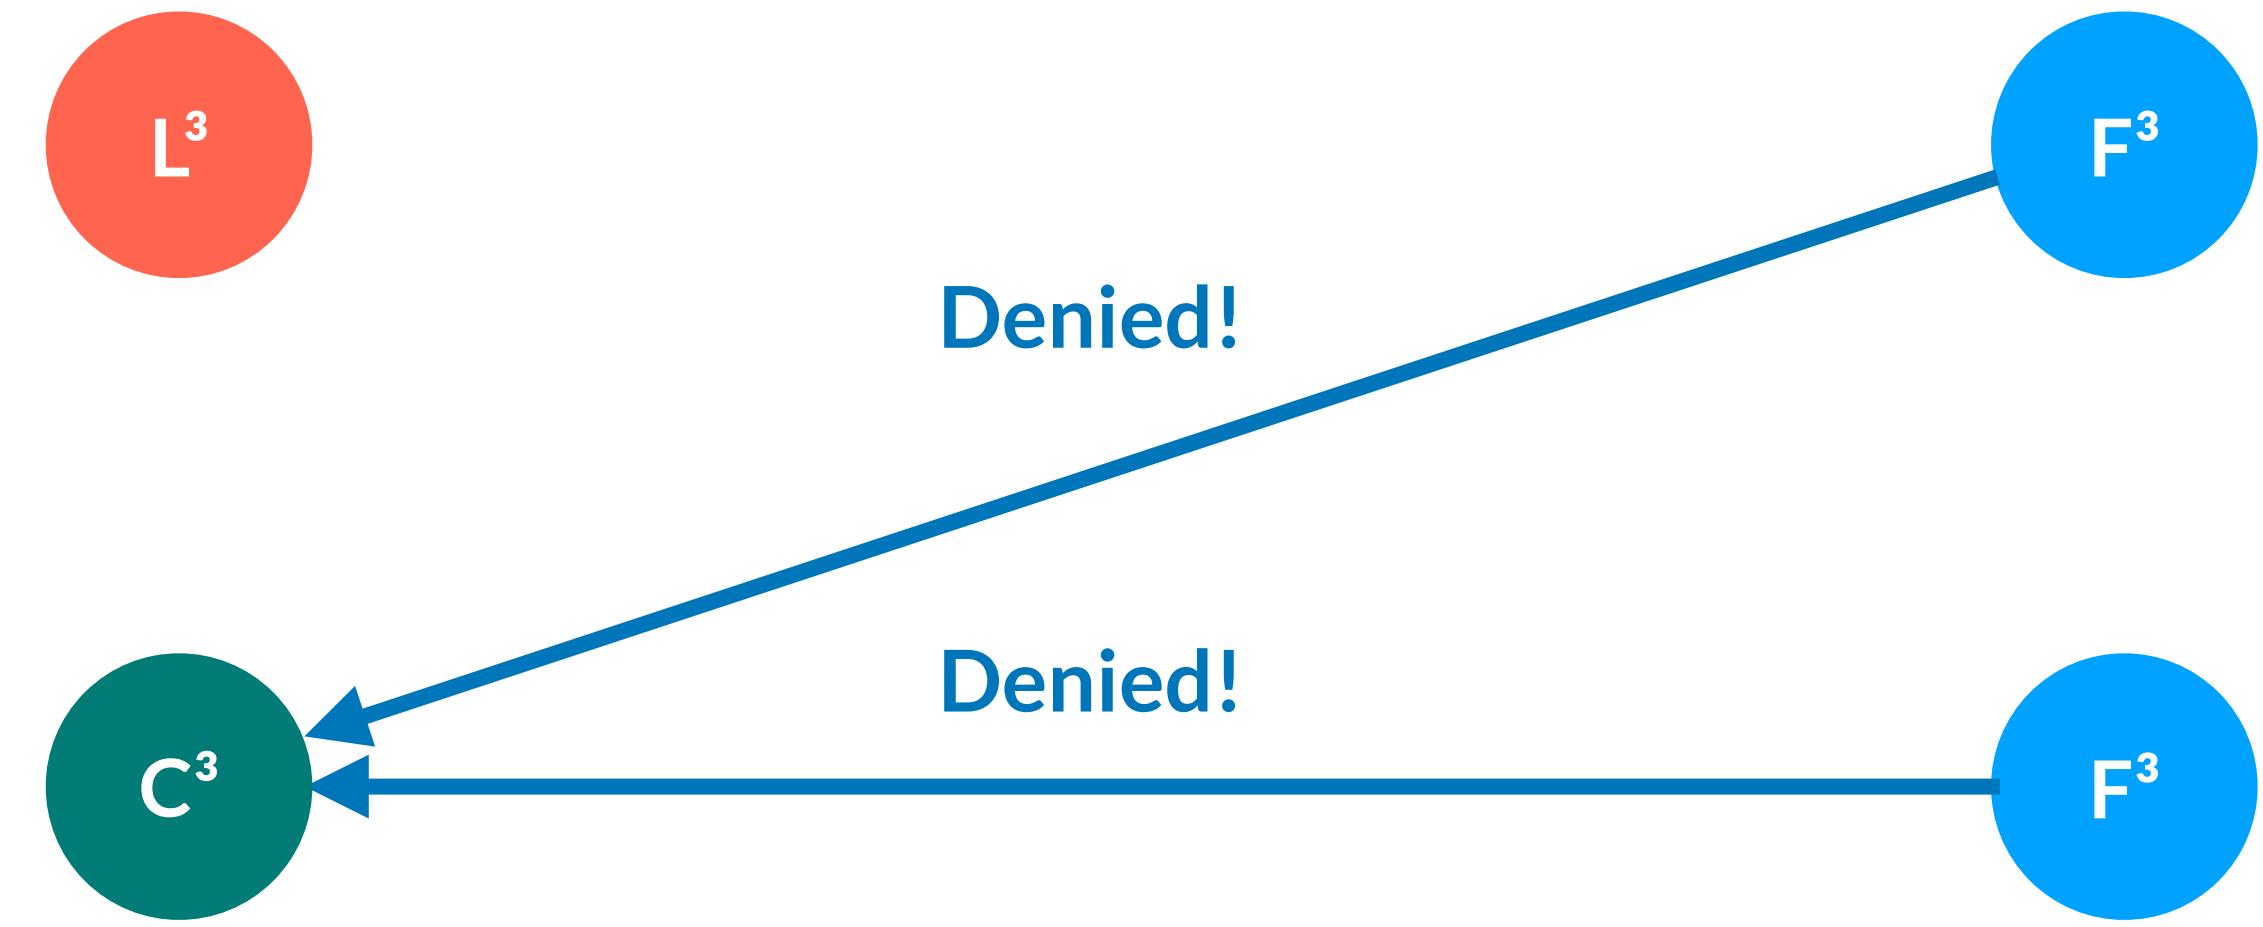
Leader becomes elected, but there is still one candidate!



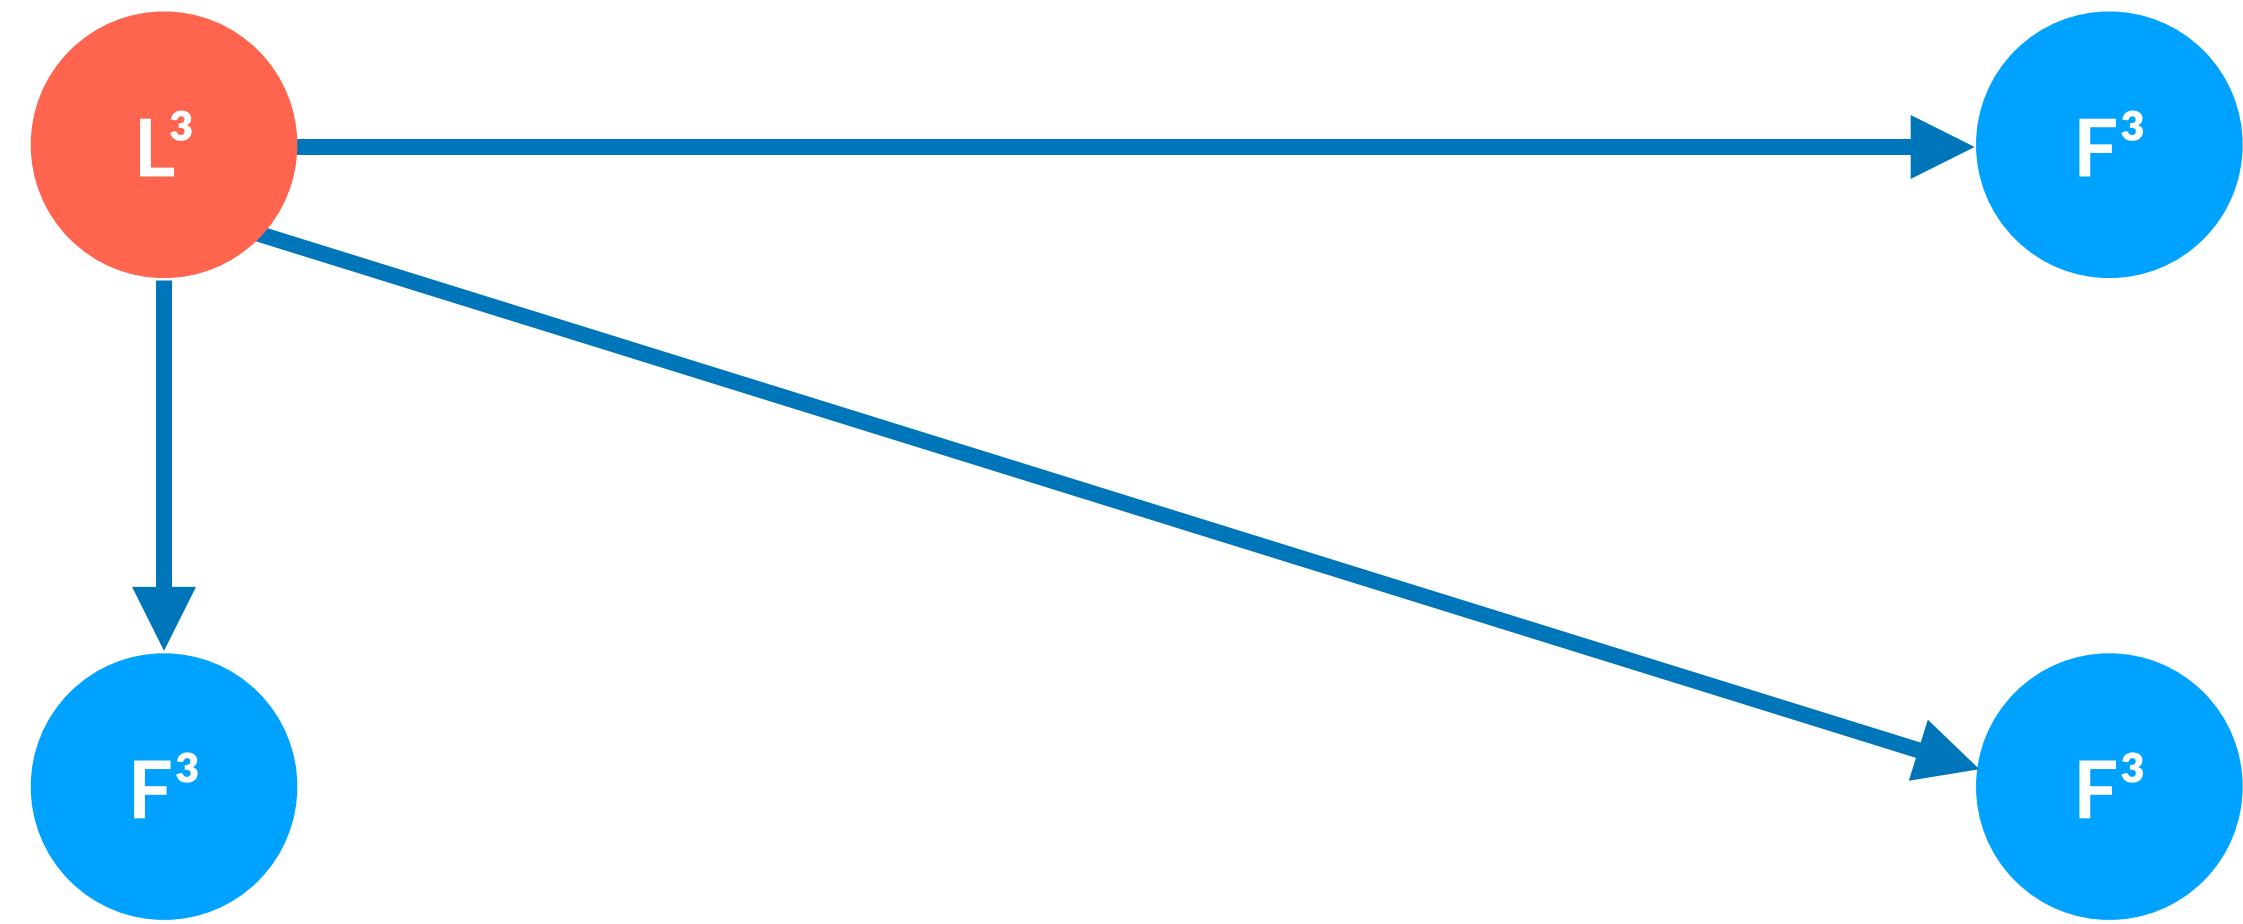
Leader becomes elected, but there is still one candidate!



It will not try to increase its term and run for leader



It will not try to increase its term and run for leader



Leader 3 will assert itself as the leader

Majority Based Failure

Raft is majority based

- If you have 3 systems and 2 goes down, you are out!
- If you have 5 systems and 3 goes down, you are out!
- KRaft operates only as three nodes currently

Requirements



Hardware Requirements

- Do not co-locate other applications due to memory page cache pollution and will degrade performance
- Performant drive is required
- Storage capacity will need to be calculated by the expected messages per day and retention
- Slower networks can degrade the rate in which messages are produced

Cloud Requirements

- Analyze by data retention
- Analyze performance need by the producers
- If low latency is required, SSD should be considered
- Ephemeral Storage may be required (Elastic Block Storage)

Kafka Guarantees



Kafka Guarantees

- Messages sent by a producer to a particular topic partition will be appended in the order they are sent.
- A consumer instance sees records in the order they are stored in the log.
- For a topic with replication factor N, we will tolerate up to N-1 server failures without losing any records committed to the log.



Kafka CLI

Creating a Topic

Creating the topic orders with replication factor 2 and 4 partitions

```
$ /usr/bin/kafka-topics --create --bootstrap-server localhost:9092 \  
--replication-factor 2 --partitions 4 --topic orders
```

Listing Topics

List the topics using one of the brokers

```
$ /usr/bin/kafka-topics --bootstrap-server localhost:9092 --list
```

Sending a Message

Send a message to list of brokers for a particular topic

```
$ /usr/local/kafka/bin/kafka-console-producer.sh \
--broker-list <kafka-broker>:9092 --topic <topic>
> Hello
> I
> am
> sending
> six
> messages
```

Receiving a Message

Receiving the messages from the topic that was posted by the CLI producer

```
$ /usr/local/kafka/bin/kafka-console-consumer.sh \
--bootstrap-server <kafka-broker>:9092 \
--topic <topic> \
--from-beginning

> Hello
> I
> am
> sending
> six
> messages
```

Showing Distributed Partitions

Showing how partitions are distributed

```
$ /usr/local/kafka/bin/kafka-topics.sh --describe \  
--topic <topic-name> \  
--bootstrap-server localhost:9092
```

Kafka Programming Producers

Establishing Properties

- Construct a `java.util.Properties` object
- Provide two or more locations where the bootstrap servers are located
- Provide a Serializer for the key
- Provide a Serializer for the value

```
Properties properties = new Properties();
properties.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
    "broker1:9092,broker2:9092");
properties.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
    StringSerializer.class);
properties.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
    IntegerSerializer.class);
```

Create a Producer Object

- Construct a `org.apache.kafka.clients.producer.KafkaProducer` object
- A Kafka Producer is thread-safe
- Inject the Properties into the Producer object

```
KafkaProducer<String, Integer> producer = new KafkaProducer<>(properties);
```

Creating a Record/Message

- Create a `org.apache.kafka.clients.producer.ProducerRecord`
- Accepts many parameters but the main ones are:
 - Topic
 - Key (if applicable)
 - Value

```
ProducerRecord<String, Integer> producerRecord =  
    new ProducerRecord<>("my_orders", state, amount);
```

Sending a message

- Send the record by calling send on the Producer
- Returns a Future object to process the results on another Thread.
- The Future will contain RecordMetadata, an object that has information about your send.

```
Future<RecordMetadata> send = producer.send(producerRecord);
```

RecordMetadata

Contains information about your send including the messages

```
if (metadata.hasOffset()) {  
    System.out.format("offset: %d\n",  
        metadata.offset());  
}  
System.out.format("partition: %d\n",  
    metadata.partition());  
System.out.format("timestamp: %d\n",  
    metadata.timestamp());  
System.out.format("topic: %s\n", metadata.topic());  
System.out.format("toString: %s\n",  
    metadata.toString());
```

Sending with a Callback

- Alternately, you can send with a Callback (lambda)
- Callback is an interface, can be used as a lambda
- If RecordMetadata is null there was an error, if Exception is null the send was successful

```
producer.send(producerRecord, new Callback() {  
    @Override  
    public void onCompletion(RecordMetadata metadata,  
                            Exception e) {  
        ...  
    }  
}
```

Using a Closure to capture Key and Value

- RecordMetadata does not have information on key and value
- Using a closure you can obtain that in the block of your lambda

```
producer.send(producerRecord, (metadata, e) -> {
    if (metadata != null) {
        System.out.println(producerRecord.key());
        System.out.println(producerRecord.value());
        System.out.println(metadata.offset());
    }
})
```

Be a good citizen, close your resources

- When you need to terminate, flush messages from the bufferpool
- Close the Producer

```
producer.flush();
producer.close();
```

Closing Resources in a Shutdown Hook

- `Runtime.getRuntime().addShutdownHook(...)` will listen for SIGTERM (CTRL+C)
- This would make an excellent place to flush and close, and close any loops that you may have created

```
Runtime.getRuntime().addShutdownHook(new Thread(() -> {
    done.set(true);
    producer.flush();
    producer.close();
}));
```

Lab: Producers

Acknowledgements & Retries

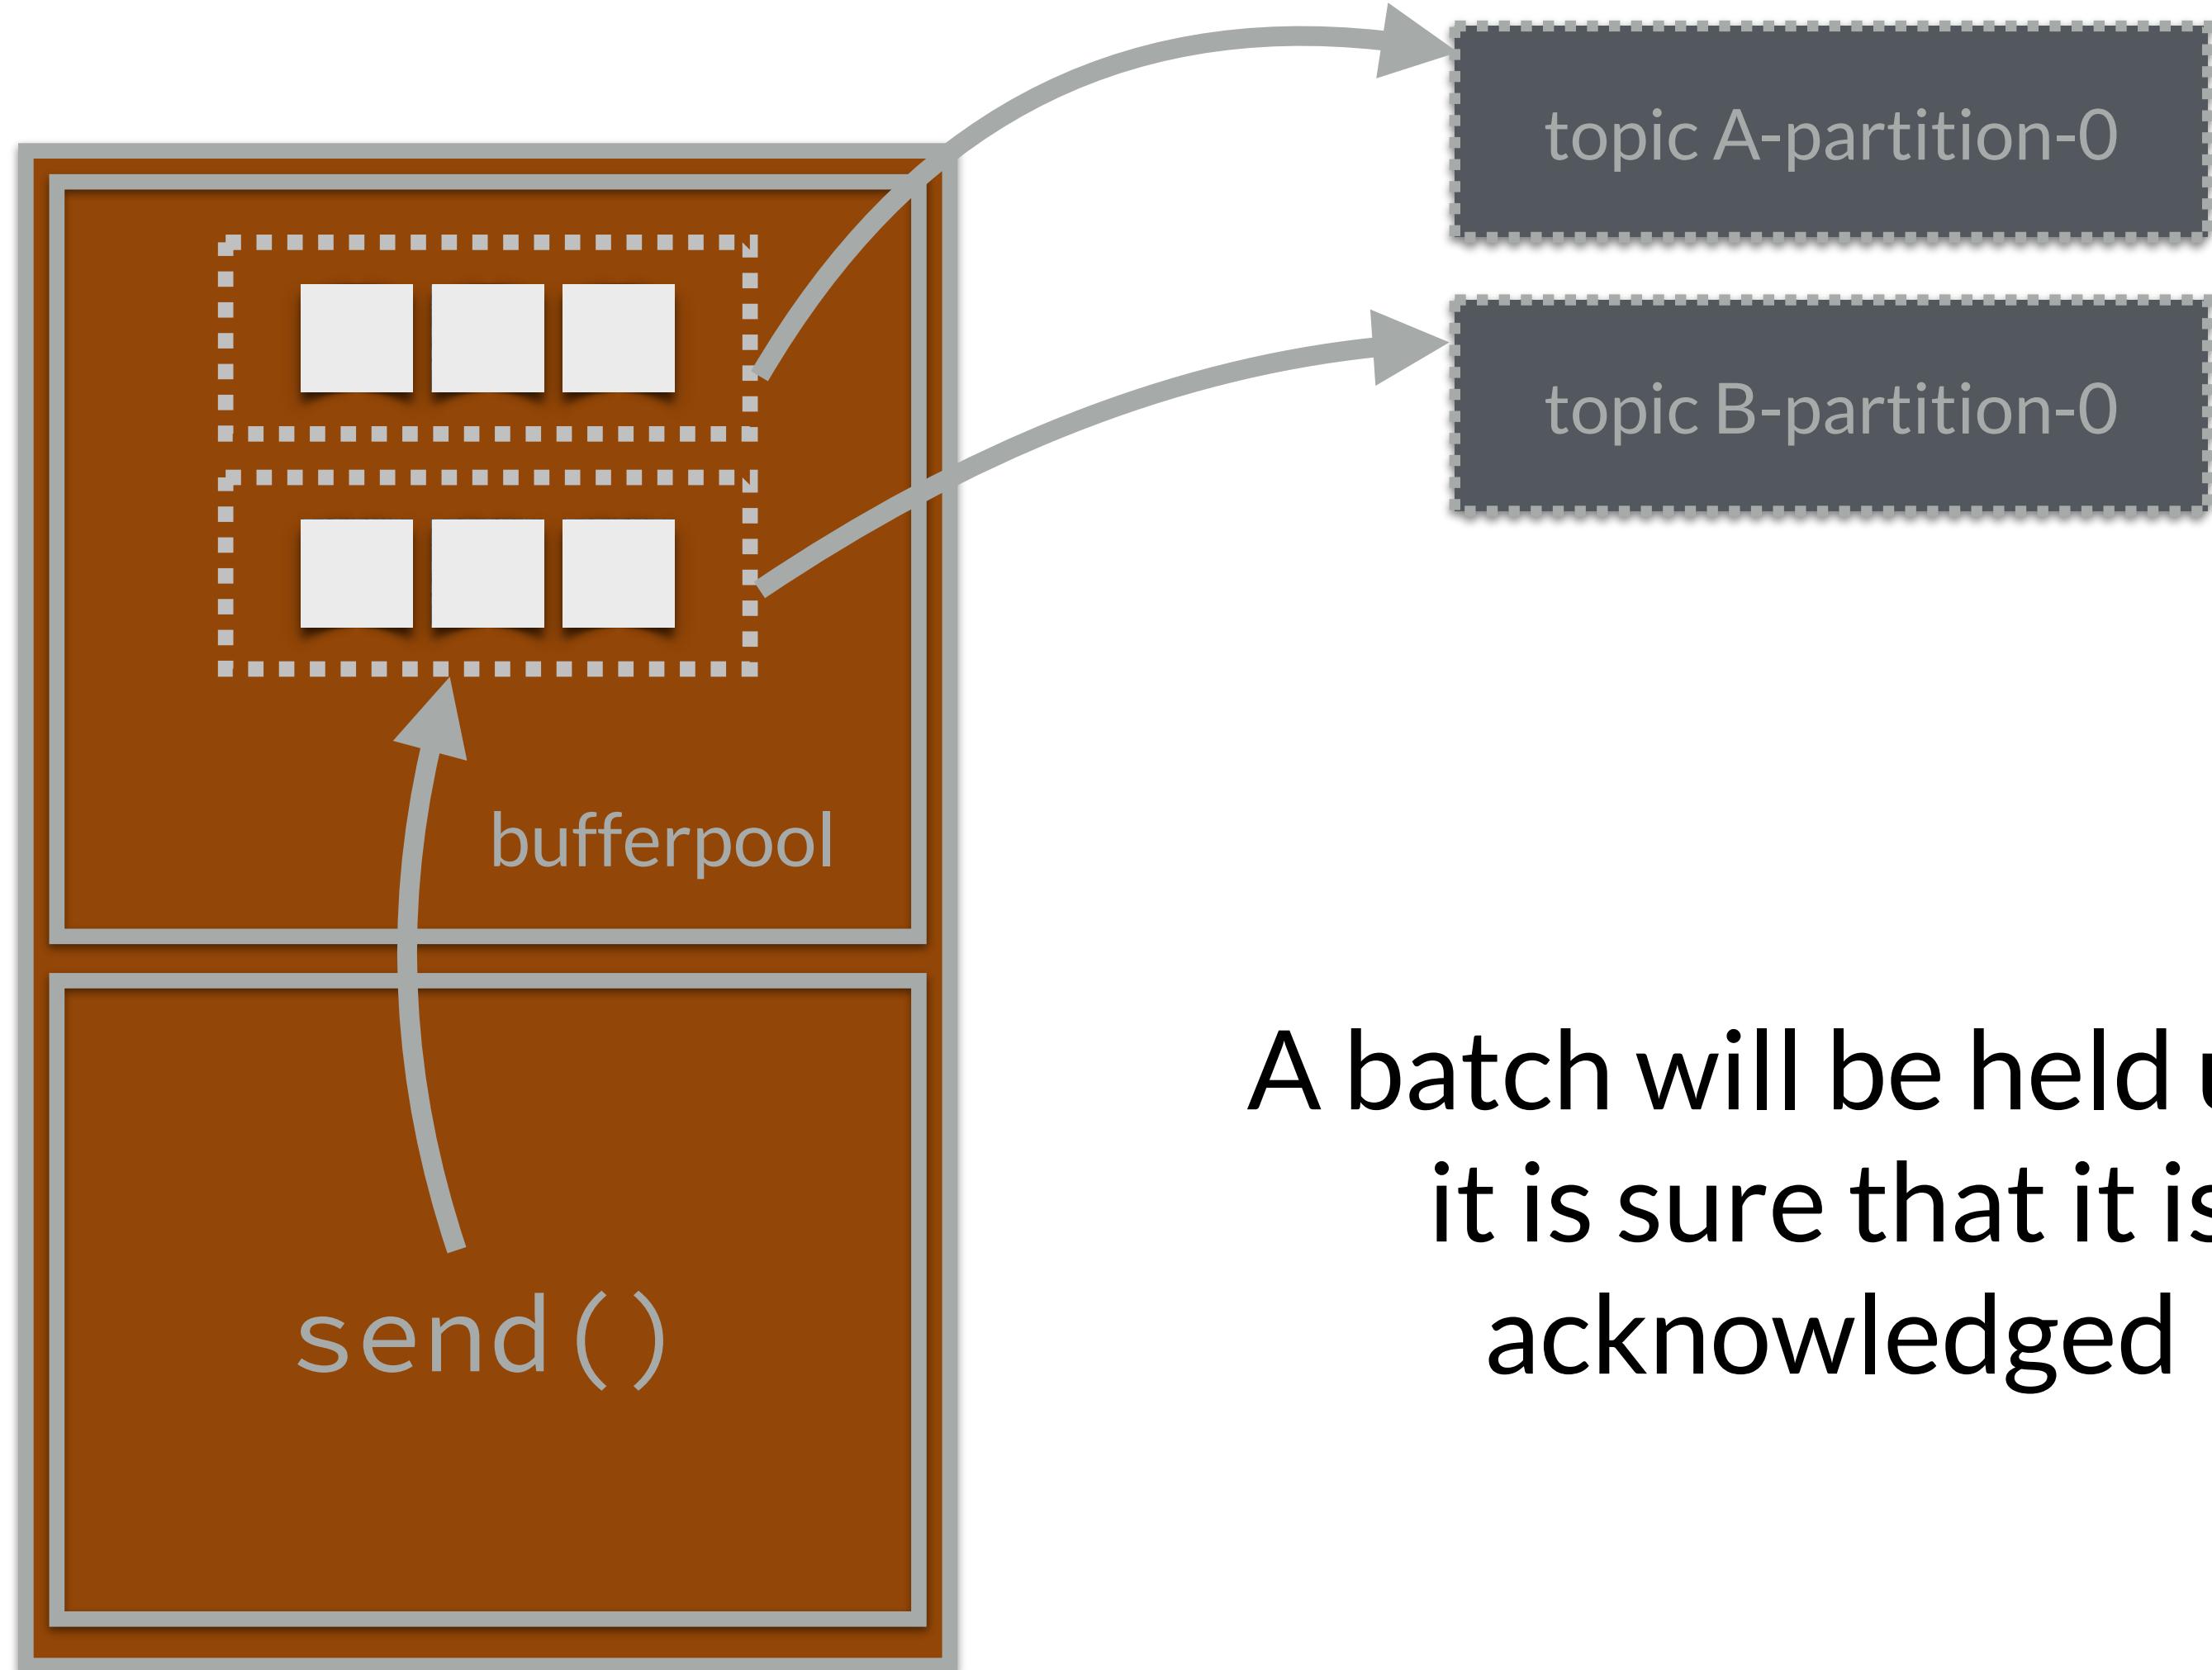
Acks

`acks` controls how many partition replicas must receive the record before the write is considered a success.

```
Properties properties = new Properties();
properties.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
    "localhost:9092");
properties.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
    StringSerializer.class);
properties.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
    IntegerSerializer.class);
properties.put(ProducerConfig.ACKS_CONFIG, "all");
properties.put(ProducerConfig.RETRIES_CONFIG, 20);
```

Acks

acks	description
0	No acknowledgment, assume all is well
1	The leader will respond with an acknowledgement
all	All replicas must acknowledge. Higher latency, safest




A batch will be held until
it is sure that it is
acknowledged

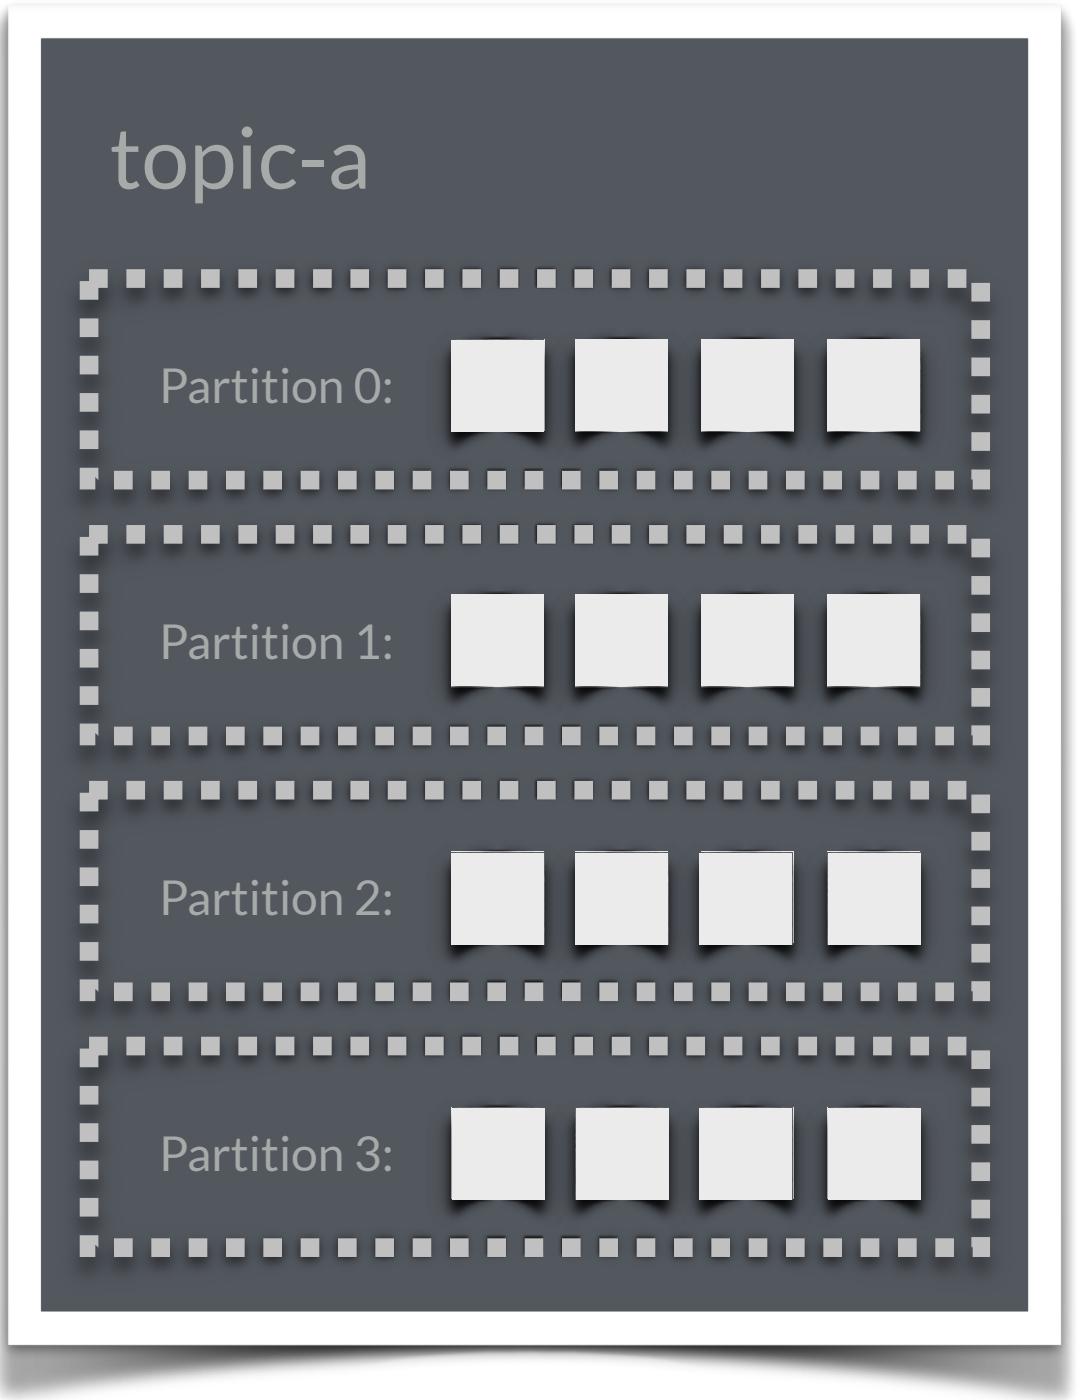
Lab: Ack Producers

Kafka Programming Consumers & Groups

Kafka Consumer Groups

- Consumers are typically done as a group
- A single consumer will end up inefficient with large amounts of data
- A consumer may never catch up
- **Every consumer should be on it's own machine, instance, pod**





group-id: marketing
client-id: 12942

group-id: marketing
client-id: 12941

group-id: marketing
client-id: 12940

group-id: marketing
client-id: 12939

Kafka's Goal

Kafka scales to large amount of different consumers without affecting performance

Kafka Consumer Threading

- There are no multiple consumers that belong to the same group in one thread.
One consumer per one thread.
- There are not multiple threads running one consumer, **One consumer per one thread.**

__consumer_offsets

- Topic on Kafka brokers that contain information about consumer and their offsets, stored in Kafka's data directory

Establishing Properties

- Construct a `java.util.Properties` object
- Provide two or more locations where the bootstrap servers are located
- Provide a "Team Name", officially called a `group.id`
- Provide a DeSerializer for the key
- Provide a DeSerializer for the value

```
Properties properties = new Properties();
properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG,
    "localhost:9092");
properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");
properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
    "org.apache.kafka.common.serialization.StringDeserializer");
properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
    "org.apache.kafka.common.serialization.IntegerDeserializer");
```

Create a Consumer Object

- Construct a `org.apache.kafka.clients.consumer.KafkaConsumer` object
- A Kafka Consumer is **not thread-safe**
- Inject the Properties into the Consumer object

```
KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```

Processing Messages

- Use the consumer that you have constructed, and call poll with pulse time
- The poll is a max sleep time, if a message is ready it will download a batch, ConsumerRecords
- Iterate through each record in the batch with a for-loop, process the message

```
while (!done.get()) {  
    ConsumerRecords<String, String> records =  
        consumer.poll(Duration.of(500, ChronoUnit.MILLISECONDS));  
    for (ConsumerRecord<String, String> record : records) {  
        System.out.format("offset: %d\n", record.offset());  
        System.out.format("partition: %d\n", record.partition());  
        System.out.format("timestamp: %d\n", record.timestamp());  
        System.out.format("timeStampType: %s\n", record.timestampType());  
        System.out.format("topic: %s\n", record.topic());  
        System.out.format("key: %s\n", record.key());  
        System.out.format("value: %s\n", record.value());  
    }  
}
```

Be a good citizen, close your resources

- When you need to terminate, let the *Kafka group coordinator* know you are done
- Close the Consumer

```
consumer.close();
```

Consumer Offset Reset

Consumer Offset Reset

- When a consumer is assigned it must determine what offset for the partition to start
- Position is set according to your consumer's reset policy, either earliest or latest
- This is only applicable if there is no valid offset in the `__consumer_offset__` topic

```
Properties properties = new Properties();
...
properties.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```

Manual Commits

Manual Commits

- Consumers will commit automatically at minimum 5 seconds
- If your consumer fails within those 5 seconds another consumer may get those messages
- Might be too long before commit, in such case you may want to handle commits yourself
- Turn off your auto-commit by setting `ENABLE_AUTO_COMMIT_CONFIG` to false

```
Properties properties = new Properties();
...
properties.put(ConsumerConfig.ENABLE_AUTO_COMMIT_CONFIG, false);
KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```

Synchronous Commit

- After turning off auto-commit
- Committing and blocking synchronously yourself
- Synchronous commits will block

Synchronous Commit

- After turning off auto-commit
- Committing and blocking synchronously yourself
- Synchronous commits will block

```
KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);  
...  
consumer.commitSync(); //Block
```

Asynchronous Commit

- After turning off auto-commit
- Committing and letting it commit asynchronously
- Advantage: Speed, Higher throughput
- Disadvantage: You will get error reposts asynchronously

```
consumer.commitAsync((offsets, exception) -> {  
    //offsets  
    //exception  
});
```

Using both kinds of commits

- You can commit asynchronously within the loop
- Then commit synchronously outside the loop
- Reason, if you are out of the loop, likely you want to shut down

```
Runtime.getRuntime().addShutdownHook(new Thread(() -> done.set(true)));

while (!done.get()) {
    ConsumerRecords<String, String> records = consumer.poll(Duration.of(500, ChronoUnit.MILLIS));
    for (ConsumerRecord<String, String> record : records) {
        ...
    }
    consumer.commitAsync((offsets, exception) -> {
        //offsets
        //exception
    });
}
consumer.commitSync(); //Block
consumer.close();
```

Consumer Rebalancing

Consumer Rebalance


- When one partition is moved from one consumer to another, this is known as a *rebalance*.
- A way to mitigate when either consumers go down, or when consumers are added.
- Although unavoidable, this will cause an unfortunate pause, and it ***will lose state***.

Groups and Heartbeats

- At regular intervals, consumers will send heartbeats to the broker, to let it know it is alive and still reading data
- Heartbeats are sent to a Kafka broker called the *group coordinator using out-of-band thread*
- consumer polls and consumes a record, where commit the offset

Leaving the Group

- When a consumer leaves the group, it lets the *group coordinator* know it is done
- The Kafka broker then triggers a rebalance of the group.



Consumer Rebalance Listener

- You can implement a ConsumerRebalanceListener and be alerted when a rebalance has occurred
- The rebalance listener requires that you implement onPartitionsRevoked, and onPartitionsAssigned

```
consumer.subscribe(Collections.singletonList("my_orders"),
    new ConsumerRebalanceListener() {
        @Override
        public void onPartitionsRevoked(Collection<TopicPartition> collection) {
            System.out.println("Partition revoked:" +
                collectionTopicPartitionToString(collection));
            consumer.commitAsync();
        }

        @Override
        public void onPartitionsAssigned(Collection<TopicPartition> collection) {
            System.out.println("Partition assigned:" +
                collectionTopicPartitionToString(collection));
        }
    });
}
```

Lab: Consumer

Retention Settings

Retention

- Durable Storage over a Period of Time
- Can either be configured in time or size
- Once reached messages are expunged

Retention Formula

Who administers Kafka can update the default retention time or size

```
log.retention.hours = 120  
log.retention.bytes = 3221225472
```

Default Hours are set to (168 hours) or one week

You may also use **log.retention.minutes** or **log.retention.ms**, whichever is smallest wins.

You can also configure retention per topic, like the following

```
bin/kafka-configs --bootstrap-server localhost:2181 --entity-type topics --entity-name my-topic --alter --add-config retention.ms=128000
```


Compaction

Compaction

- A form of retention where messages of the same key where only the latest message will be retained.
- Compaction is performed by a *cleaner thread*

kafka broker: 0

Partition 0:




[0] [1] [2] [3] [4] [5]



kafka broker: 0


Partition 0:



[3] [5]

kafka broker: 0

Partition 0:



clean

dirty

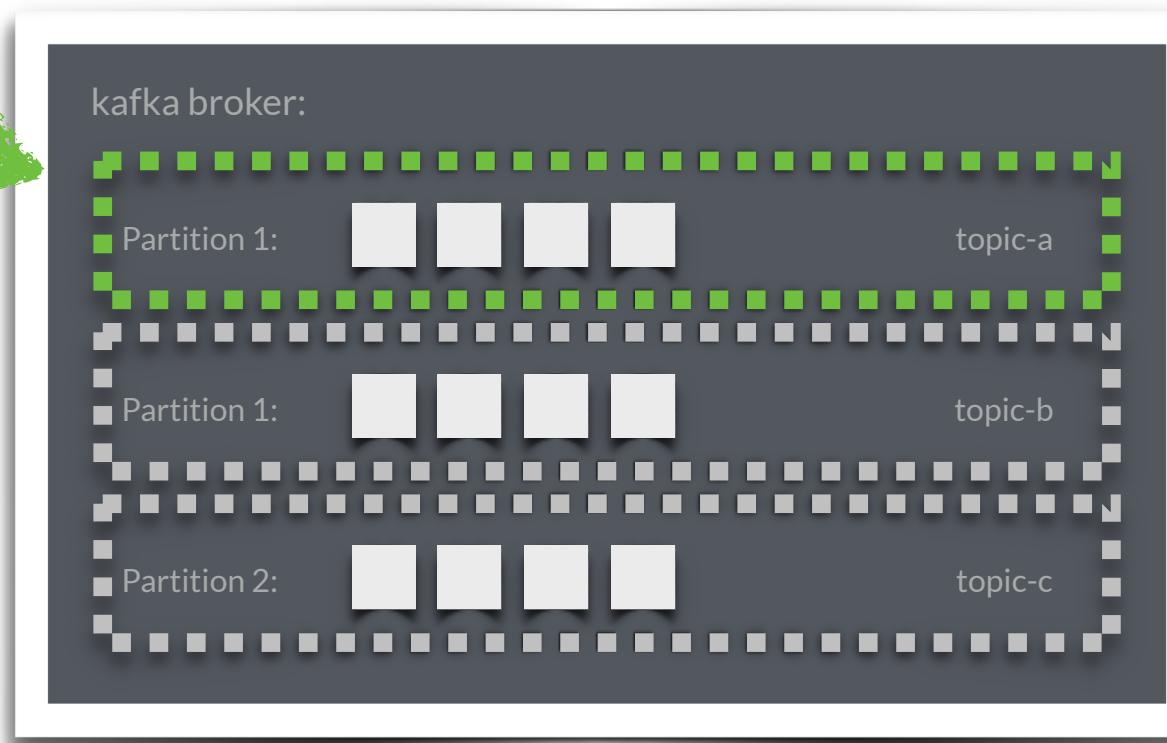

Kafka will start compacting when 50% of the topic contains dirty records

Resiliency


ISR (In sync replicas)

- Given a leader partition, an in sync replica is one that has been kept up to date within the last 10 seconds
- This is configurable
- During a crash, the closes ISR will retain control for failover


Leader




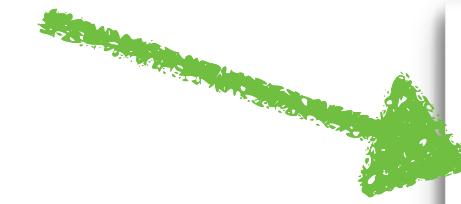
ISR




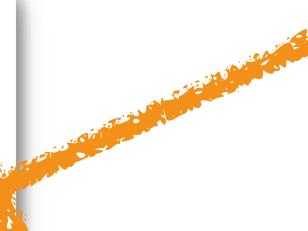
ISR




Leader



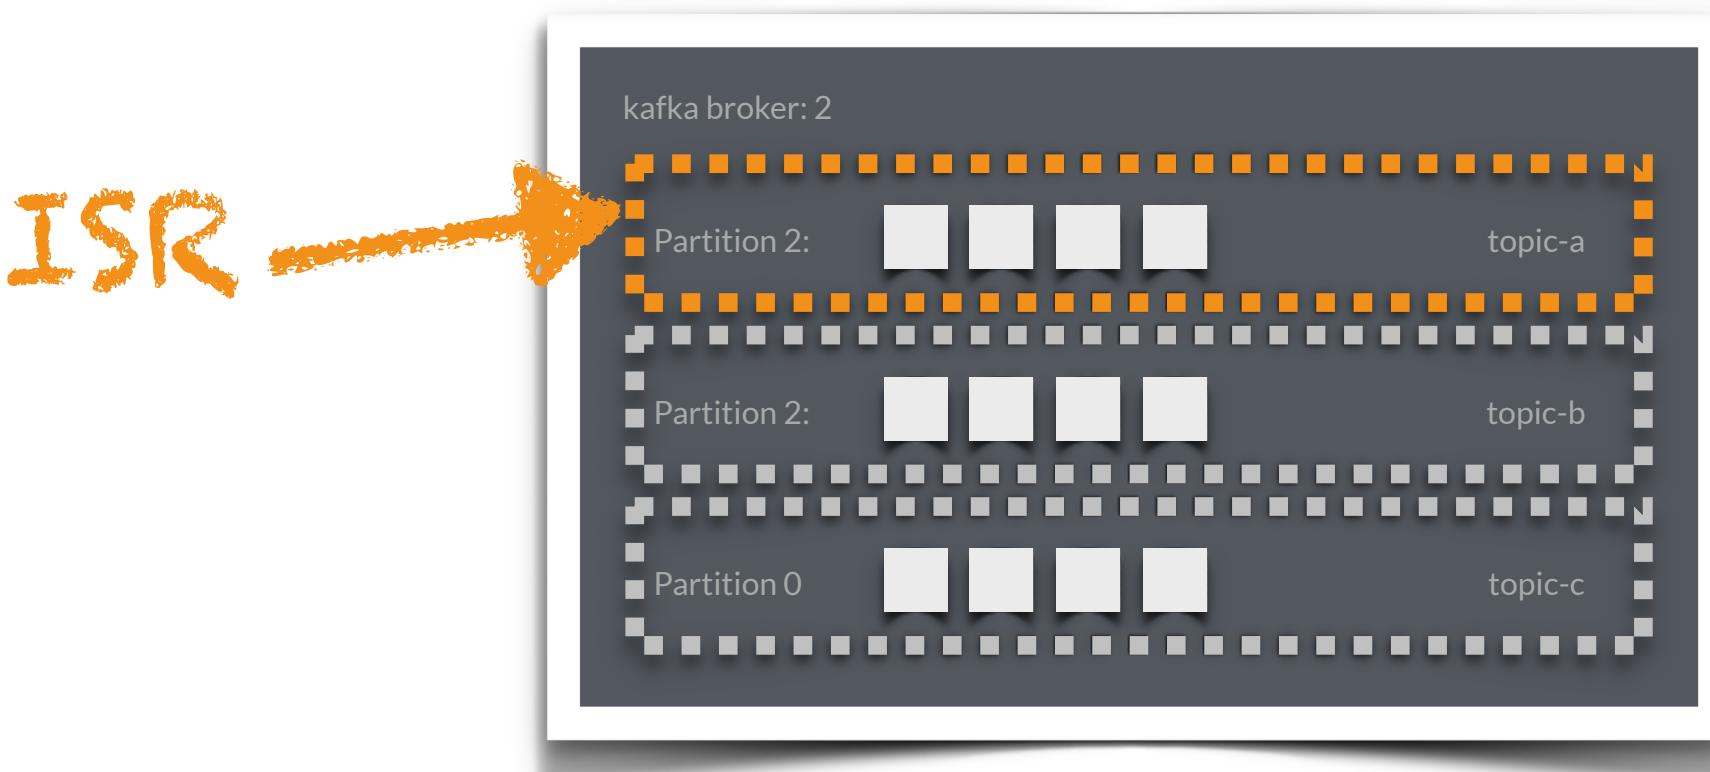

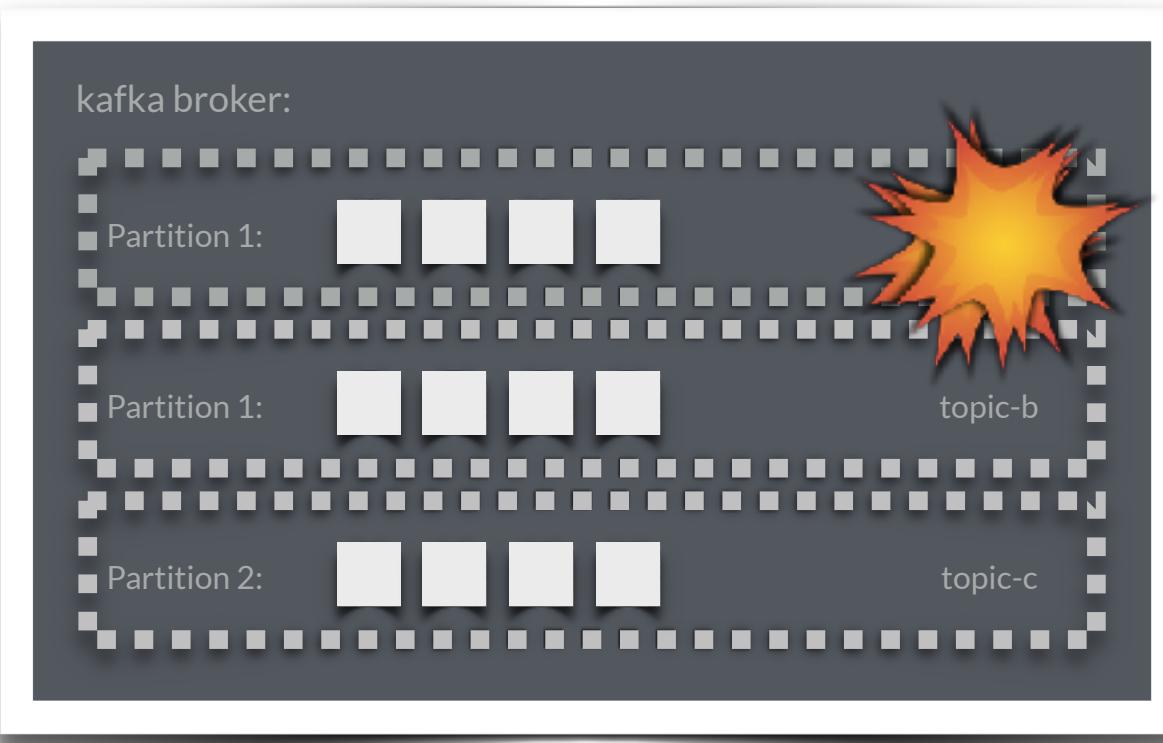
ISR



ISR



Leader



Streaming Concepts



Stream Processing

Endless supply of data

“Replayable”

Real Time Processing for Fraud Detection, High End Sales Detection, Internet of Things, and More.


Will require built-in accumulator table to perform real time data processing, like counting, and grouping

NY

100.00

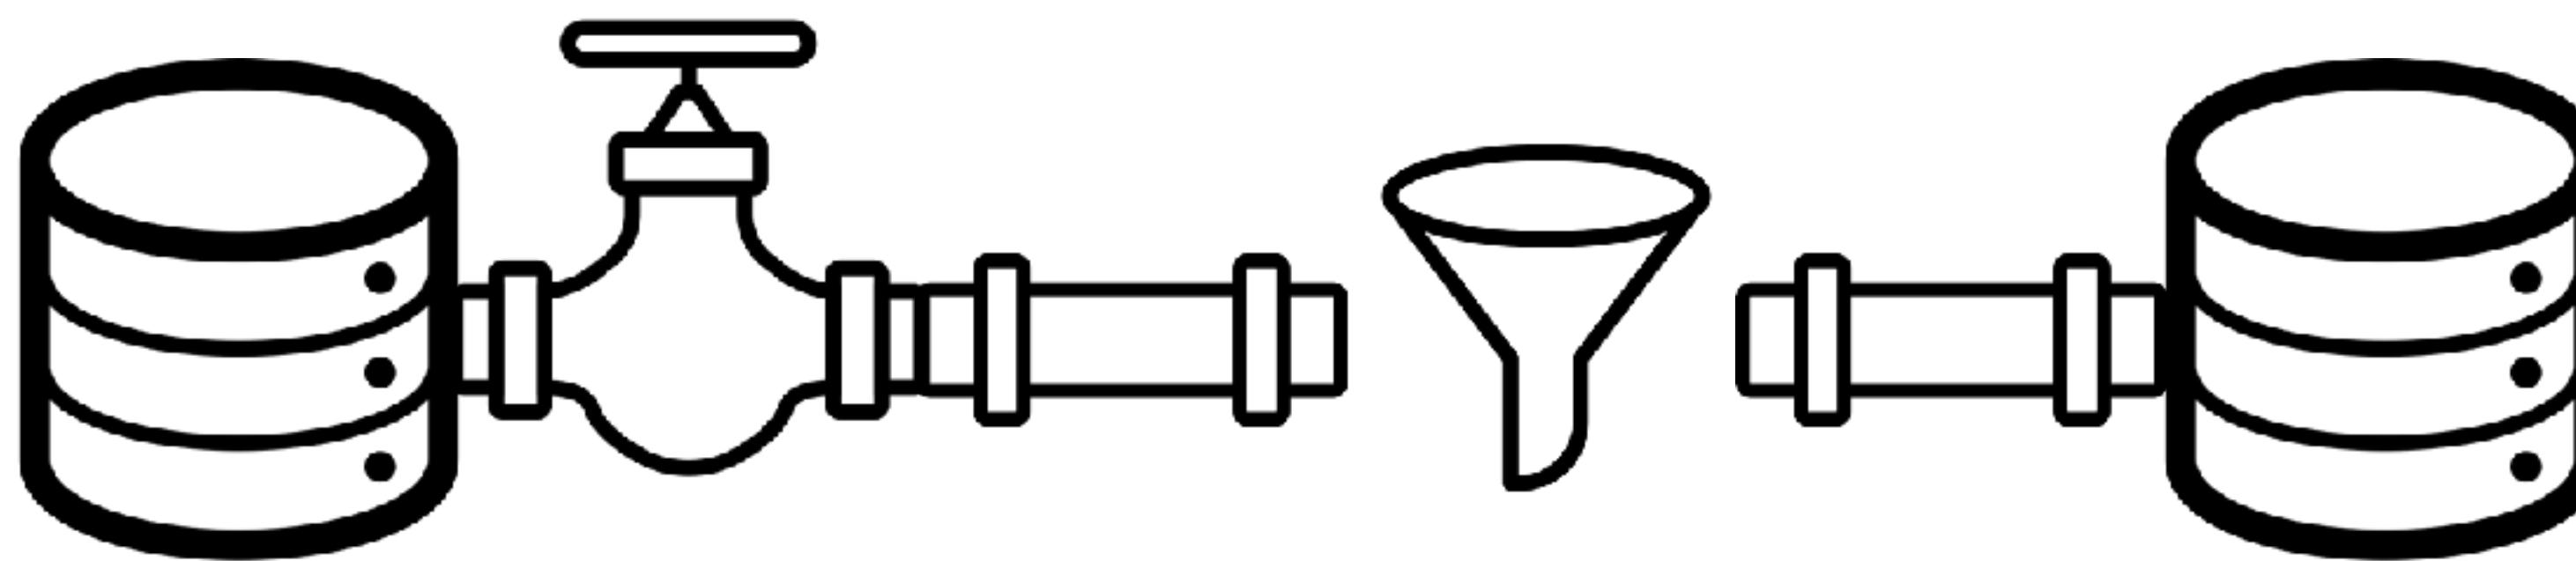
Filtering

NY
.....
100.00



value > 10000

OH
.....
20000

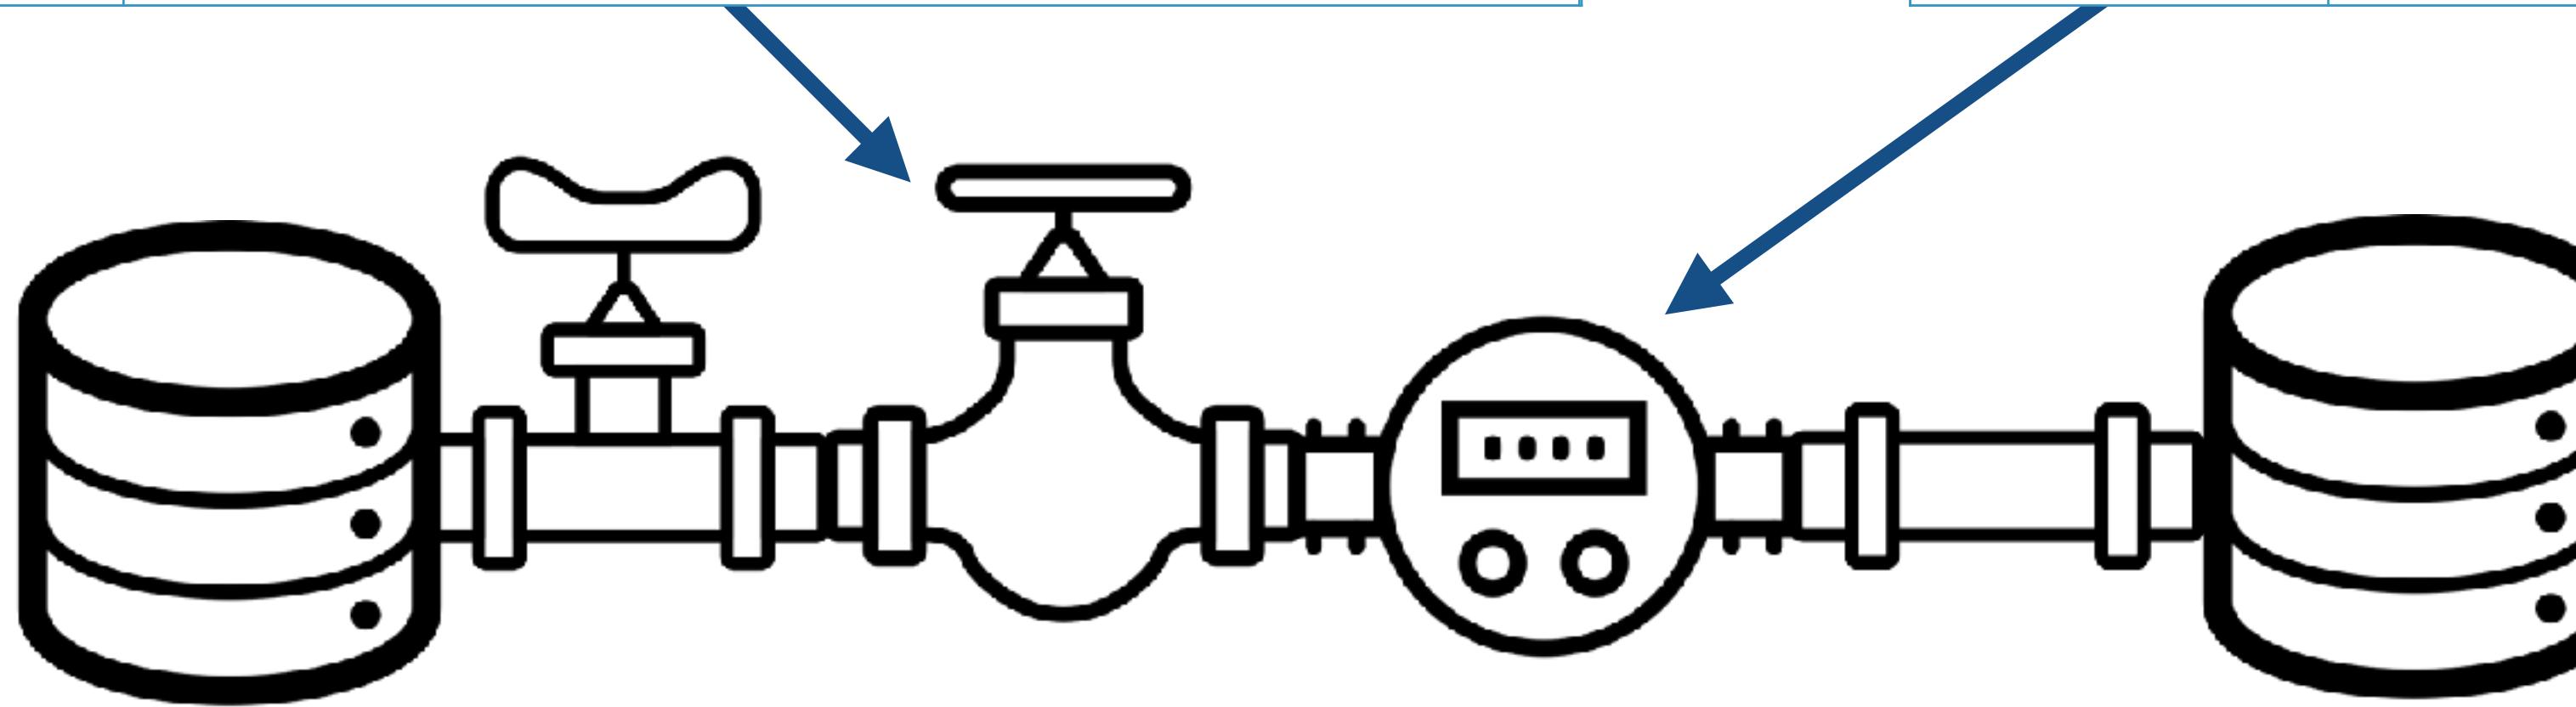


value > 10000

Aggregating

Key	Value
OH	100.0

Key	Value
OH	100.0

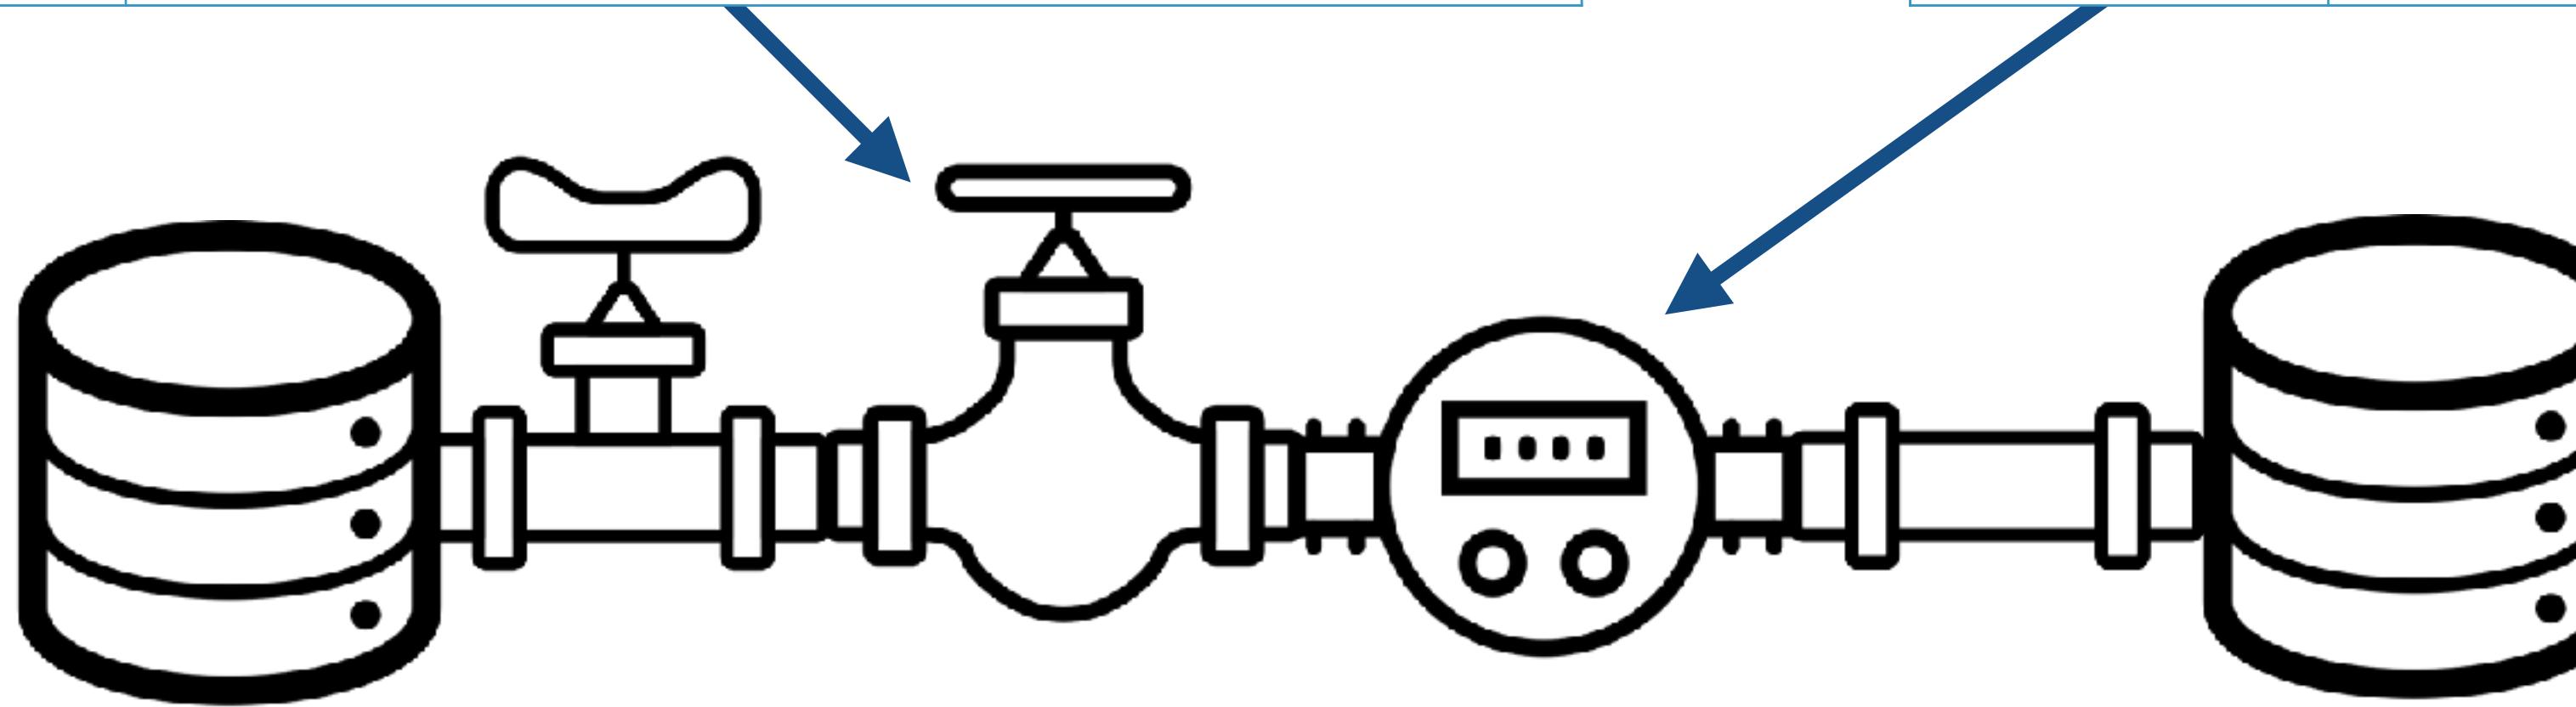


OH
.....
100.00

OH
.....
100.00

Key	Value
OH	100.0
NY	20.0

Key	Value
OH	100.0
NY	20.0

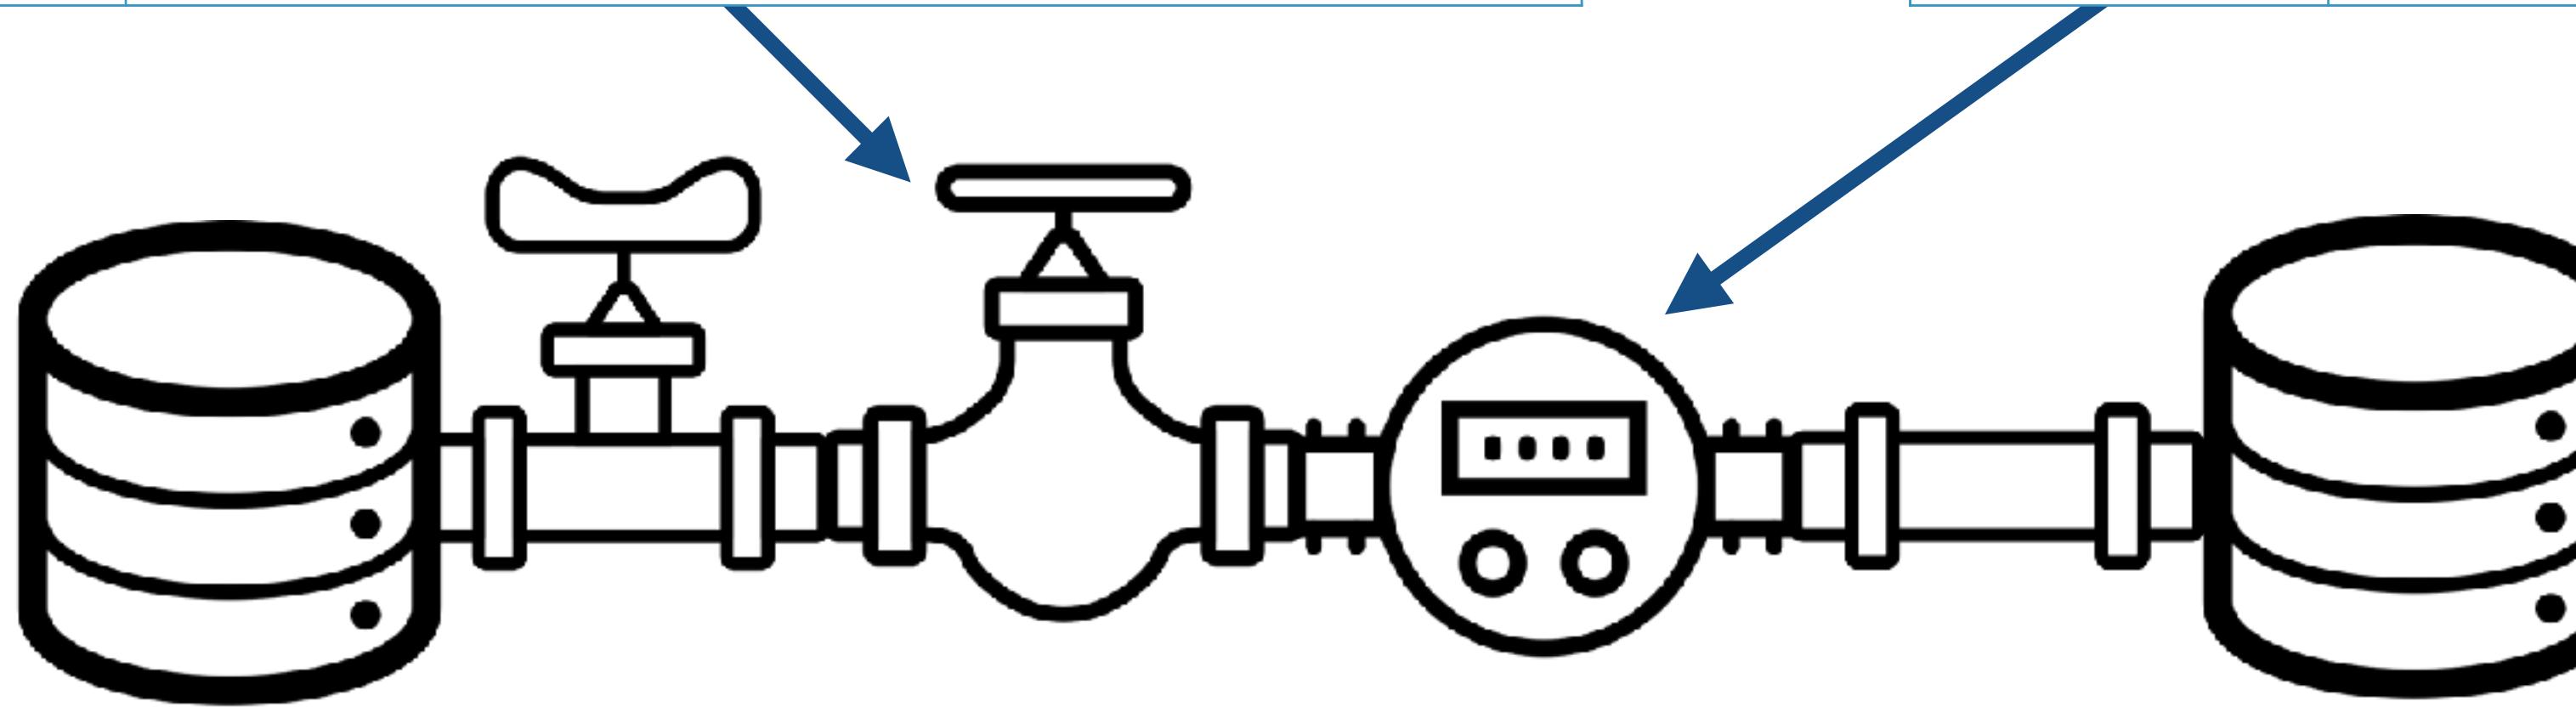


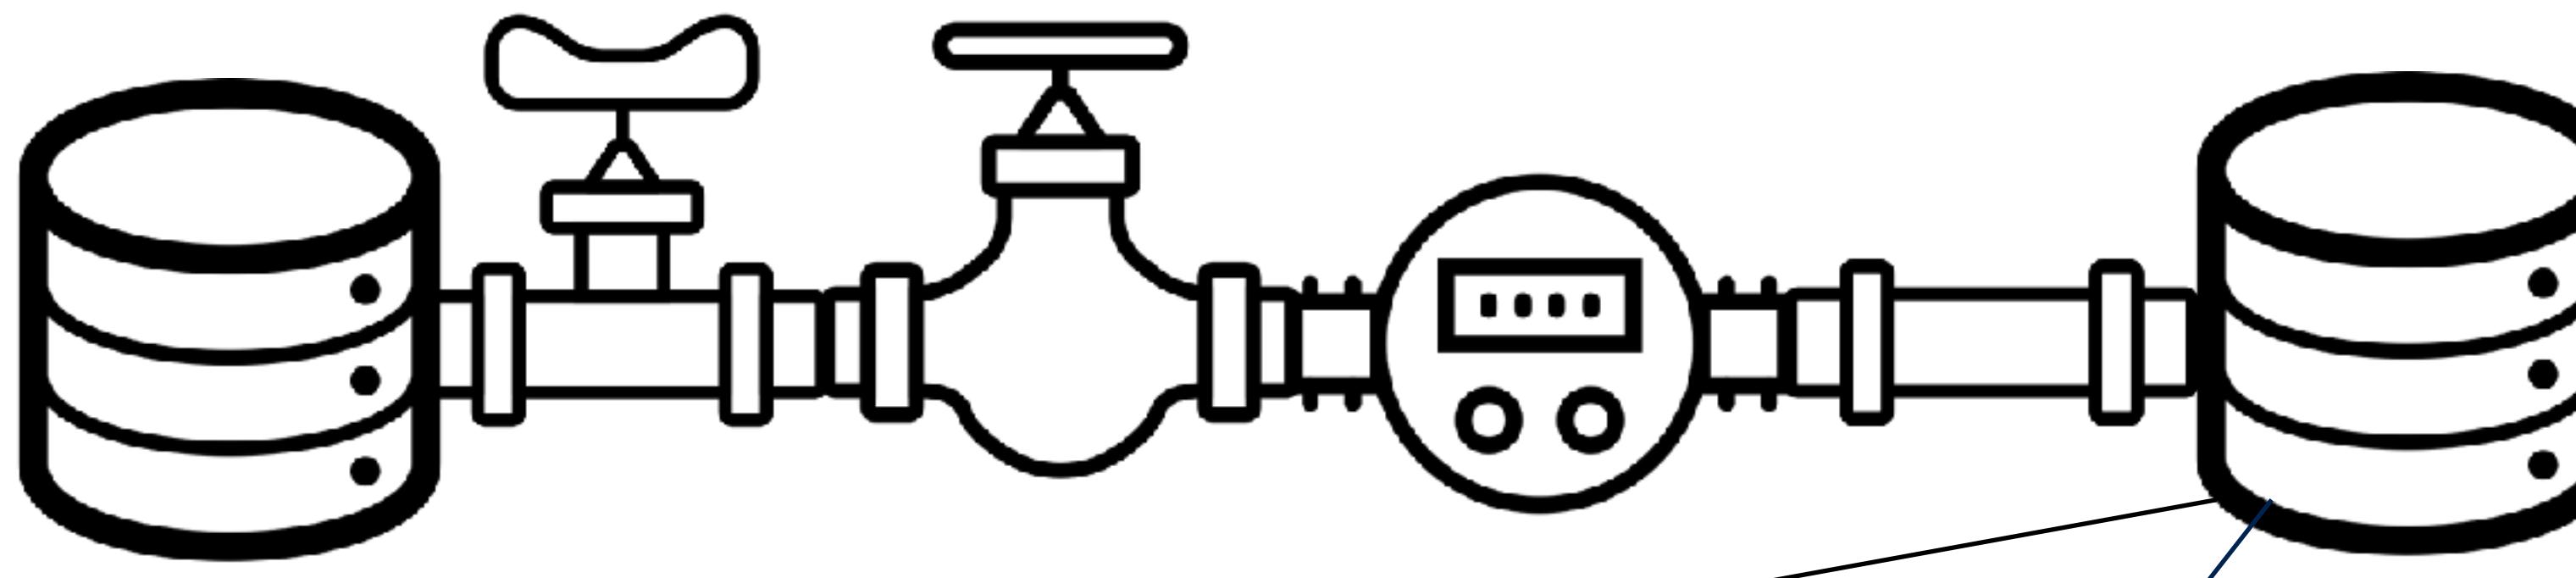
NY
20

NY
20

Key	Value
OH	100.0
NY	20.0, 60.0

Key	Value
OH	100.0
NY	80.0



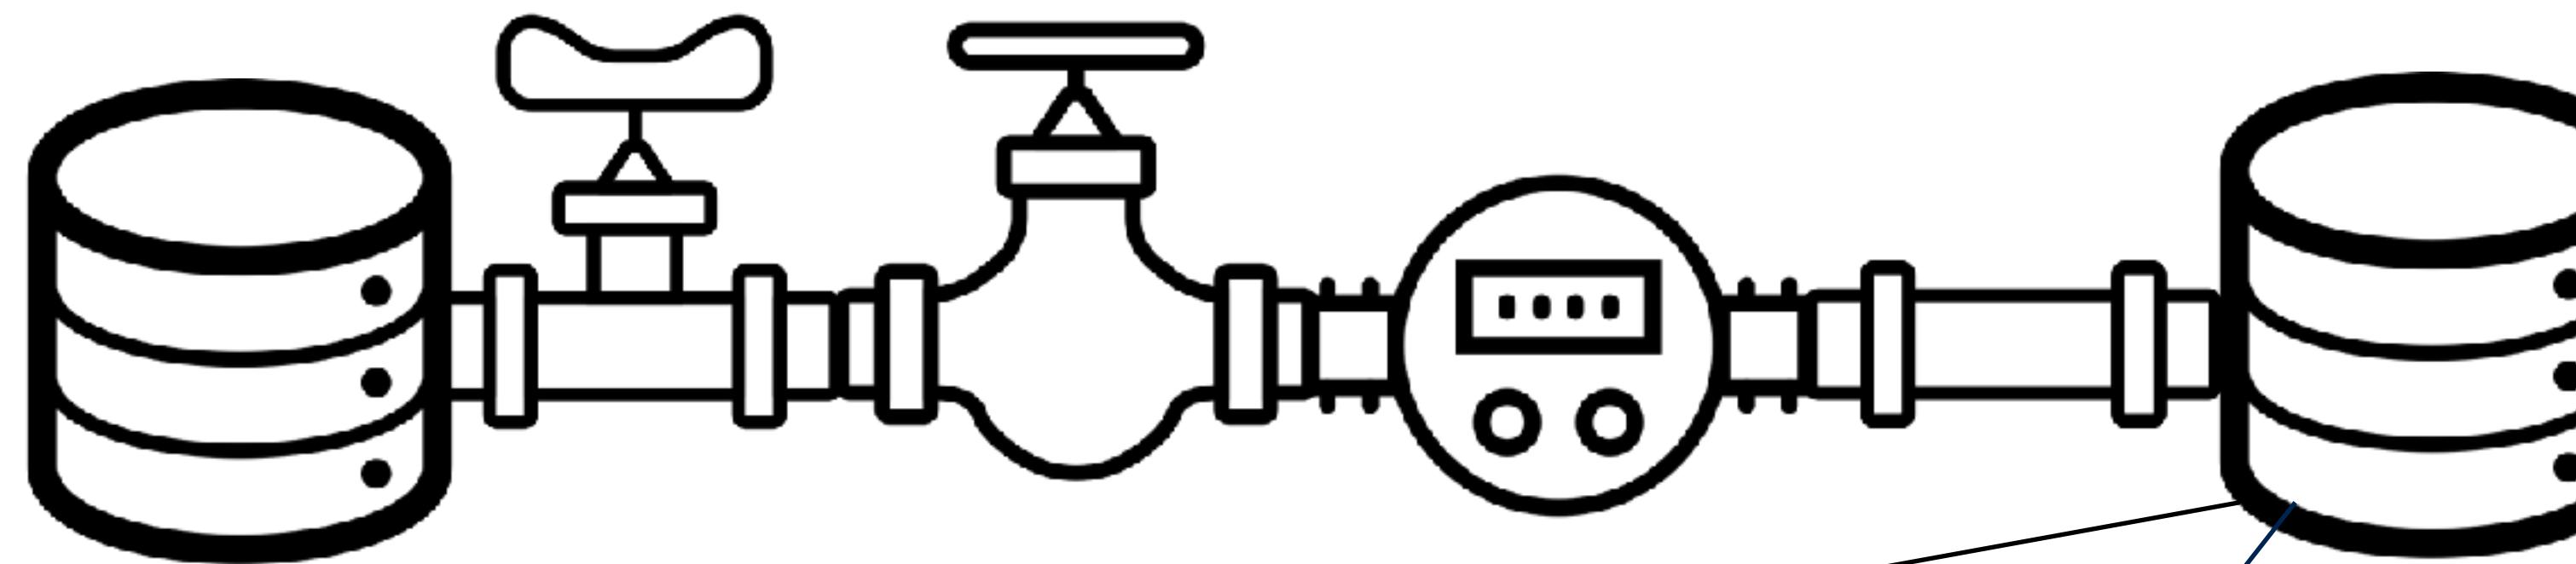


Offset	Key	Value
0	OH	100.0

Partition 0

Offset	Key	Value
0	NY	20.0
1	NY	80.0

Partition 1



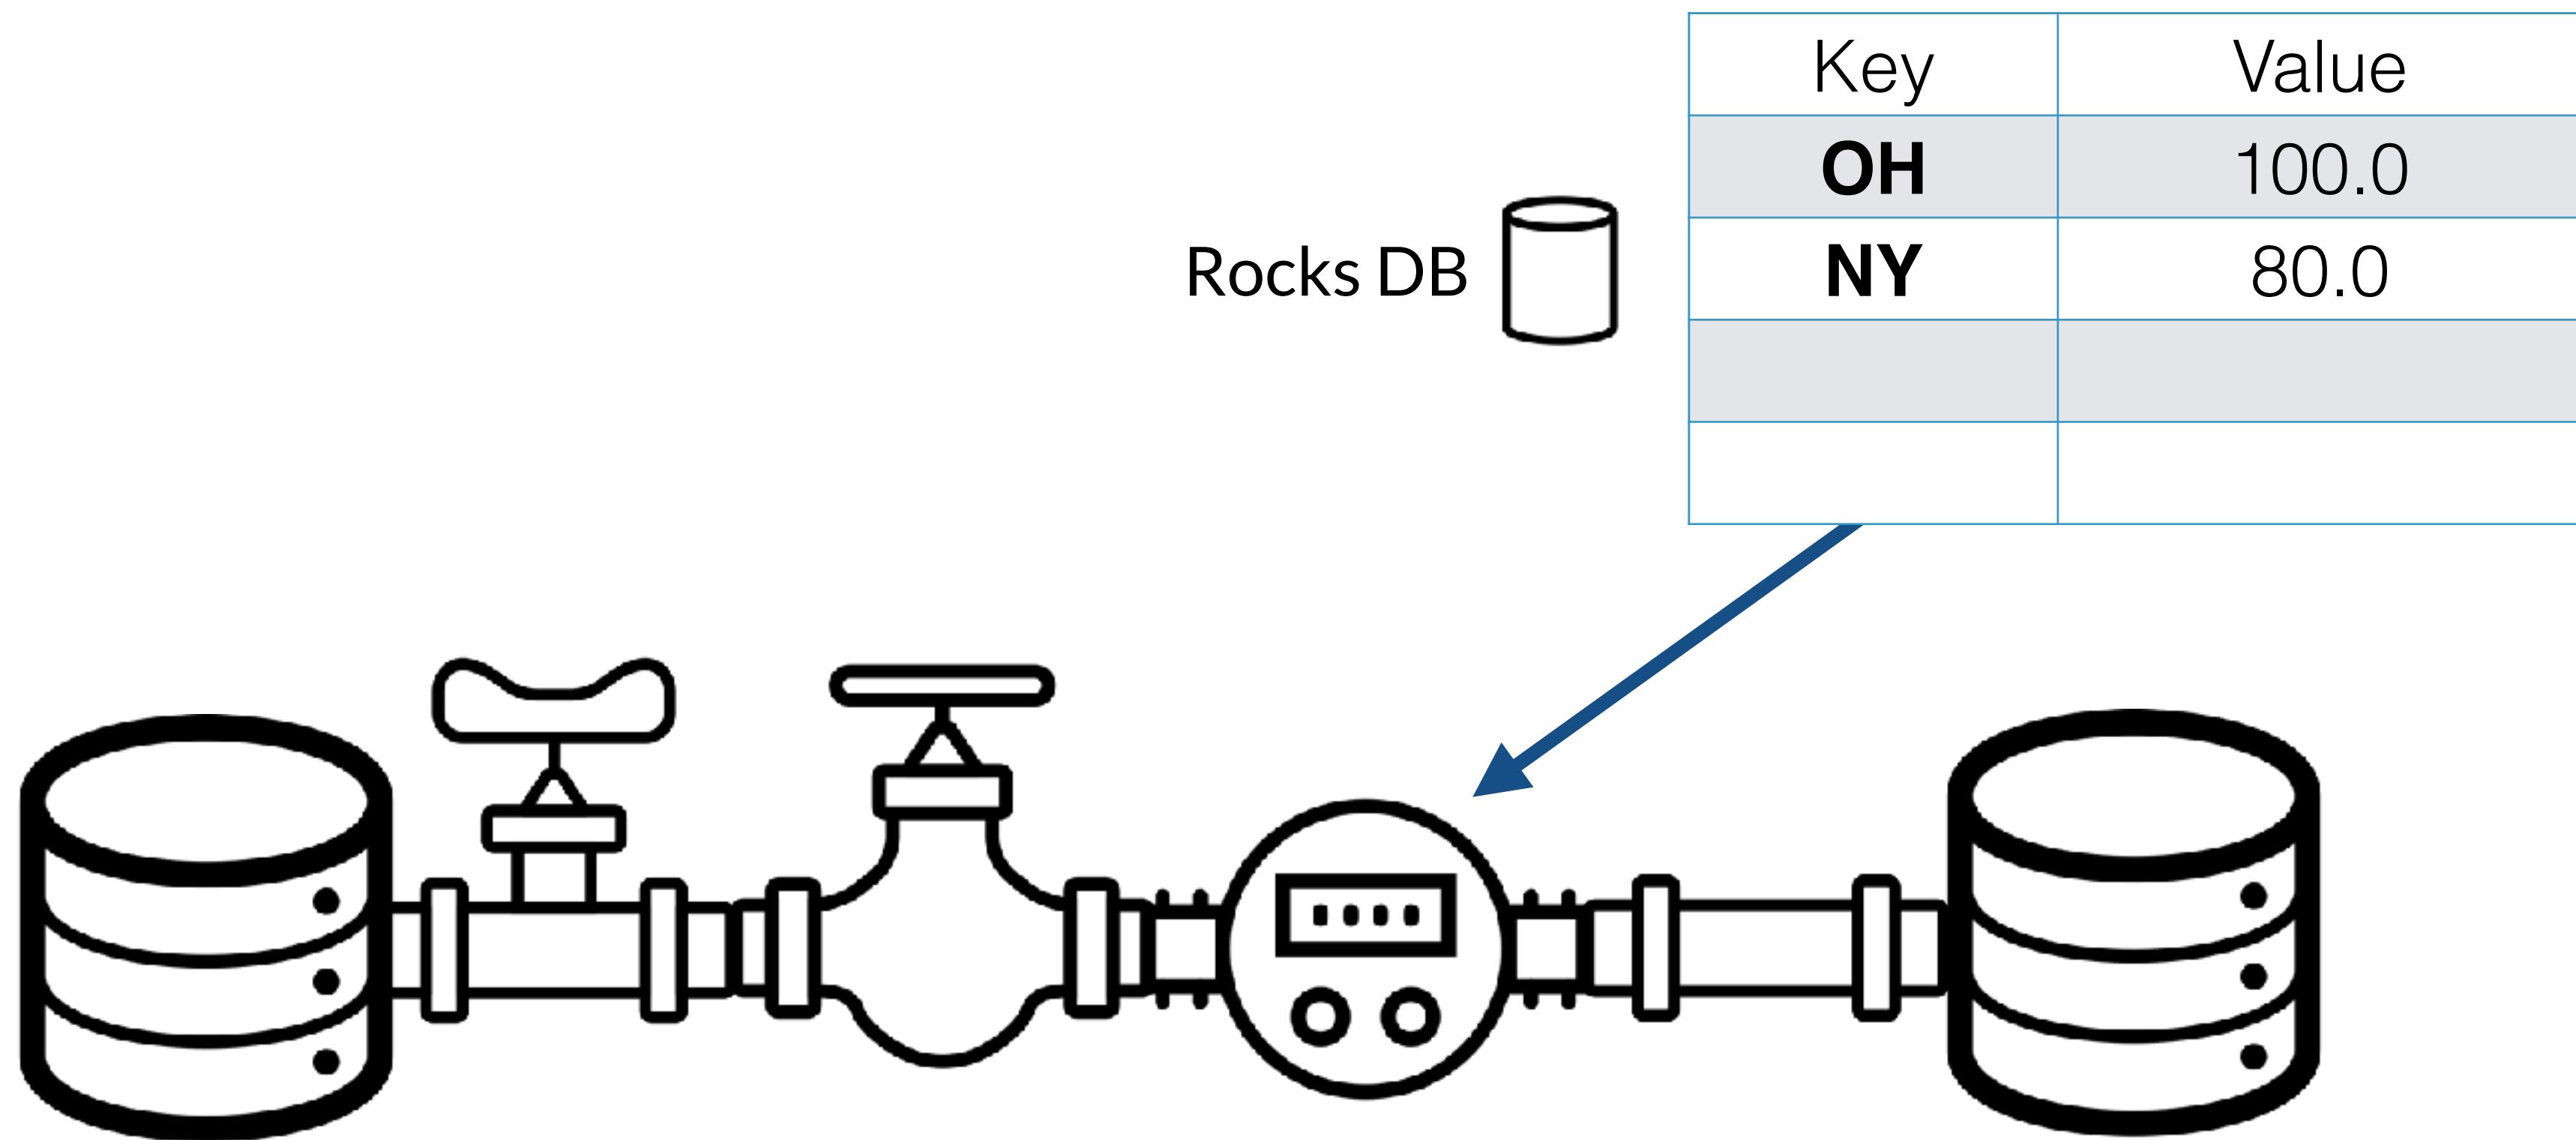
Offset	Key	Value
0	OH	100.0

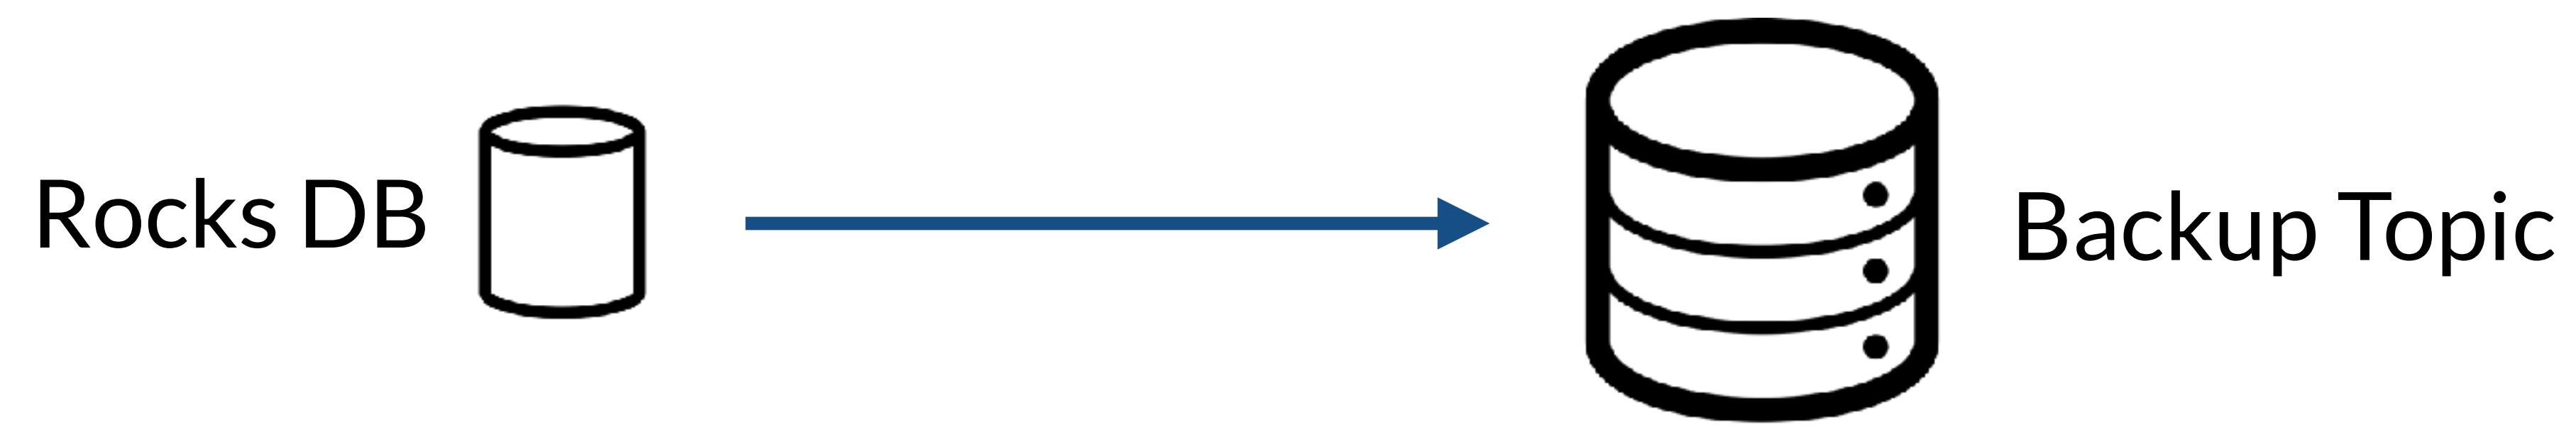
Partition 0


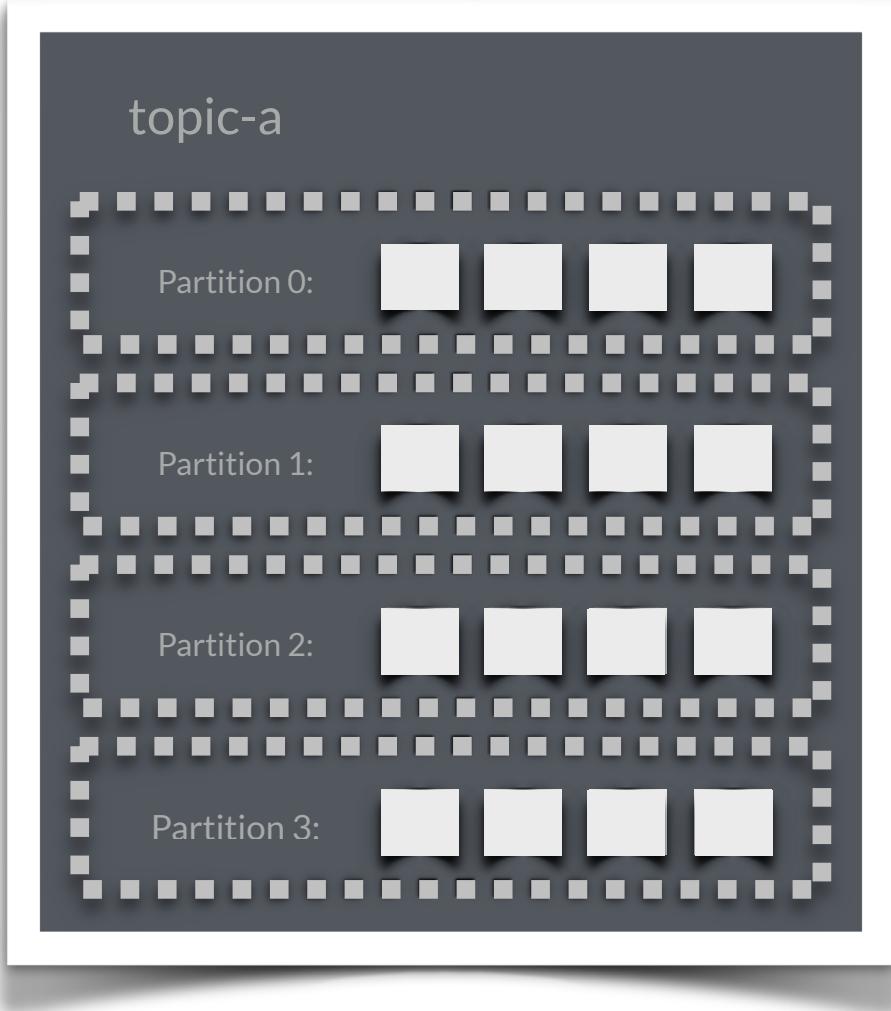
Offset	Key	Value
1	NY	80.0

Partition 1

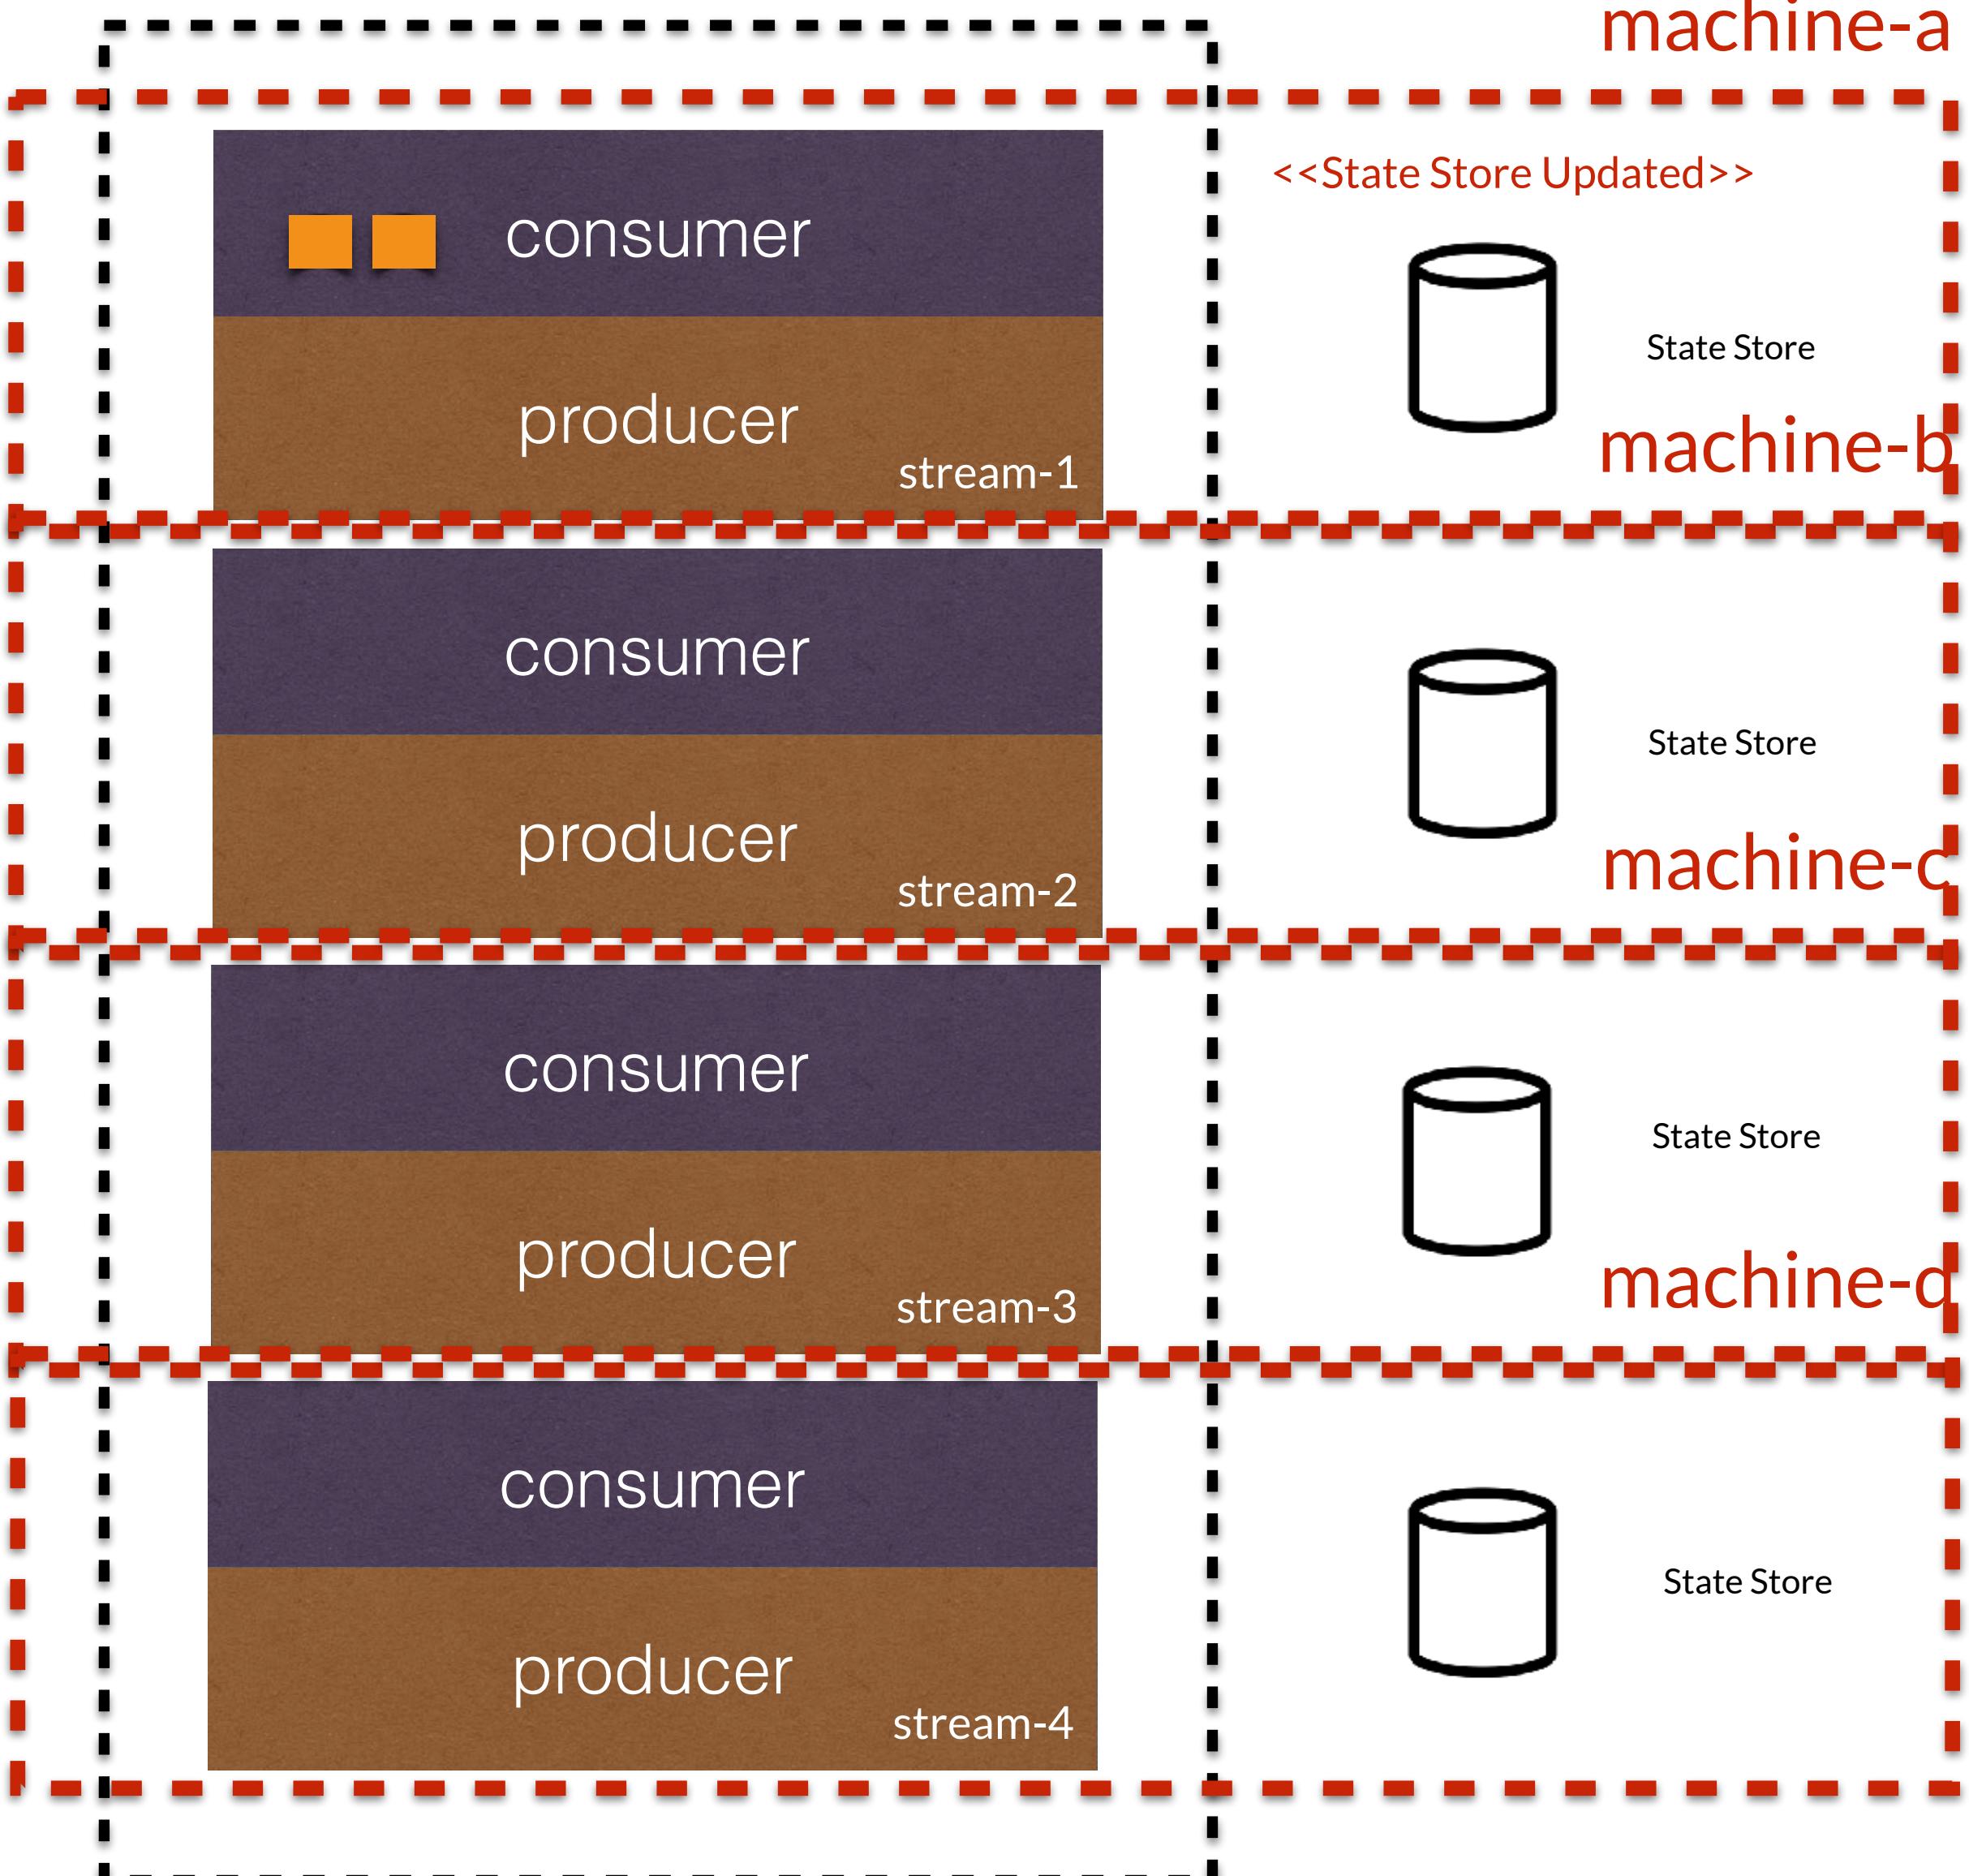
After Compaction




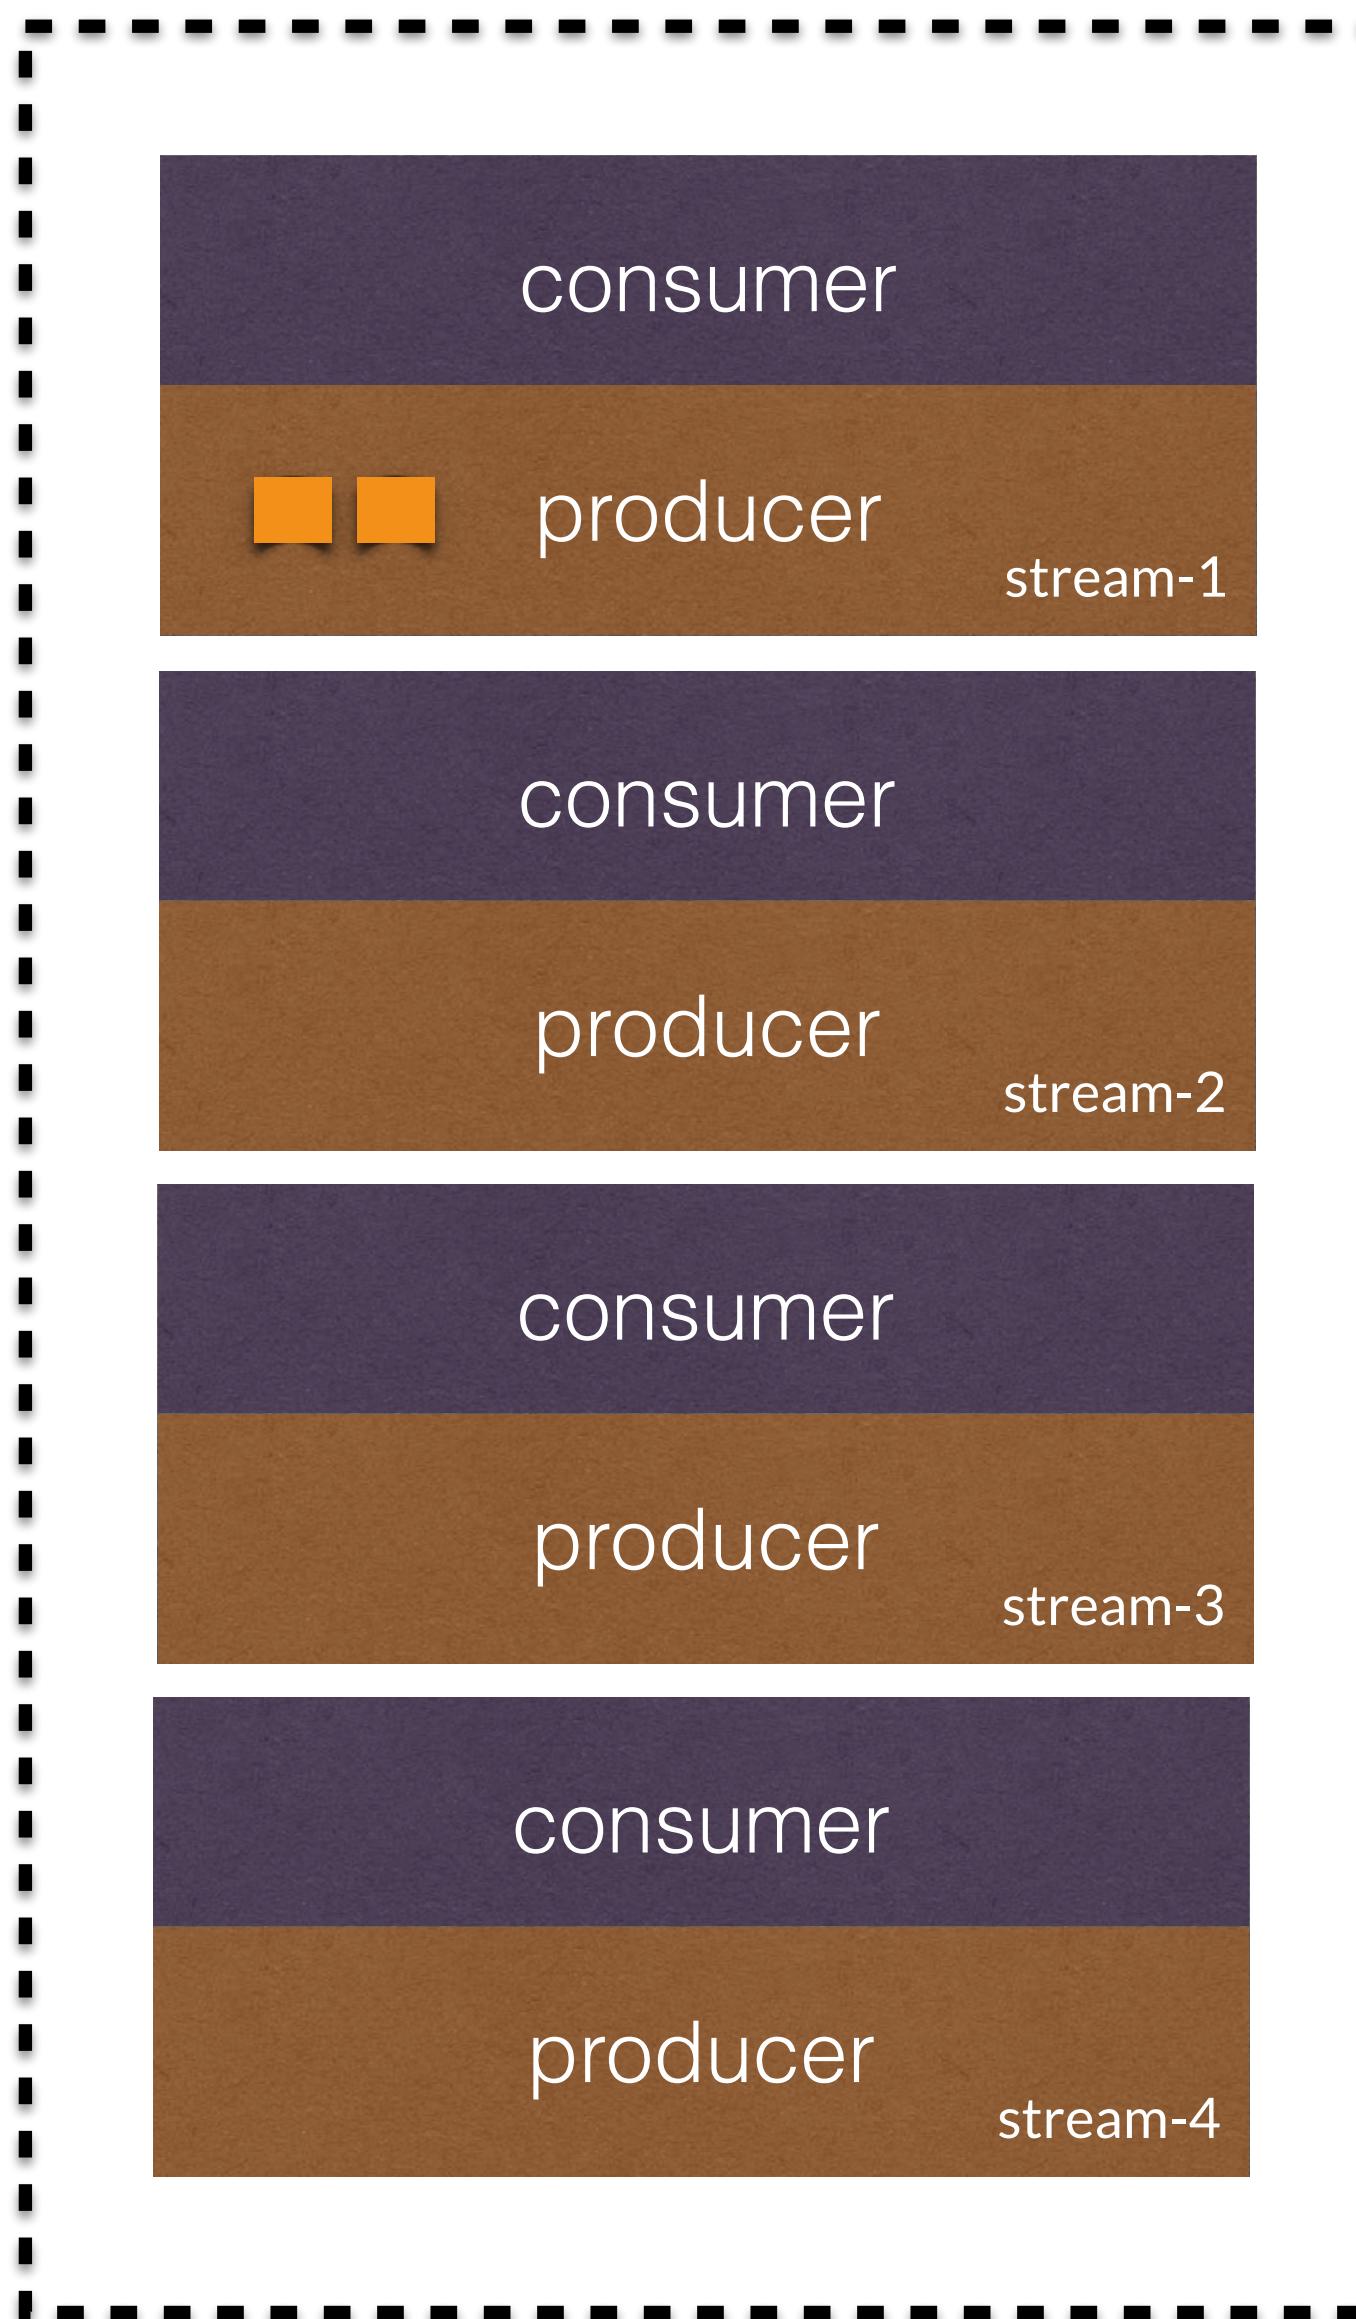
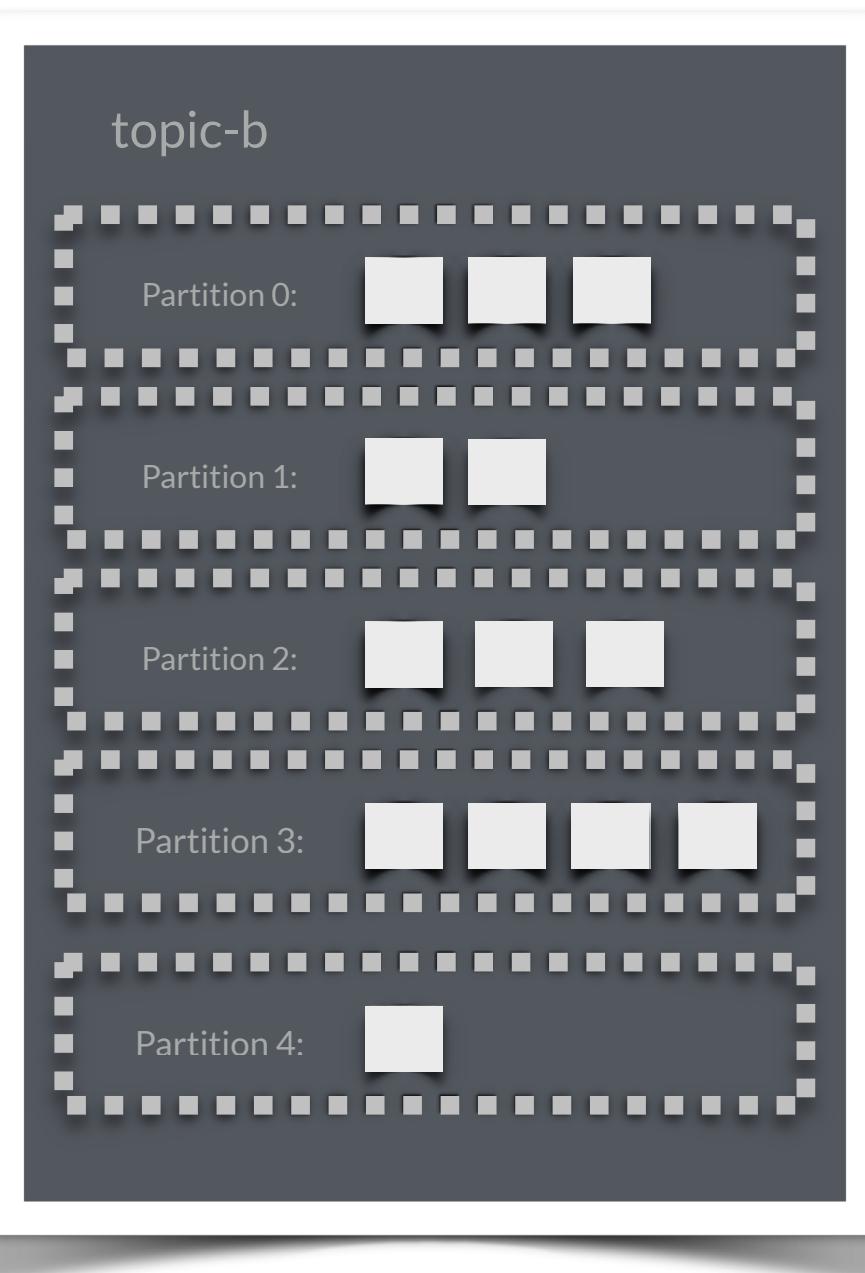
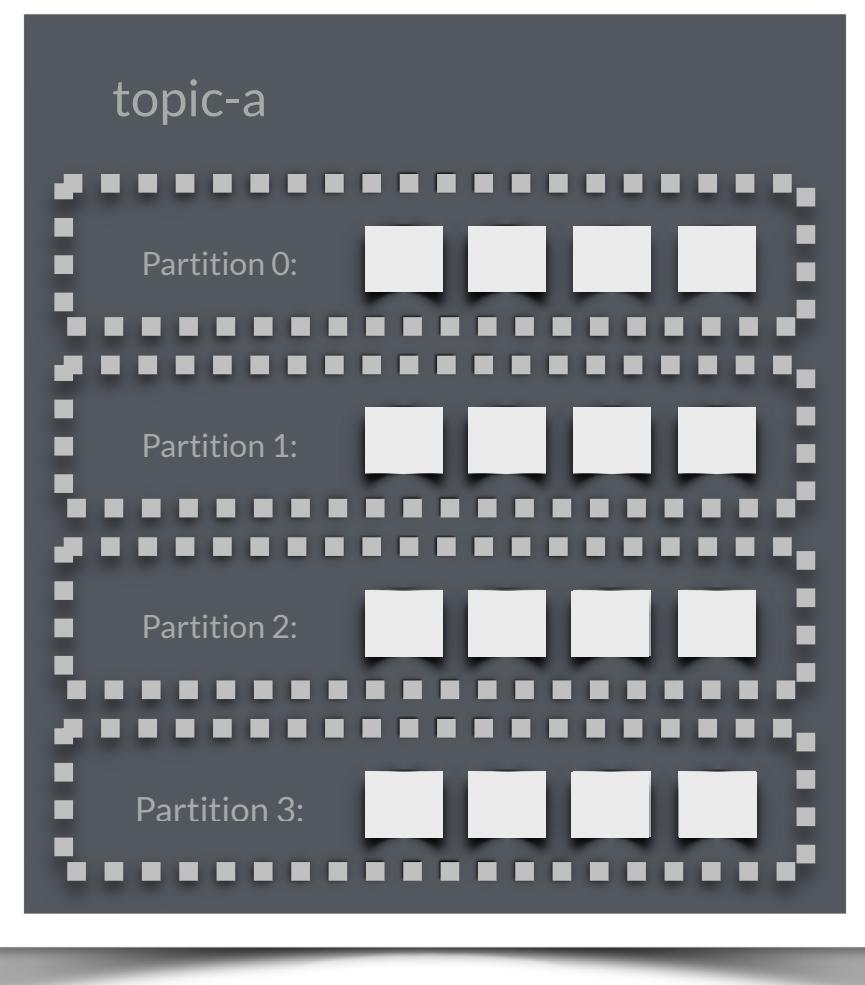




application.id = “my-stream”

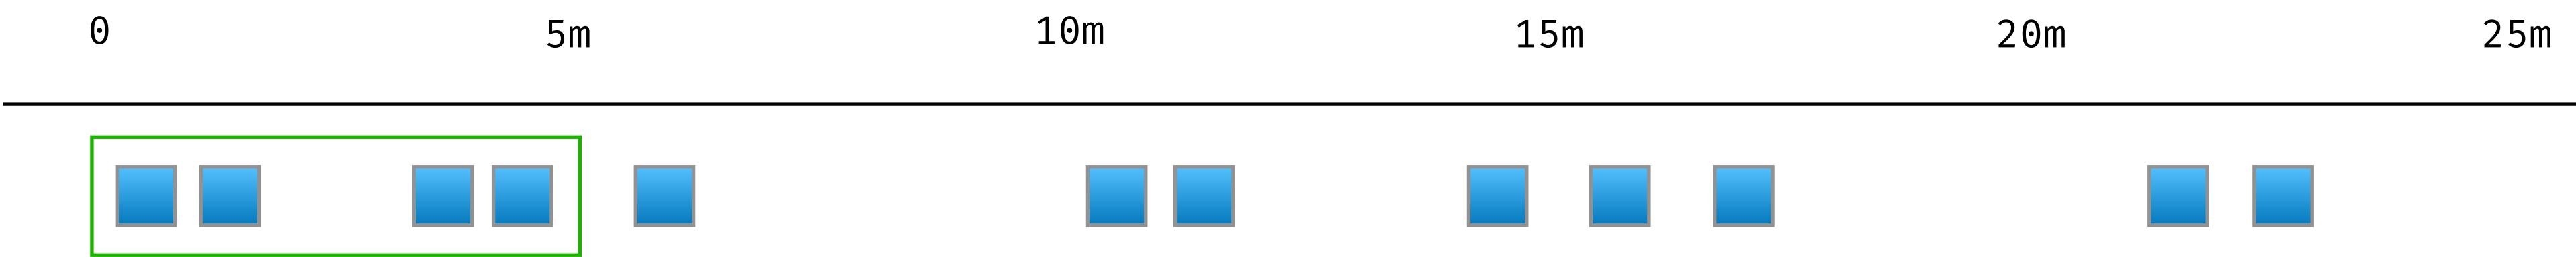


application.id = “my-stream”

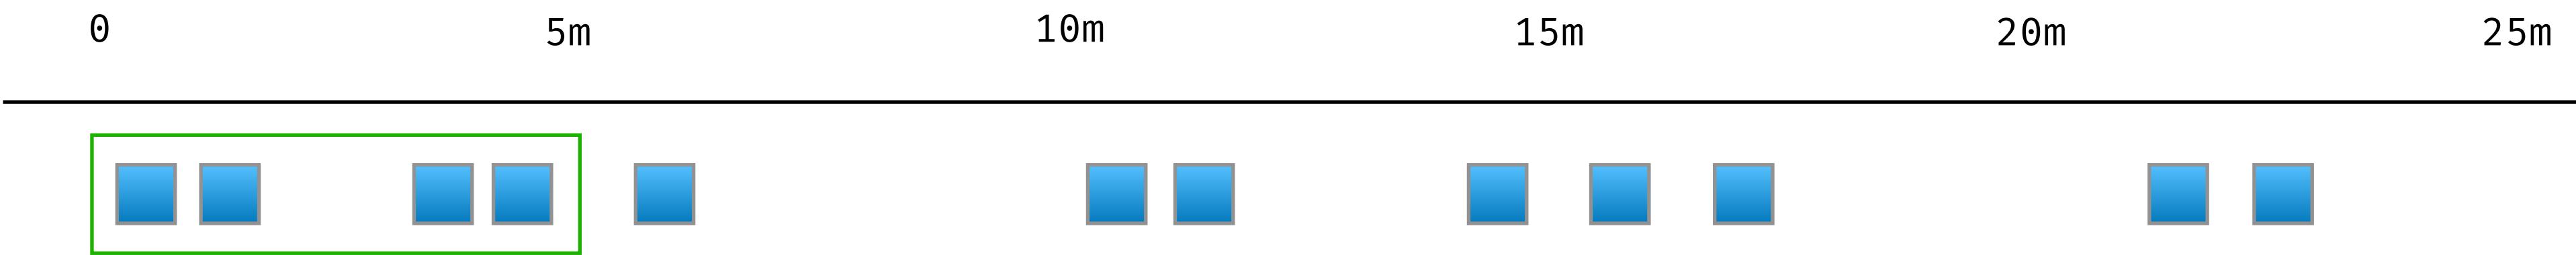


Windowing

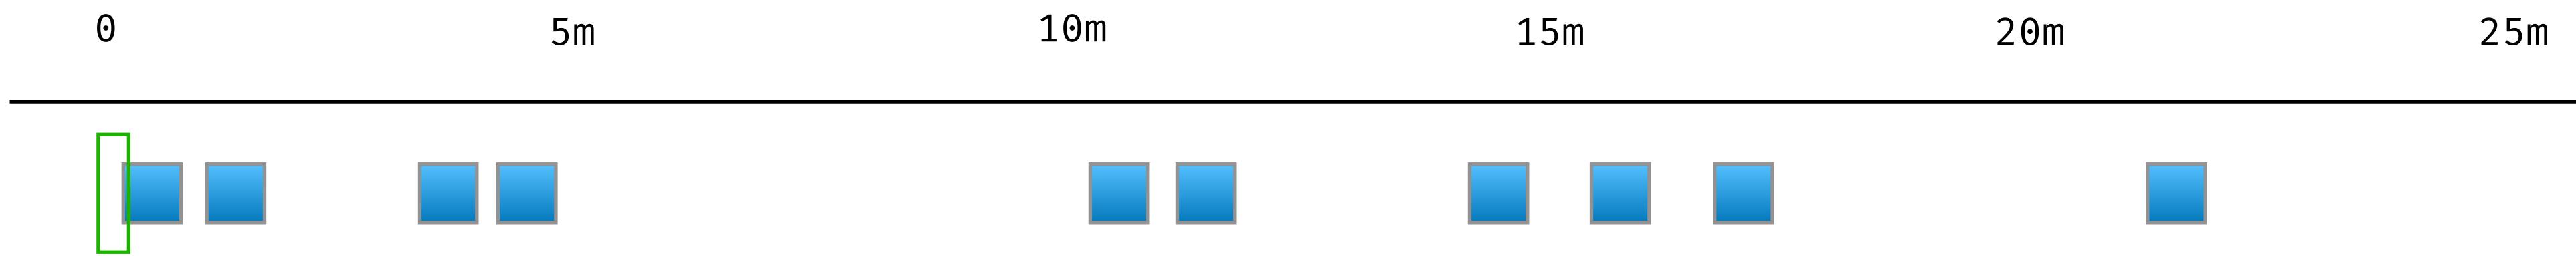
Tumbling Window



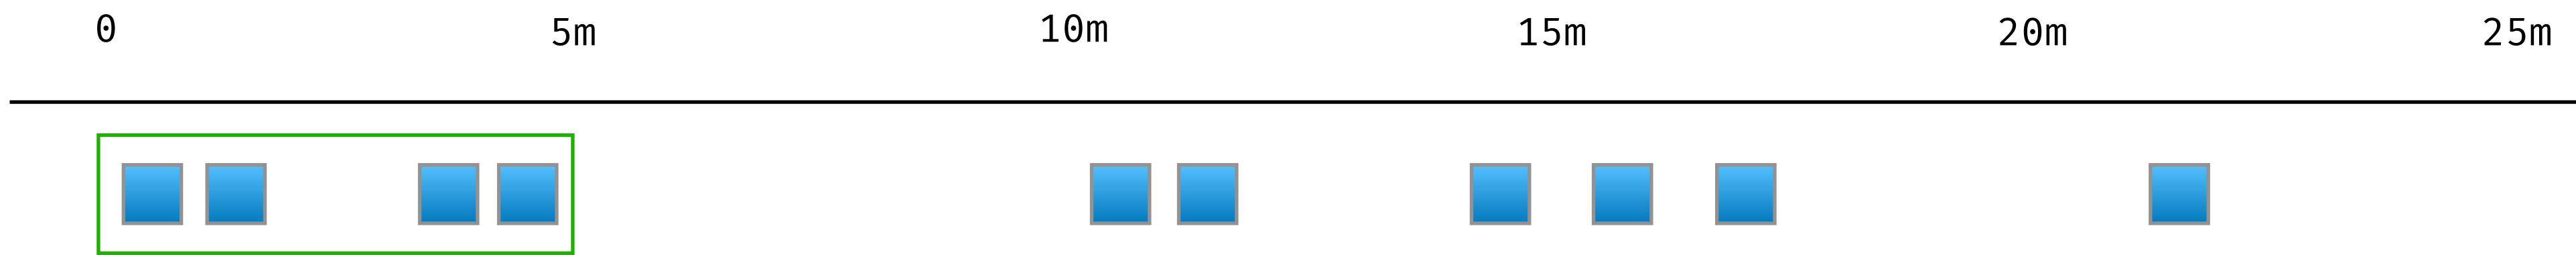
Hopping Window



Session Window

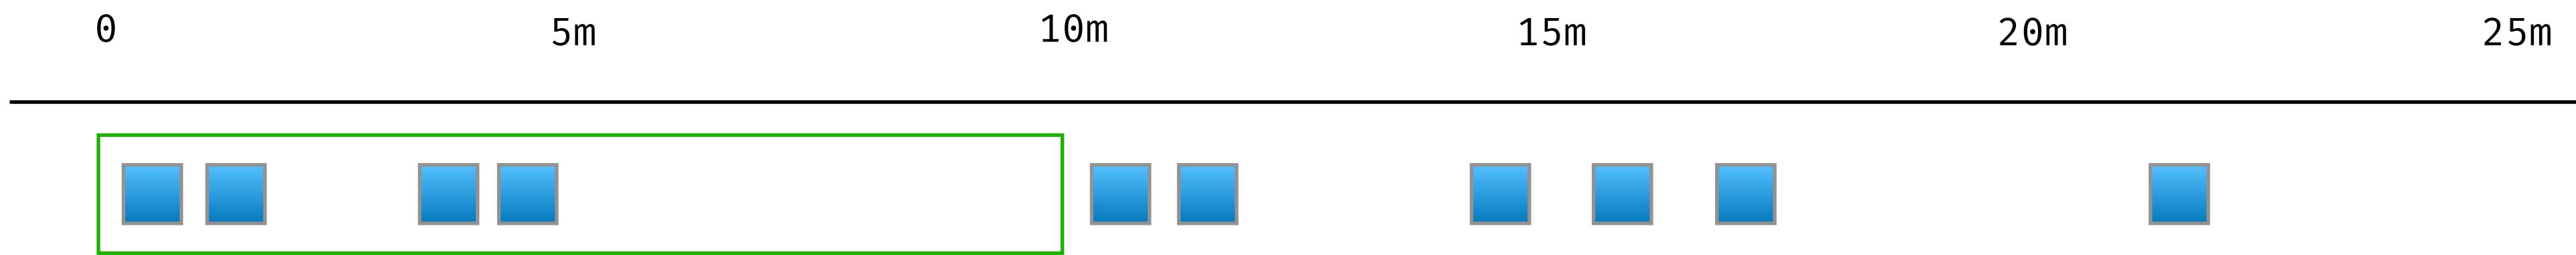


Session Window




the last element has occurred; now we wait

Session Window



waited 5 minutes = make it a window

Session Window



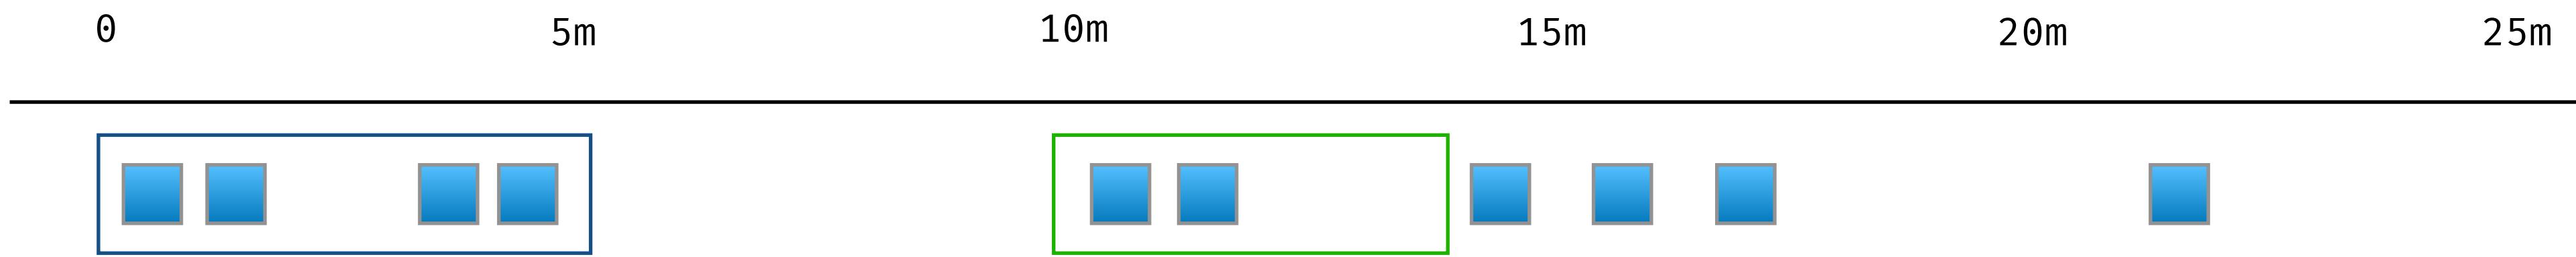
we have a new element; we start a new window

Session Window



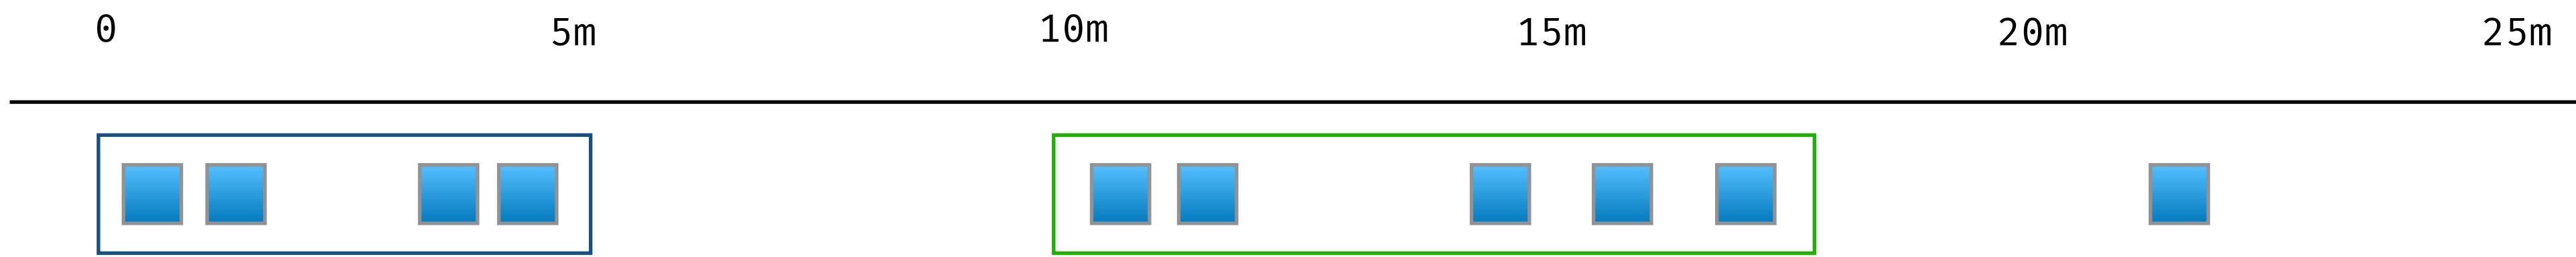
we wait for 5 minutes

Session Window




2 minutes elapsed; no windowing yet

Session Window



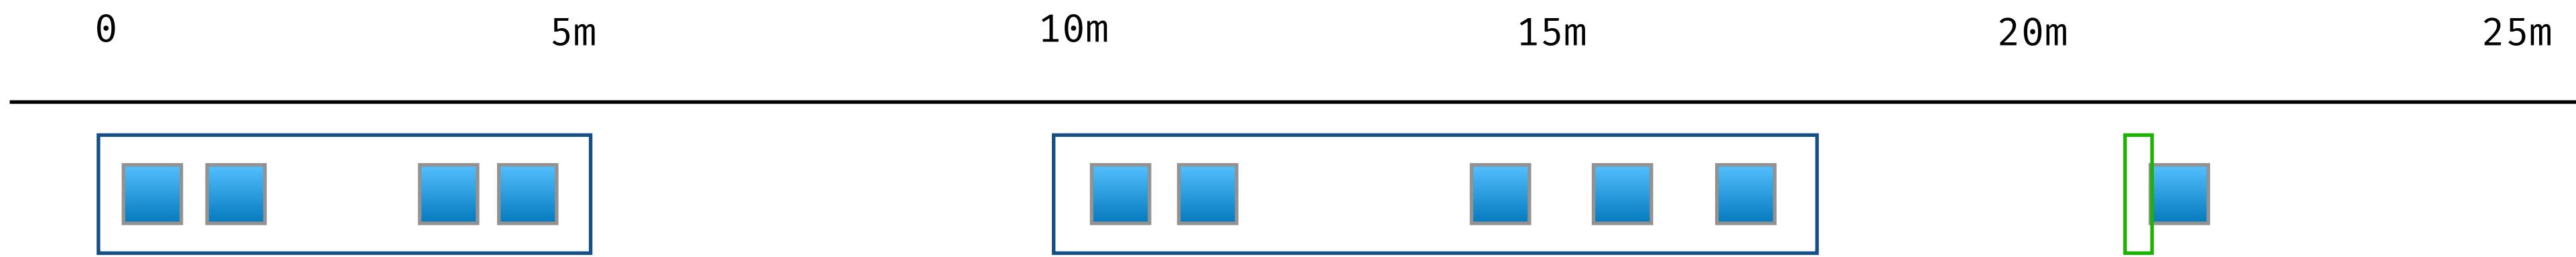
we wait five minutes

Session Window



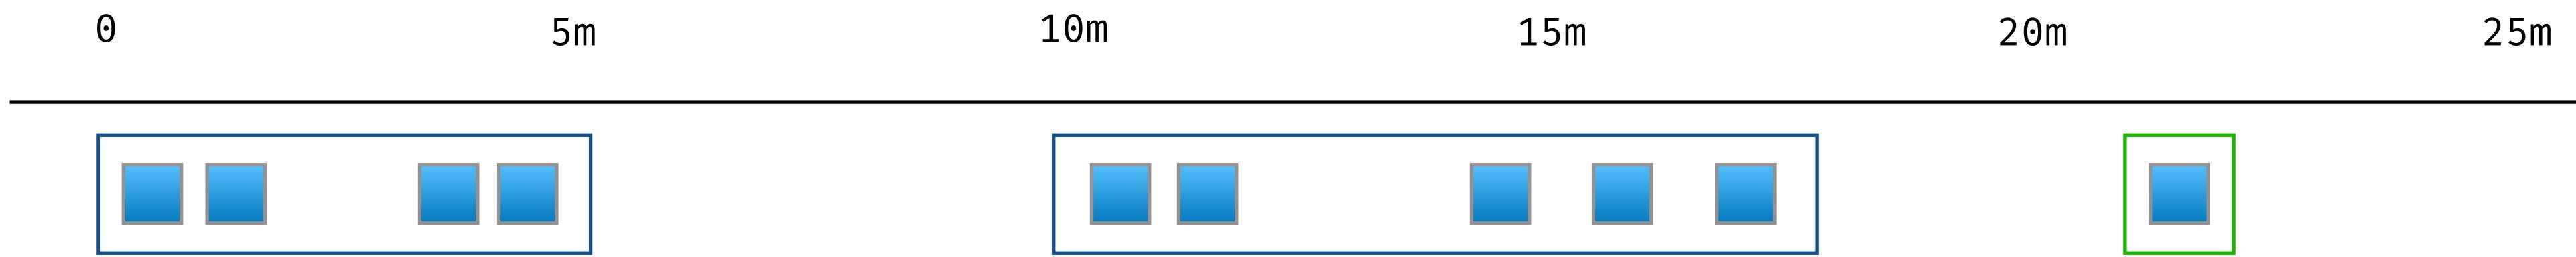
no new elements; we make it a window

Session Window



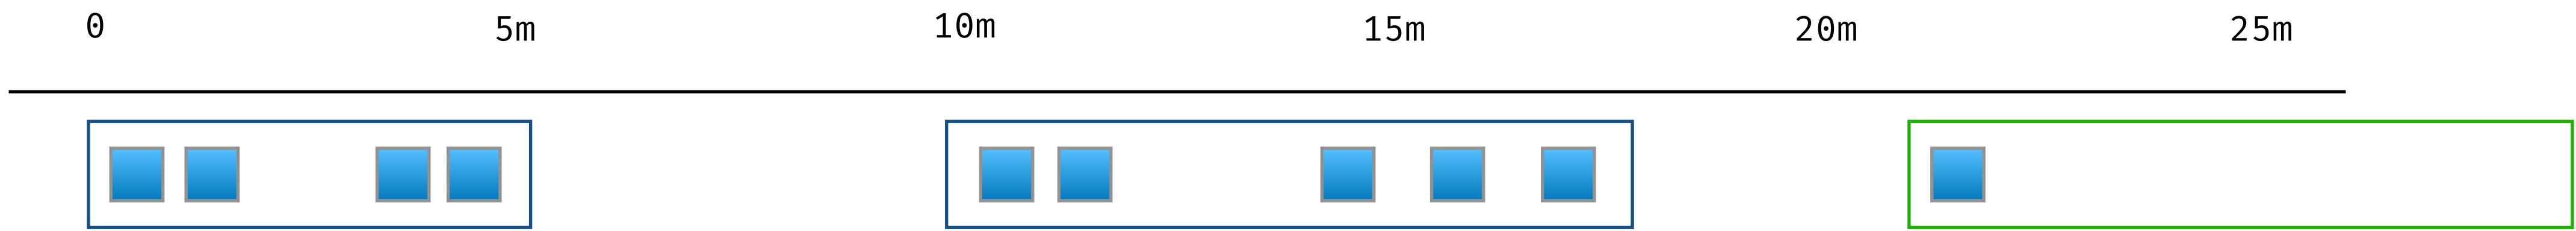
we see something new

Session Window



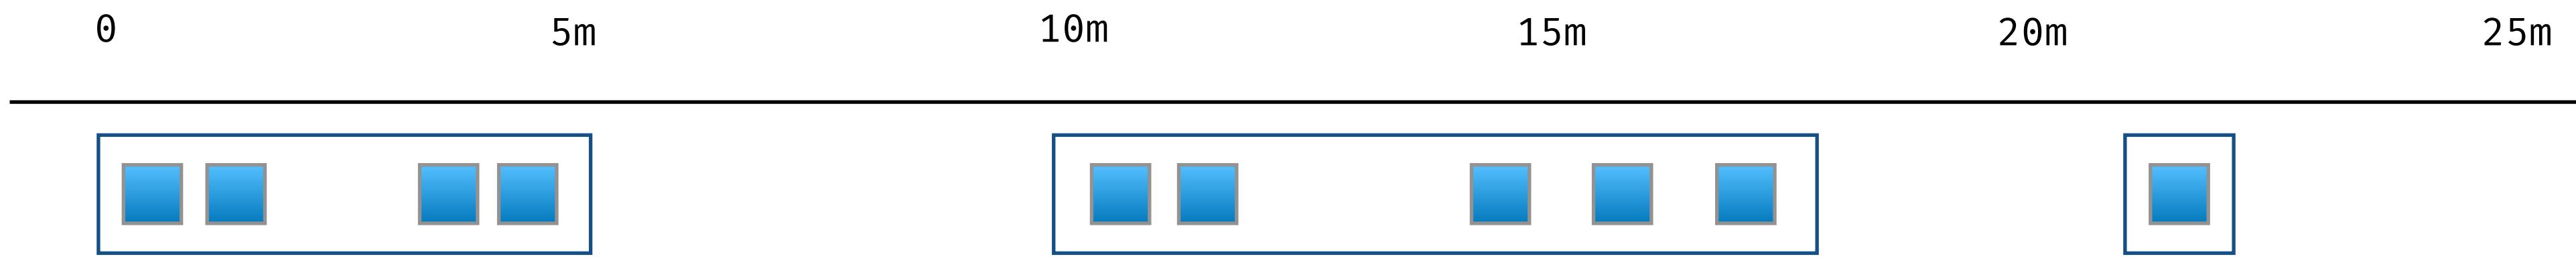
last element; we wait five minutes

Session Window



nothing else; we wait 5 minutes

Session Window



that's a new window

Stream API

Establishing Properties

- Create an `application.id` that represents your application group or "team"
- Serde is combination of Serializer/Deserializer
- Every stream application and KSQL app (later) is a consumer-producer

```
Properties props = new Properties();
props.put(StreamsConfig.APPLICATION_ID_CONFIG, "my_stream_app");
props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(StreamsConfig.DEFAULT_KEY_SERDE_CLASS_CONFIG, Serdes.String().getClass());
props.put(StreamsConfig.DEFAULT_VALUE_SERDE_CLASS_CONFIG, Serdes.Integer().getClass());
```

Create a Stream Builder

- Always start with a StreamBuilder object
- This is the GoF builder pattern, where we will create a Topology object that represents our data pipeline

```
StreamsBuilder builder = new StreamsBuilder();
```

Read from a Topic

- Recall there are two representations of how to ingest data from a topic
 - Stream - Flow of data
 - Table - An aggregation; Or a database
- How do you read the topic is up to you
- Result is either a stream or table object that can take functional programming DSL.

```
KStream<String, Integer> stream = builder.stream("my_topic");
```

```
KTable<String, Integer> table = builder.table("my_topic");
```

Standard Functional Programming

- map
- filter
- flatMap
- groupBy
- reduce
- window
- join
- leftJoin
- outerJoin

Mapping

- Given a Stream

```
(1, "Hello"), (2, "Zoom"), (3, "Fold")
```

- Applying map

```
stream.map((key, value) -> new KeyValue<>(key + 1, value + "!"));
```

- Resulting in

```
(2, "Hello!"), (3, "Zoom!"), (4, "Fold!")
```

Filtering

- Given a Stream

```
(1, "Hello"), (2, "Zoom"), (3, "Fold"), (4, "Past")
```

- Applying filter

```
stream.filter((key, value) -> key % 2 == 0);
```

- Resulting in

```
(2, "Zoom"), (4, "Past")
```

FlatMap

- Given a Stream

```
(1, "One"), (2, "Two"), (3, "Three"), (4, "Four")
```

- Applying flatMap

```
stream.flatMap((key, value) -> Arrays.asList(  
    new KeyValue(key, value),  
    new KeyValue(key * 100, value + " Hundred")))
```

- Resulting in

```
(1, "One"), (100, "One Hundred"), (2, "Two"), (200, "Two Hundred"), (3,  
"Three"), (300, "Three Hundred"), (4, "Four"), (400, "Four Hundred")
```

Peek

- Peers into the Stream and view what goes through

```
stream1.peek((key, value) ->  
    System.out.printf("key: %s, value: %d", key, value));
```

GroupBy

- Given a Stream

```
(1, "One"), (2, "Two"), (3, "Three"), (4, "Four")
```

- Applying groupBy

```
stream.groupBy((key, value) -> key % 2 == 0 ? "Even", "Odd")
```

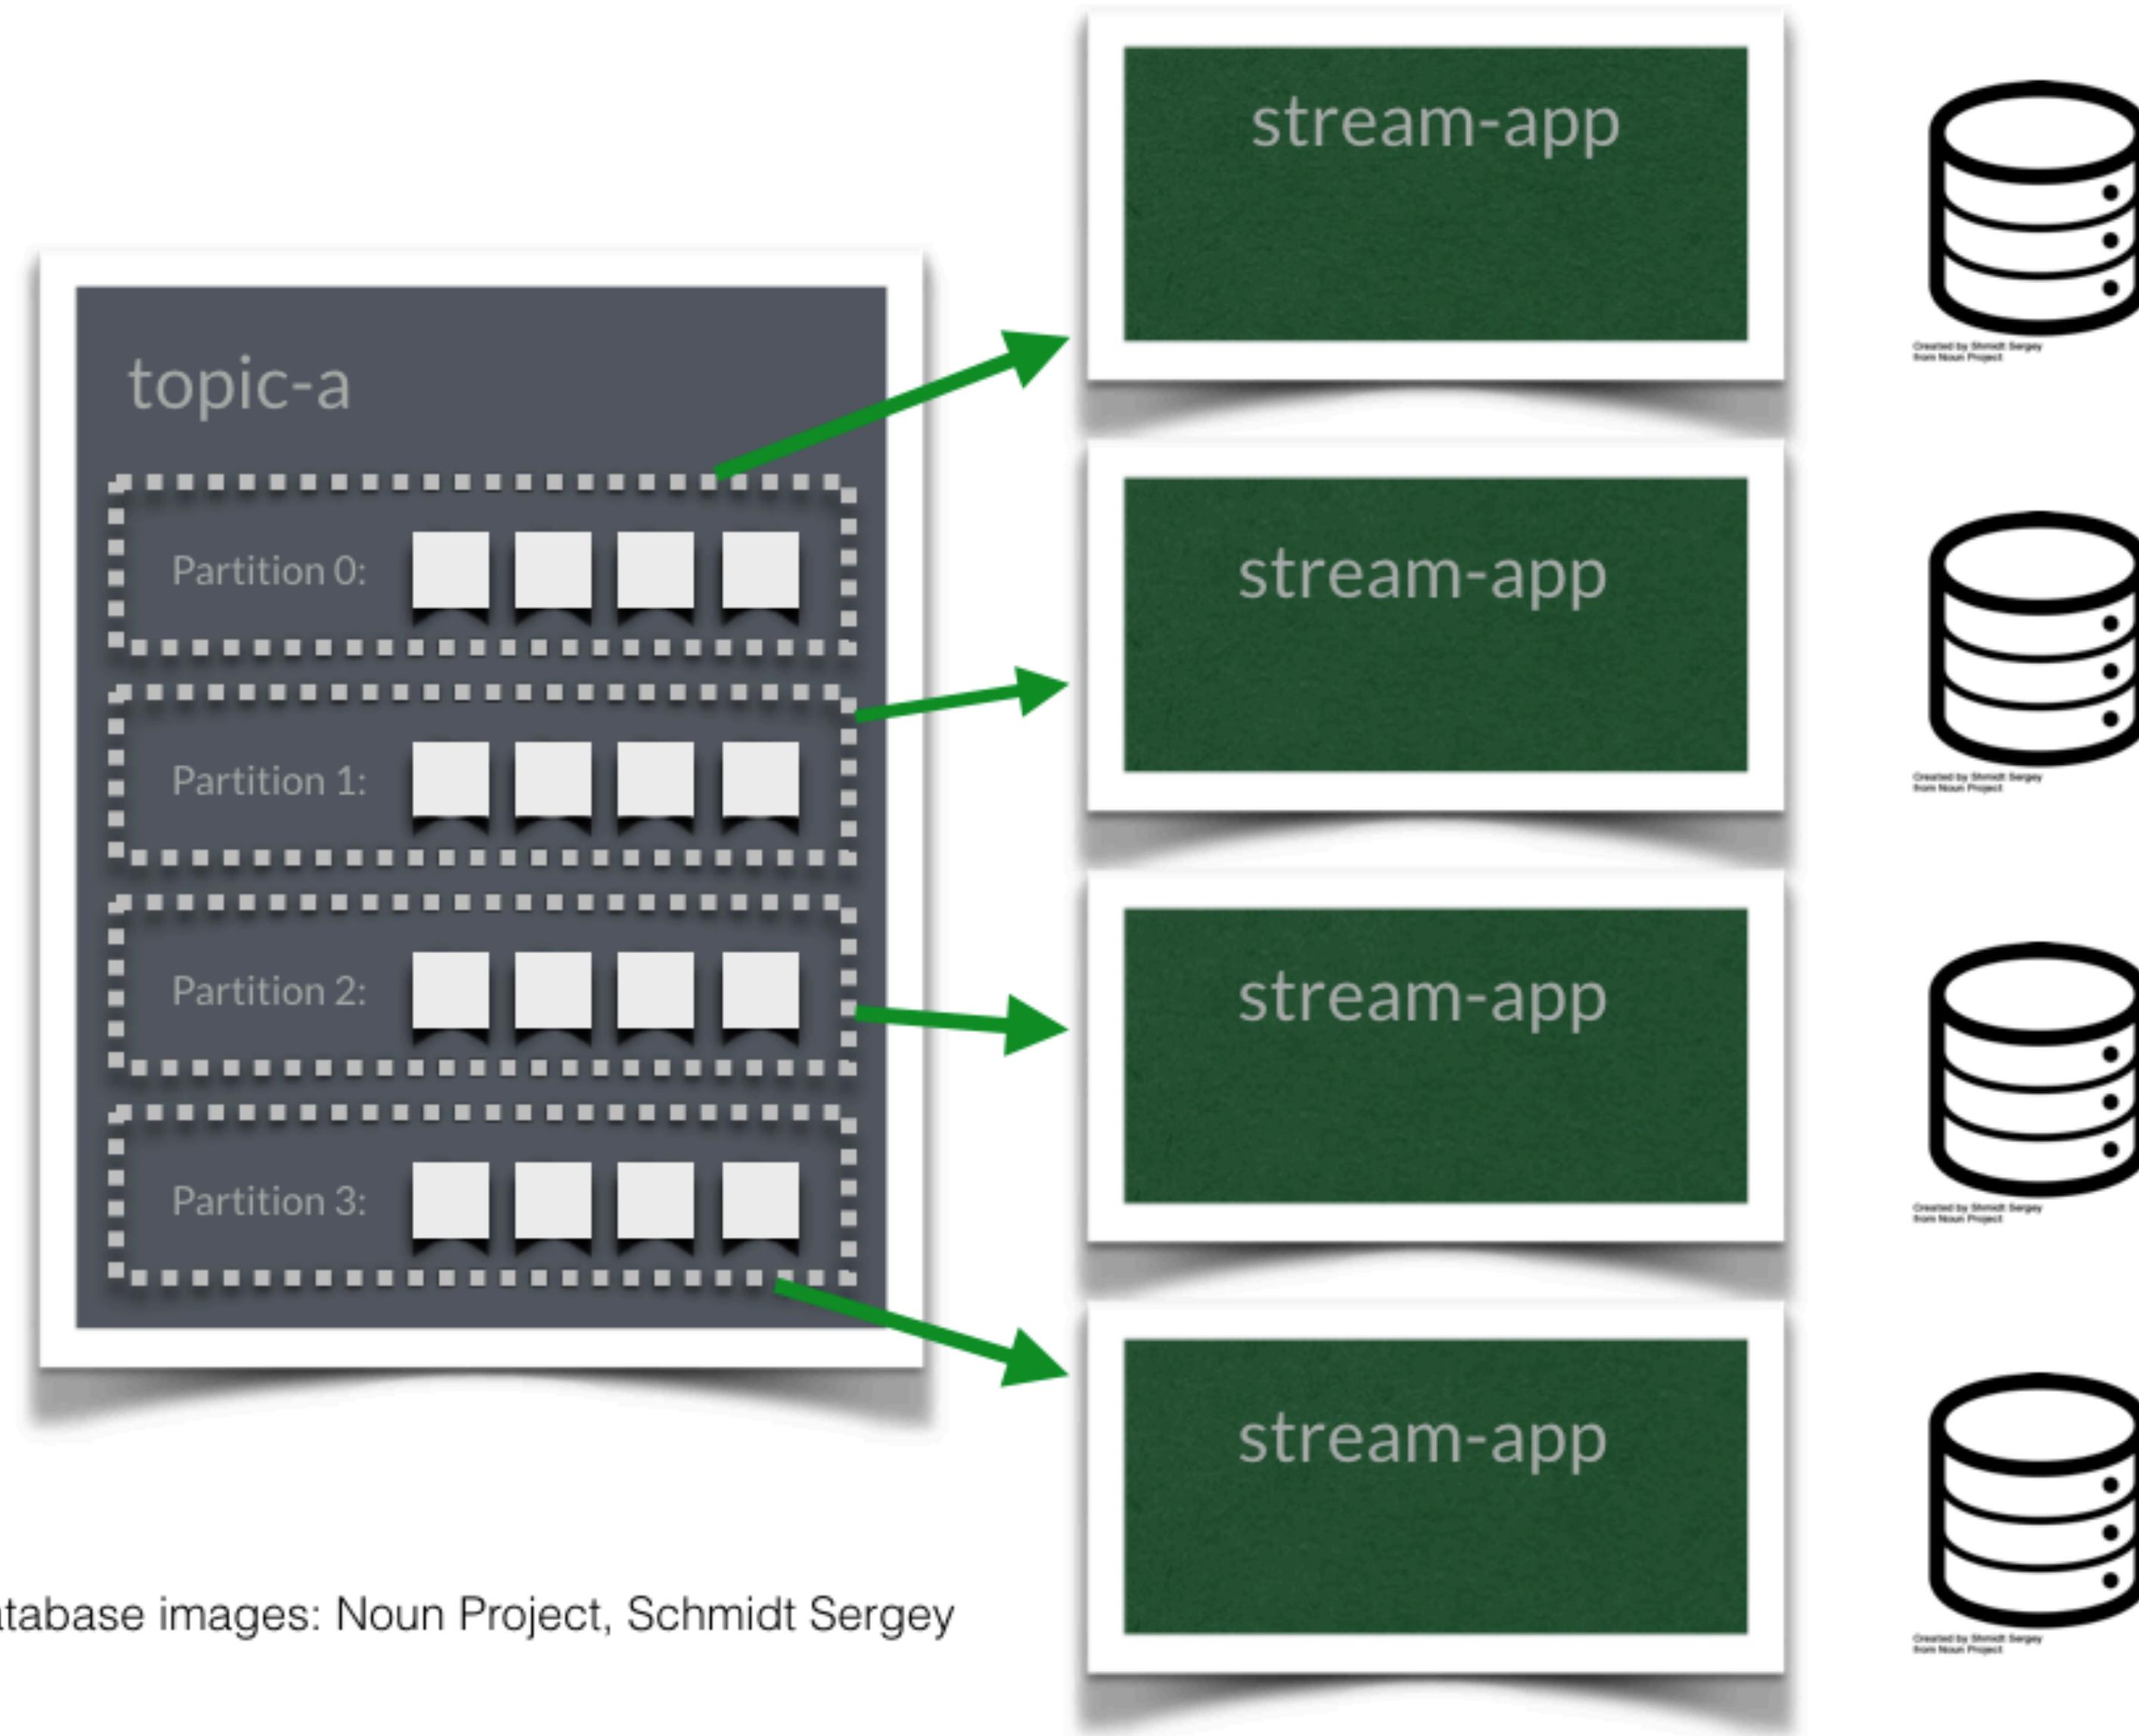
- Resulting in KGroupedStream

Once we have KGroupedStream

- `groupBy` creates a `KGroupedStream` object
- You have to decide what to do with this group
 - `count`
 - `aggregate`
 - `reduce`
 - `windowBy`
- Once you call any of the above, you have a `KTable`

KTable

- Represents a changelog, where the oldest records are not important
- Ideal for counters, and state
- KTable values are stored at each application's ephemeral storage
- Most ephemeral storage is backed by Rocks DB
- All streaming is then backed up as a topic for resilience



Database images: Noun Project, Schmidt Sergey

GroupByKey

- Given a Stream

```
(1, "One"), (2, "Two"), (3, "Three"), (4, "Four")
```

- Applying groupByKey

```
stream.groupByKey()
```

- Resulting in KGroupedStream
- Is Synonymous with the Following

```
stream.groupBy((key, value) -> key)
```

Count

- Aggregates the count of what has been grouped
- Records with null key or value are ignored
- Returns KTable that represents the latest rolling count
- Will maintain the count in ephemeral storage, (like RocksDB)

Count

- Given a Stream

```
(1, "One"), (2, "Two"), (3, "Three"), (4, "Four")
```

- Applying groupBy with count

```
stream.groupBy((key, value) -> key % 2 == 0 ? "Even" : "Odd").count();
```

- Will result in the RocksDB ephemeral storage

```
("Even", 2)  
("Odd", 2)
```

Reduce

- Given a Stream

```
(1, "One"), (2, "Two"), (3, "Three"), (4, "Four")
```

- Applying groupBy with count

```
stream.groupBy((key, value) -> key % 2 == 0 ? "Even" : "Odd")
    .reduce((value1, value2) -> value1 + "," + value2);
```

- Will result in the RocksDB ephemeral storage

```
("Even", "Two,Four")
("Odd", "One,Three")
```

Aggregate

- Given a Stream

```
(1, "One"), (2, "Two"), (3, "Three"), (4, "Four")
```

- Applying groupBy with count

```
groupedStream.aggregate(() -> "0:Zero",
(nextKey, nextValue, aggString) -> aggString + nextKey + ":" + nextValue ",");
```

- Will result in the RocksDB ephemeral storage

```
("Even", "0:Zero,2:Two,4:Four")
("Odd", "0:Zero,1:One,3:Three")
```

Hopping Windows

- Creating a Time Window

```
import java.util.concurrent.TimeUnit;
import org.apache.kafka.streams.kstream.TimeWindows;

long windowSizeMs = TimeUnit.MINUTES.toMillis(5); // 5 * 60 * 1000L
long advanceMs =    TimeUnit.MINUTES.toMillis(1); // 1 * 60 * 1000L
TimeWindows.of(windowSizeMs).advanceBy(advanceMs);
```

- Applying the Window

```
streams.groupBy(...).windowedBy(TimeWindows.of(windowSizeMs).advanceB
y(advanceMs));
```

Session Windows

- Creating a Time Window

```
import java.util.concurrent.TimeUnit;
import org.apache.kafka.streams.kstream.SessionWindows;

SessionWindows.with(TimeUnit.MINUTES.toMillis(5));
```

- Applying the Window

```
streams.groupBy(...).windowedBy(SessionWindows.with(TimeUnit.MINUTES.
toMillis(5)));
```

Dump results to a topic

- Dump the results to a topic using to

```
KStream<String, Integer> stream = builder.stream("my_topic");  
stream.filter(...).to("new_topic");
```

- Dump the results to a topic using through to post to topic and continue:

```
KStream<String, Integer> stream = builder.stream("my_topic");  
stream.filter(...).through("new_topic").flatMap(...).to("other_topic")
```

Build the Topology and Stream

- Once you create your stream, you can build a Topology object:

```
Topology topology = builder.build();
```

- Create a Kafka Streams object with both Topology and Properties you created earlier

```
KafkaStreams streams = new KafkaStreams(topology, props);
```

- And Run...

```
streams.start();
```

Adding a Shutdown Hook

- As always, be a good citizen, properly shutdown resources

```
Runtime.getRuntime().addShutdownHook(new Thread(streams::close));
```

Lab: Streams

KSQLDB

KSQLDB

Open source streaming SQL engine for Apache Kafka

Provides SQL interface for stream processing on Kafka

Scalable, elastic, fault-tolerant, and real-time.

KSQL

Supports a wide range of streaming operations, including:

- Data filtering
- Transformations
- Aggregations
- Joins
- Windowing
- Sessionization

KSQL Server



KSQL CLI

KSQL Architecture

KSQL Server communicates with the broker

KSQL servers are run separately from the KSQL CLI client and Kafka brokers.

You can deploy servers on remote machines, VMs, or containers

KSQL CLI interacts with KSQL Servers via a REST API

KSQL Server runs in two modes: CLI, Headless

KSQLDB

Kafka Streams

Producer/Consumer API

Starting in CLI Mode

```
/bin/ksql http://localhost:8088
```

Starting in Headless Mode

```
/bin/ksql-start-server \  
server_path/etc/ksql/ksql-server.properties \  
--queries-file /path/to/queries.sql
```



Tapping a Stream Based Topic

```
CREATE STREAM pageviews \
(viewtime BIGINT, \
userid VARCHAR, \
pageid VARCHAR) \
WITH (KAFKA_TOPIC='pageviews', \
VALUE_FORMAT='DELIMITED');
```

Tapping a Table Based Topic

```
CREATE TABLE users \
  (registertime BIGINT, \
  userid VARCHAR, \
  gender VARCHAR, \
  regionid VARCHAR) \
WITH (KAFKA_TOPIC = 'users', \
      VALUE_FORMAT='JSON', \
      KEY = 'userid');
```

Showing the Streams

SHOW STREAMS;

Stream Name	Kafka Topic	Format
PAGEVIEWS	pageviews	DELIMITED

Showing the Tables

```
SHOW TABLES;
```

Table Name	Kafka Topic	Format
USERS	users	JSON

Supported SQL Types

- BOOLEAN
- INTEGER
- BIGINT
- DOUBLE
- VARCHAR (or STRING)
- ARRAY<ArrayType> (JSON and AVRO only. Index starts from 0)
- MAP<VARCHAR, ValueType> (JSON and AVRO only)
- STRUCT<FieldName FieldType, ... > (JSON and AVRO only)

Various Operators

Scalar Operators:

ABS, ARRAYCONTAINS, CEIL, CONCAT,
EXTRACTJSONFIELD, FLOOR, GEO_DISTANCE, IFNULL, LC
ASE, LEN, MASK, RANDOM, ROUND,
STRINGTOTIMESTAMP, SUBSTRING, TIMESTAMPTOSTRING,
TRIM, UCASE

Aggregator Operators:

COUNT, MAX, MIN, SUM, TOPK, TOPKDISTINCT, WINDOWS
TART, WINDOWEND

Displays Information

SHOW TOPICS	Show all topics
SHOW <TOPIC>	Display Topic
SHOW STREAMS	Show Streams
SHOW TABLES	Show Tables
SHOW FUNCTIONS	Show Available Functions
SHOW QUERIES	Show Persistent Queries

Describing

DESCRIBE

List columns in stream or table

DESCRIBE EXTENDED

Describe in detail stream or table; runtime statistics, and queries that populate the stream or table

Non-Persistent Pull Query

```
SELECT * FROM pageviews  
WHERE ROWTIME >= 1510923225000  
AND ROWTIME <= 1510923228000;
```

A `LIMIT` can be used to limit the number of rows returned.
Once the limit is reached the query will terminate. This is a pull query since we are want a one time pull of an exact set

Non-Persistent Push Query

```
SELECT * FROM pageviews EMIT CHANGES;
```

The **EMIT CHANGES** is used for Push Queries. When there are new results, they will *pushed* to you. This is the only time, a push exists in the Kafka ecosystem.

Persistent Query

```
CREATE TABLE users_female AS \  
SELECT userid, gender, regionid FROM users \  
WHERE gender='FEMALE' EMIT CHANGES;
```

What makes this a persistent query is the form

```
CREATE (TABLE|STREAM) AS SELECT
```

Once executed, it will continuously run, until terminated
with **TERMINATE** command

Persistent v. Non Persistent

	Persistent	Non Persistent
Where does data go?	Will store into a topic	Will display on screen
How do I stop it?	Find query id using SHOW QUERIES and TERMINATE <id>	CTRL + C
How to create?	CREATE STREAM AS SELECT ...	SELECT

Selecting Customized Fields

```
CREATE STREAM pageviews_transformed \
  WITH (TIMESTAMP='viewtime', \
        PARTITIONS=5, \
        VALUE_FORMAT='JSON') AS \
  SELECT viewtime, \
         userid, \
         pageid, \
         TIMESTAMPTOSTRING \
           (viewtime, 'yyyy-MM-dd HH:mm:ss.SSS') \
             AS timestamp \
  FROM pageviews \
  PARTITION BY userid;
```

Start from the Beginning

```
SET 'auto.offset.reset' = 'earliest';
```

Tumbling Window

```
SELECT item_id, SUM(quantity)  
FROM orders  
WINDOW TUMBLING (SIZE 20 SECONDS)  
GROUP BY item_id;
```

Hopping Window

```
SELECT item_id, SUM(quantity)
FROM orders
WINDOW HOPPING (SIZE 20 SECONDS,
                  ADVANCE BY 5 SECONDS)
GROUP BY item_id;
```

Session Window

```
SELECT item_id, SUM(quantity)  
FROM orders  
WINDOW SESSION (20 SECONDS)  
GROUP BY item_id;
```

Lab: KSQLDB

Conclusion

Agenda

- ★ Understand Kafka
- ★ Understand Producer
- ★ Understand Consumer
- ★ Understand Rebalancing
- ★ Understand Streaming
- ★ Understand Tables
- ★ Understand KSQL